


SME3006 Machine Learning – 2025 Fall

ML Basics and Data preprocessing




INHA UNIVERSITY

Types of ML



Machine learning system roadmap




Terminology



- > Dataset
 - Instance, observation
 - Data sample, sample space
 - Feature, feature vector
- > Learning, training
 - Training data, training sample, training set
 - Validation set, test set
 - Ground-truth
 - Label, label space
- > Model
 - Inputs, input vectors
 - Outputs, targets

Terminology

> MNIST dataset




> CIFAR-10/100 dataset



Terminology


- > Sample space
 - coin tossing: head or tail (0 or 1)
 - MNIST: $256^{28 \times 28}$ (gray)
 - CIFAR10: $256^{32 \times 32 \times 3}$ (RGB)
- > Label space
 - MNIST: {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
 - CIFAR10: {airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck}

Terminology



Terminology

- > Dataset splitting: training / validation / testing



Terminology

> Features

| Features | | | | Target |
|----------|------------|-----|---------|----------------|
| Person | Name | Age | Income | Marital status |
| 1 | Jane Doe | 24 | 81,200 | Single |
| 2 | John Smith | 41 | 121,000 | Married |

ML algorithm:
learns mapping from
features to target




| Features | | | | | | | | Label |
|------------|--------|---------|------|----------|----------------|----------------|---------------|----------|
| date | lat | long | temp | humidity | cloud_coverage | wind_direction | atmp_pressure | rainfall |
| 2021-09-09 | 49.71N | 82.16W | 74 | 20 | 3 | N | 18.6 | .01 |
| 2021-09-09 | 32.71N | 117.16W | 82 | 42 | 6 | SW | 29.94 | .23 |

Example

Terminology

> Features

- nose shape, whisker, pupils, ear shape, muzzle length





Terminology

> Generalization

- Over-training, overfitting
- Underfitting
- Interpolation, extrapolation
- Uncertainty
- Inductive bias

Terminology

> Generalization



Terminology

> Uncertainty






Fig. 1: A schematic view of main differences between aleatoric and epistemic uncertainties.



Terminology

> Inductive bias


- set of assumptions a learning algorithm makes in order to generalize from limited data to unseen cases



Restriction bias: Learning is restricted to a certain class of functions.

Preference bias: Prefer certain functions, but learning isn't restricted to them

Hypothesis Space




Terminology

- > Distribution
 - Independent and Identically distributed (i.i.d)
 - Distribution shift, covariate shift
 - Out-of-distribution (OOD)
- > Accuracy
 - Sensitivity, specificity, precision, recall, F1
 - Mean square error(MSE), mean average error(MAE)
 - Bias-variance tradeoff

Terminology

- > Independent and Identically distributed (i.i.d)



Terminology

> Distribution shift, covariate shift, out-of-distribution





Image Classification

Train




dog

Test




Reinforcement Learning


Train



Test



Terminology




Accuracy = $\frac{TP + TN}{TP + TN + FP + FN}$

Specificity = $\frac{TN}{TN + FP}$

Precision = $\frac{TP}{TP + FP}$

Recall = $\frac{TP}{TP + FN}$



How many retrieved items are relevant?

Precision = $\frac{\text{true positives}}{\text{true positives} + \text{false positives}}$

How many relevant items are retrieved?



Recall = $\frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Terminology



| Epoch | Prediction | Target |
|-------|------------|-----------|
| 1 | [0, 4, 9] | [3, 5, 7] |
| 2 | [2, 4, 2] | [3, 5, 7] |
| 3 | [3, 5, 6] | [3, 5, 7] |

$$MSE = \frac{1}{N} \sum_i^N (pred_i - target_i)^2$$
$$RMSE = \sqrt{\frac{1}{N} \sum_i^N (pred_i - target_i)^2}$$
$$MAE = \frac{1}{N} \sum_i^N |(pred_i - target_i)|$$



Terminology


> Bias-variance tradeoff



Terminology

- > Task
 - Classification
 - Regression
 - Clustering
 - Dimensionality reduction
 - Generative task

Terminology



General principles in ML


- > No Free Lunch Theorem (NFL)
 - There is no universally best learning algorithm for all problems
- > Curse of Dimensionality
 - As the number of dimensions increases, the volume of the space grows exponentially, making the data sparse
- > Occam's Razor
 - Among competing hypotheses, the one with the fewest assumptions should be selected.
- > Bias-Variance Tradeoff
 - Balancing simplicity and flexibility is essential for generalization

Python packages

- > Numpy
 - numerical python.
 - array and matrix operations, linear algebra, random number generation
- > Pandas
 - handling structured data
 - reading CSVs, data cleaning, filtering
- > Matplotlib
 - basic plotting library
- > Sklearn (Scikit-learn)
 - a machine learning library
 - classification, regression, clustering, preprocessing, model evaluation
- > Seaborn, Scipy, TensorFlow, PyTorch, MLflow, Optuna, ...

Numpy basics

```
import numpy as np  
  
array1 = np.array([5,4,9])  
array2 = np.array([[4.1,2.0,6.7], [0.3,9.4,2.2]])  
array3 = np.array([[[1,0, 1],[0, 0, 1]], [[0, 1, 1], [...]], .... ]])  
  
print(array1.shape)  
print(array2.shape)
```



Numpy basics

- > ndarray
 - int8, int16, int32, float16, float32, string, object...
 - but, has to be same format in one array

- > list
 - multiple types can be combined

```
array1 = np.array([1, 2.0, 3.0])
array2 = np.array([1,2,3])

print(array1.dtype)
print(array2.dtype)

list1 = [1, 2, 'hi']

array1_2 = array1.astype('float64')
```

Numpy basics

```
array = np.arange(10)
print(array)
-> [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
array = np.zeros((2,3), dtype='float32')
print(array)
-> [[0. 0. 0.]
     [0. 0. 0.]]
```

```
array = np.ones((1,2), dtype ='int32')
print(array)
-> [[1 1]]
```

Numpy basics

```
array1 = np.arange(20)
array2 = array1.reshape(4,5)
array3 = array2.reshape(-1,2)
array4 = array1[2:5]
array5 = array2[[1,2], [2,3]]
array6 = array1[array1>7]
```

```
array = np.array([3, 1, 9, 5])
sorted = np.sort(array)
index = np.argsort(array)
```

Numpy basics

```
A = np.array([[1, 2, 3], [1, 1, 1], [3, 2, 1]])  
B = np.array([[2, 1, 2], [1, 2, 4], [2, 1, 0]])
```

```
A+B  
A*B  
A@B, np.matmul(A,B)  
np.transpose(A), A.T
```

```
np.linalg.det(A)  
np.linalg.inv(A)
```

```
np.append(A, [[3,4,5]], axis=0)  
np.vstack([A, A])  
np.hstack([A, A])  
A.flatten()  
B = np.expand_dims(A, axis = 0)  
A = np.squeeze(B)
```

```
np.random.rand(2,3)  
np.where(A>1, A, 0)
```

Pandas basics

https://pandas.pydata.org/pandas-docs/stable/user_guide/10min.html

Series: 1D labeled array holding data of any type

DataFrame: 2D data structure that holds data like a table

```
import pandas as pd

s = pd.Series([1, 3, 5, np.nan, 6, 8])
df = pd.DataFrame({"A": 1.0, "B": pd.Timestamp("20250702"), "C": np.array([2]*2),
"D": pd.Categorical(["train", "test"])})


df.head(n) # view the top rows of the frame      # n is the number of rows
df.tail() # view the bottom rows of the frame
df.index # 0, 1
df.columns # A, B, C, D

df.to_numpy()
```

Pandas basics

```
dates = pd.date_range("20130101", periods=6)
df = pd.DataFrame(np.random.randn(6,4), index = dates, columns=list("ABCD"))

df["A"]
df.loc[dates[0]]
df.iloc[3]
df.iloc[3:5, 0:2]
```

Pandas basics

```
df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
df1.loc[dates[0] : dates[1], "E"] = 1
df1

pd.isna(df1)

df1.dropna(how="any")
df.dropna(thresh=4)
df.dropna(subset=["C"])
```

Working with missing data

```
from io import StringIO
csv_data = \
'''A,B,C,D
1.0,2.0,3.0,4.0
5.0, 6.0,,8.0
10.0,11.0,12.0,'''
df = pd.read_csv(StringIO(csv_data))



df.isnull().sum()

df.fillna(5) # replace NA with a scalar value
df.ffill() # fill gaps forward
df.bfill() # fill gaps backward



df.fillna(df.mean()) #
```

Working with missing data

> Interpolation



> Extrapolation



Working with missing data

```
from sklearn.impute import SimpleImputer

imr = SimpleImputer(missing_values= np.nan, strategy='mean') #median,
most_frequent
imr = imr.fit(df.values)
imputed_data = imr.transform(df.values)

from sklearn.impute import KNNImputer

kimr = KNNImputer() # basic n_neighbors = 5
kimr.fit_transform(df.values)

df.interpolate('linear')
df.interpolate('quadratic')
df.interpolate('cubic')
df.interpolate(method='spline', order = 2)
df.interpolate(method='polynomial", order = 3)
```

Categorical data

```
df = pd.DataFrame([['green', 'M', 10.1, 'class2'],
                   ['red', 'L', 13.5, 'class1'],
                   ['blue', 'XL', 15.3, 'class2']])
df.columns = ['color', 'size', 'price', 'classlabel']

size_mapping = {'M':1, 'L':2, 'XL':3}
df['size'] = df['size'].map(size_mapping)

inv_size_mapping = {v: k for k, v in size_mapping.items()}
df['size'].map(inv_size_mapping)
```

Categorical data

- > Label encoding (integer encoding)
 - ['cat', 'dog', 'bird'] -> [0, 1, 2]
 - simple and memory efficient
 - Implies ordinal relationship which can confuse the models
- > One-hot encoding
 - cat: [1, 0, 0]
 - dog: [0, 1, 0]
 - bird: [0, 0, 1]
 - No ordinal assumption
 - High dimensionality (memory and computation cost)

Categorical data

| Color |
|-------|
| Red |
| Green |
| Blue |
| Red |
| Green |
| Red |
| Blue |

Label
encoding

| |
|---|
| 2 |
| 1 |
| 0 |
| 2 |
| 1 |
| 2 |
| 0 |

One-Hot
encoding

| Dummy_r | Dummy_b | Dummy_g |
|---------|---------|---------|
| 1 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| 0 | 0 | 1 |
| 1 | 0 | 0 |
| 0 | 1 | 0 |

Categorical data

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder

order_enc = col_trans = ColumnTransformer([('ord_enc', OrdinalEncoder(dtype = int), ['color'])])
X_trans = col_trans.fit_transform(df)
X_trans

X = df[['color', 'size', 'price']].values
c_transf = ColumnTransformer([('onehot', OneHotEncoder(dtype=int), [0]), ('nothing', 'passthrough', [1, 2])])


c_transf.fit_transform(X)

pd.get_dummies(df[['price', 'color', 'size']])
```

Dataset splitting

> Training, validation, and test sets

- A validation set performs the initial testing on the model as it is being trained



Dataset splitting

> https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

```
from sklearn.model_selection import train_test_split

df_wine = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data', header=None)
X, y = df_wine.iloc[:, 1:].values, df_wine.iloc[:, 0].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0,
stratify = y)
```

Data scaling

> Normalization

- Linear scaling: when the feature is mostly uniformly distributed

$$x' = \frac{x - x_{min}}{x_{max} - x_{min}}$$

- Z-score scaling: when the feature is normally distributed

$$x' = \frac{x - \mu}{\sigma}$$

- Log scaling: when the feature distribution is heavy skewed

$$x' = \log(x)$$

- Clipping: when the feature contains extreme outliers

$$x' = clip(x, min, max)$$

Data scaling

> <https://scikit-learn.org/stable/api/sklearn.preprocessing.html>

```
from sklearn.datasets import load_iris
import pandas as pd


iris = load_iris()
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

from sklearn.preprocessing import StandardScaler, MinMaxScaler


scaler = StandardScaler() # MinMaxScaler, RobustScaler, MaxAbsScaler
scaler.fit(iris_df)
iris_scaled = scaler.transform(iris_df)
```

Visualization


> https://matplotlib.org/stable/plot_types/index.html




`plot(x, y)`




`scatter(x, y)`




`bar(x, height)`




`stem(x, y)`




`fill_between(x, y1, y2)`




`hist(x)`




`boxplot(X)`




`errorbar(x, y, yerr, xerr)`




`quiver(X, Y, U, V)`




`streamplot(X, Y, U, V)`




`imshow(Z)`




`pcolormesh(X, Y, Z)`



`contour(X, Y, Z)`



`contourf(X, Y, Z)`



`hist2d(x, y)`

Recommended reading

- > <https://developers.google.com/machine-learning/crash-course/numerical-data>
- > <https://amueller.github.io/COMS4995-s20/slides/aml-02-matplotlib/#p1>

Reference

- > 알고리즘 중심의 머신러닝 가이드, Chapter 2
- > 머신러닝 교과서 파이토치편, Chapter 4
- > 파이썬 머신러닝 완벽가이드 Chapter 1, 2
- > UC Berkeley, Concise Machine Learning, Chapter 1

- > Python for data analysis — ML Engineering
- > Recap: Data preprocessing — ML Engineering
- > Python for scientific computing — ML Engineering