

SME3006 Machine Learning – 2025 Fall

Linear and Logistic Regression




INHA UNIVERSITY

Supervised learning

> House size vs Price


- We want to predict the prices of other houses as a function of the size



Size	Price
2480	500
560	110
1400	290
1033	200
600	100
...	...
3456	600

Supervised learning

- > Given a training set, to learn a function $h: X \rightarrow Y$
 - h is a good predictor (hypothesis)
- > When the target value is continuous, we call the problem a regression (*when y can take on a small number of discrete values \rightarrow classification)



Linear regression



- > To make our housing example more interesting, let's consider a richer dataset in which we also know the number of bedrooms

Living area	#bedrooms	Prices
2104	3	400
1600	3	330
2400	3	369
1416	2	232
3000	4	540
...

- > Approximate y as a linear function of x
 - $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$
 - θ_i are the parameters (weights)

Linear regression

- > In general, $h(x) = \sum_{i=0}^d \theta_i x_i = \theta^\top x$ where $x_0 = 1$
- > To learn the parameters θ , we define a function that measures, for each value of θ , how close the h is to the corresponding y
 - We define the cost function
 - $J(\theta) = \frac{1}{2} \sum_{i=1}^n (h_\theta(x^{(i)}) - y^{(i)})^2$
 - We want to find θ that minimizes the cost function



The normal equations

- > Matrix form

$$\begin{aligned}x_{11}\theta_1 + \cdots + x_{1d}\theta_d &= y_1 \\x_{21}\theta_1 + \cdots + x_{2d}\theta_d &= y_2 \\\cdots \\x_{n1}\theta_1 + \cdots + x_{nd}\theta_d &= y_n\end{aligned}$$

$$> X = \begin{bmatrix} x_{11} & \cdots & x_{1d} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{nd} \end{bmatrix} = \begin{bmatrix} -(x^{(1)})^\top - \\ \vdots \\ -(x^{(n)})^\top - \end{bmatrix}, \quad y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}, \quad \theta = \begin{bmatrix} \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}$$

$$> \text{Then, } X\theta - y = \begin{bmatrix} h_\theta(x^{(1)}) - y^{(1)} \\ \vdots \\ h_\theta(x^{(n)}) - y^{(n)} \end{bmatrix}$$

$$> \frac{1}{2} \|X\theta - y\|^2 = \frac{1}{2} \sum_{i=1}^n (h_\theta(x^{(i)}) - y^{(i)})^2 = J(\theta)$$

The normal equations

- > minimize $J = \frac{1}{2} \|X\theta - y\|^2 = \frac{1}{2}(X\theta - y)^\top(X\theta - y)$
- > $\frac{d}{d\theta} \frac{1}{2}(X\theta - y)^\top(X\theta - y) = \frac{d}{d\theta} (\theta^\top X^\top X\theta - \theta^\top X^\top y - y^\top X\theta + y^\top y)$ $= \frac{d}{d\theta} \frac{1}{2}(\theta^\top X^\top X\theta - 2y^\top X\theta)$ $= \frac{1}{2}(2X^\top X\theta - 2X^\top y)$ $= X^\top(X\theta - y) = 0$
- > Normal equation: $X^\top X\theta = X^\top y$
 - $\theta = (X^\top X)^{-1}X^\top y$ (pseudo-inverse)

Linear regression

- > To minimize $J(\theta)$, we use search algorithm
- > Start with some ‘initial guess’ and repeatedly changes θ to make $J(\theta)$ smaller

- > Gradient descent:


$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta), \quad \alpha \text{ is the learning rate}$$

$$\begin{aligned}> \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \left[\frac{1}{2} (h_\theta(x) - y) \right]^2 \\&= (h_\theta(x) - y) \frac{\partial}{\partial \theta_j} (h_\theta(x) - y) \\&= (h_\theta(x) - y) \frac{\partial}{\partial \theta_j} (\theta^\top x - y) \\&= (h_\theta(x) - y)x_j\end{aligned}$$

Gradient descent

> Gradient descent:

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$





Gradient descent

- > Gradient descent:


$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

update is simultaneously performed
for all values of $j = 0, \dots, d$



Gradient descent

- > Gradient descent is susceptible to local minima in general



- > For a quadratic cost function, which is convex, gradient descent always converges to the global minimum (assuming the learning rate is not too large)

Gradient descent

- > Batch gradient

$$\theta \leftarrow \theta + \alpha \sum_{i=1}^n \left(y^{(i)} - h_\theta(x^{(i)}) \right) x^{(i)}$$

- > Stochastic gradient

for i = 1 to n, {


$$\theta \leftarrow \theta + \alpha \left(y^{(i)} - h_\theta(x^{(i)}) \right) x^{(i)}$$

}

- > Whereas batch gradient descent has to scan through the entire training set, stochastic gradient descent can start making progress right away
 - However, SGD may never converge to the minimum, and the parameters θ will keep oscillating around the minimum of $J(\theta)$


Gradient descent

Batch Gradient Descent




All training samples

Stochastic Gradient Descent



one training sample

Mini-Batch Gradient Descent



part of training sample

Code

> Dataset generation

```
import numpy as np
import matplotlib.pyplot as plt


np.random.seed(0)

X = 2*np.random.rand(100,1)
y = 4 + 3*X+np.random.randn(100,1)
```

Code

> Normal equation


```
X_b = np.c_[np.ones((100,1)), X]  
theta_best = np.linalg.inv(X_b.T@(X_b))@(X_b.T)@(y)  
  
X_new = np.array([[0], [2]])  
X_new_b = np.hstack((np.ones((2,1)), X_new))  
y_predict = X_new_b.dot(theta_best)
```



Code

> Gradient descent


```
alpha = 0.02 # learning rate # test 0.02, 0.1, 0.4
n_iterations = 1000
m = 100
theta= np.random.randn(2,1)
for iteration in range(n_iterations):
    gradients = 1/m * X_b.T@(X_b@theta-y)
    theta = theta - alpha * gradients
```



Code

> Stochastic gradient descent

```
alpha = 0.02; n_epochs = 50
theta = np.random.randn(2,1) # random initialization
for epoch in range(n_epochs):
    for i in range(m):
        random_index = np.random.randint(m)
        xi = X_b[random_index:random_index+1]
        yi = y[random_index:random_index+1]
        gradients = 2*xi.T@(xi@theta -yi)
        theta = theta - alpha * gradients
```




Probabilistic interpretation

- > Why linear regression with the least-squares cost function?
- > In perspective of probability, assume that
 - $y^{(i)} = \theta^\top x^{(i)} + \epsilon^{(i)}$
 - where ϵ is an error that captures unmodeled effects or random noise
 - we assume that ϵ is i.i.d. Gaussian
 - $\epsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$
- > Gaussian distribution and its density function
 - $p(\epsilon^{(i)}) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(\epsilon^{(i)})^2}{2\sigma^2}\right)$

Probabilistic interpretation

> Gaussian (Normal) distribution



$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

$$p(D|\theta) \propto p(\theta|D)p(D)$$


posterior likelihood prior

$$f(\mu|x=a) = \mathcal{N}(a, \sigma^2)$$

- Why Gaussian?
 - Central limit theorem
 - tractability (linear combination, conditional distribution, marginal distribution)
 - maximum entropy

Probabilistic interpretation

- > $p(\epsilon^{(i)}) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(\epsilon^{(i)})^2}{2\sigma^2}\right)$
- > $p(y^{(i)}|x^{(i)}; \theta) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right)$
 - distribution of $y^{(i)}$ given $x^{(i)}$ and θ .
 - we call it likelihood $L(\theta) = L(\theta; X, y) = p(y|X; \theta)$



- > By the assumption of independence of ϵ , $L(\theta) = \prod_{i=1}^n p(y^{(i)}|x^{(i)}; \theta)$
- > Given this probabilistic model, what is a reasonable way of choosing θ ?
- > The principle of **maximum likelihood** says that we should choose θ so as to make the data as high probability as possible (i.e., we should maximize $L(\theta)$)

Probabilistic interpretation

- > Instead of maximizing $L(\theta)$, we maximize $\log(L(\theta))$ – log likelihood
 - why? because it make the problem simple
- >
$$\begin{aligned} l(\theta) &= \log(L(\theta)) \\ &= \log \prod_{i=1}^n p(y^{(i)} | x^{(i)}; \theta) \\ &= \log \prod_{i=1}^n \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right) \\ &= \log \sum_{i=1}^n \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y^{(i)} - \theta^\top x^{(i)})^2}{2\sigma^2}\right) \\ &= n \log \frac{1}{\sigma\sqrt{2\pi}} - \frac{1}{2\sigma^2} \sum_{i=1}^n (y^{(i)} - \theta^\top x^{(i)})^2 \end{aligned}$$

————— —————

constant Least squares cost function
- > Therefore, maximizing $l(\theta)$ is equivalent to minimizing LS cost function

Logistic regression


- > Regression can be used for classification
 - binary classification using logistic regression
 - use linear regression algorithm to predict y given x
 - It will generate values larger than 1 or smaller than 0 while $y \in \{0,1\}$
 - To fix this, we change the form for our hypotheses

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

which is called the logistic function or the sigmoid function

Logistic regression

- > Regression can be used for classification
 - binary classification using logistic regression



Logistic regression

- > Logistic function $g(z) = \frac{1}{1+e^{-z}}$
 - $g(z) \rightarrow 1$ as $z \rightarrow \infty$
 - $g(z) \rightarrow 0$ as $z \rightarrow -\infty$
- > Other functions that smoothly increase from 0 to 1 can also be used, but the choice of the logistic function is fairly natural

$$\begin{aligned}> g'(z) &= \frac{d}{dz} \frac{1}{1+e^{-z}} \\&= \frac{1}{(1+e^{-z})^2} (e^{-z}) \\&= \frac{1}{(1+e^{-z})^2} \left(1 - \frac{1}{1+e^{-z}}\right) \\&= g(z)g(1-z)\end{aligned}$$

Logistic regression

- > Likewise, fit θ for the logistic regression model
 - $p(y = 1|x; \theta) = h_\theta(x)$
 - $p(y = 0|x; \theta) = 1 - h_\theta(x)$
 - $p(y|x; \theta) = (h_\theta(x))^y (1 - h_\theta(x))^{1-y}$
- > Assume that the n training examples were generated independently,
 - $$\begin{aligned} L(\theta) &= p(y|X; \theta) \\ &= \prod_{i=1}^n p(y^{(i)}|x^{(i)}; \theta) \\ &= \prod_{i=1}^n (h_\theta(x^{(i)}))^{y^{(i)}} (1 - h_\theta(x^{(i)}))^{1-y^{(i)}} \end{aligned}$$
 - $$l(\theta) = \log(L(\theta)) = \sum_{i=1}^n y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)}))$$

Logistic regression

> To do gradient ascent, we take the derivative of $l(\theta)$


$$\begin{aligned} - \frac{\partial}{\partial \theta_j} l(\theta) &= \left(y \frac{1}{g(\theta^\top x)} - (1-y) \frac{1}{1-g(\theta^\top x)} \right) \frac{\partial}{\partial \theta_j} g(\theta^\top x) \\ &= \left(y \frac{1}{g(\theta^\top x)} - (1-y) \frac{1}{1-g(\theta^\top x)} \right) g(\theta^\top x)(1-g(\theta^\top x)) \frac{\partial}{\partial \theta_j} \theta^\top x \\ &= \left(y(1-g(\theta^\top x)) - (1-y)g(\theta^\top x) \right) x_j \\ &= (y - h_\theta(x))x_j \end{aligned}$$

$$\begin{aligned} - \theta_j &\leftarrow \theta_j + \alpha \nabla_{\theta_j} l(\theta) \\ \theta_j &\leftarrow \theta_j + \alpha \left(y^{(i)} - h_\theta(x^{(i)}) \right) x_j^{(i)} \end{aligned}$$

- If we compare this to the LMS update rule, it looks identical;
but this is not the same algorithm because h_θ is a nonlinear function

Another method: Newton's method

- > Our objective: find θ that maximizes $l(\theta)$
- > Newton's method: finding a zero of a function (numerical optimization)
 - To find $f(x) = 0$, $x \leftarrow x - \frac{f(x)}{f'(x)}$
 - Approximating the function f via a linear function




Another method: Newton's method

- > Our objective: find θ that maximizes $l(\theta)$
 - The maxima of l corresponds to points where its first derivative $l'(\theta)$ is 0
 - To find $l'(\theta) = 0$, $\theta \leftarrow \theta - \frac{l'(x)}{l''(x)}$
 - In multi-dimensional setting $\theta \leftarrow \theta - H^{-1}\nabla_{\theta} l(\theta)$,
- where H is hessian, $H = \frac{\partial^2 l(\theta)}{\partial \theta_i \partial \theta_j}$
- > Newton's method typically converges faster than gradient descent
 - however, one iteration is more expensive
 - still, it is usually much faster overall
 - When Newton's method is applied to maximize the logistic regression log likelihood function $l(\theta)$, the resulting method is called Fisher scoring

Generalized linear models

> Linear regression $p(y|X) \sim \mathcal{N}(\theta^\top X, \sigma^2)$


- variance is constant
- mean is a linear combination of features
- distribution is Gaussian



Generalized linear models

> $p(y|X) \sim \mathcal{N}(\theta^T X, \sigma^2)$

- What if the distribution is not normal, but some other distribution?
- What if X has a more sophisticated relationship with the mean?
- What if the variance varies with X ?



Generalized linear models

> $p(y|X) \sim \mathcal{N}(\theta^T X, \sigma^2)$

- What if the distribution is not normal, but some other distribution?
 - change the normal distribution to some other distribution from the exponential family of distributions

Exponential Family of Distributions

Gaussian

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\|x-\mu\|^2/(2\sigma^2)} \quad x \in \mathbb{R}$$

Bernoulli

$$p(x) = \alpha^x (1 - \alpha)^{1-x} \quad x \in \{0, 1\}$$

Binomial

$$p(x) = \binom{n}{x} \alpha^x (1 - \alpha)^{n-x} \quad x \in \{0, 1, 2, \dots, n\}$$

Multinomial

$$p(x) = \frac{n!}{x_1!x_2!\dots x_n!} \prod_{i=1}^n \alpha_i^{x_i} \quad x_i \in \{0, 1, 2, \dots, n\}, \sum_i x_i = n$$

Exponential

$$p(x) = \lambda e^{-\lambda x} \quad x \in \mathbb{R}^+$$

Poisson

$$p(x) = \frac{e^{-\lambda}}{x!} \lambda^x \quad x \in \{0, 1, 2, \dots\}$$

Generalized linear models

> $p(y|X) \sim \mathcal{N}(\theta^\top X, \sigma^2)$

■ What if X has a more sophisticated relationship with the mean?

- $\mu(x) = \sum_{j=1}^d \theta_d x_d \rightarrow g(\mu(x)) = \sum_{j=1}^d \theta_d x_d$
- g is called the link function (nonlinear)

Generalized linear models

> $p(y|X) \sim \mathcal{N}(\theta^T X, \sigma^2)$

■ What if the variance varies with X ?

- variance depends on the mean

$$\text{var}(y) = \phi V(\mu) \text{ where } \phi \text{ is a constant}$$

- variance function is determined by the choice of distribution

- Gaussian: $V(\mu) = 1$

- Binomial (logistic regression): $V(\mu) = \mu(1 - \mu)$

- Poisson: $V(\mu) = \mu$

Generalized linear models

> Exponential family

- it includes many of the most common distributions (normal, gamma, beta, Bernoulli, Poisson, exponential, chi-squared, Dirichlet, categorical, ...)
 - $p(y; \eta) = b(y) \exp(\eta^\top T(y) - a(\eta))$
 - $\eta = \theta^\top x$
-
- η is a natural parameter (canonical parameter)
 - $T(y)$ is the sufficient statistics (in many times, $T(y) = y$)
 - $a(\eta)$ is the log partition function (essentially, normalization constant)
 - $b(y)$ is the base measure

Generalized linear models

> Exponential family

- log-likelihood:
$$l(\theta) = \log(\prod_{i=1}^n b(y) \exp(\eta^\top T(y) - a(\eta))) \\ = \sum_{i=1}^n \log(b(y) \exp(\eta^\top T(y) - a(\eta)))$$
- log-likelihood gradient:
$$\frac{\partial}{\partial \theta_j} l(\theta) = \frac{\partial}{\partial \theta_j} \sum_{i=1}^n \eta^\top T(y) - a(\eta) \\ = \sum_{i=1}^n T(y) \frac{d\eta}{d\theta} - a'(\eta) \frac{d\eta}{d\theta} \\ = \sum_{i=1}^n \underbrace{(T(y) - a'(\eta))}_\text{residual} \underbrace{x}_\text{input}$$


Code

> Sigmoid function

```
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

def sigmoid(z):
    return 1.0 / (1.0 + np.exp(-z))

z = np.arange(-7, 7, 0.1)
sigma_z = sigmoid(z)
plt.plot(z, sigma_z)
plt.show()
```



Code

> Logistic regression -1

```
class LogisticRegressionGD:  
    def __init__(self, alpha = 0.01, n_iter = 50, random_state = 1):  
        self.alpha = alpha  
        self.n_iter = n_iter  
        self.random_state = random_state  
  
    def net_input(self, X):  
        return X@self.w_ + self.b_  
  
    def activation(self, z):  
        return 1./(1. + np.exp(-np.clip(z, -250, 250)))  
  
    def predict(self, X):  
        return np.where(self.activation(self.net_input(X)) >= 0.5, 1, 0)
```

Code

> Logistic regression -2

```
def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w_ = rgen.normal(loc=0.0, scale = 0.01, size = X.shape[1])
    self.b_ = np.float64(0.)
    self.losses_ = []
    for i in range(self.n_iter):
        net_input = self.net_input(X)
        output = self.activation(net_input)
        errors = (y-output)
        self.w_ += self.alpha *2.0 * X.T@errors / X.shape[0]
        self.b_ += self.alpha *2.0 *errors.mean()
        loss = (-y@np.log(output))- ((1-y)@np.log(1-output))/X.shape[0]
        self.losses_.append(loss)
    return self
```

Code

> Load iris dataset

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler


iris = datasets.load_iris()
X = iris.data[:, [2,3]]
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
random_state = 1, stratify = y)
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```


Code

> Load iris dataset


iris setosa



iris versicolor



iris virginica



Samples


(instances, observations)

	Sepal length	Sepal width	Petal length	Petal width	Class label
1	5.1	3.5	1.4	0.2	Setosa
2	4.9	3.0	1.4	0.2	Setosa
...					
50	6.4	3.5	4.5	1.2	Versicolor
...					
150	5.9	3.0	5.0	1.8	Virginica

Features

(attributes, measurements, dimensions)

Petal



Sepal

Class labels
(targets)


Code

- > Perform the logistic regression

```
X_train_01_subset = X_train_std[(y_train == 0) | (y_train==1)]
y_train_01_subset = y_train[(y_train == 0) | (y_train == 1)]
lrgd = LogisticRegressionGD(alpha = 0.3, n_iter = 1000, random_state = 1)
lrgd.fit(X_train_01_subset, y_train_01_subset)
plot_decision_regions(X = X_train_01_subset, y = y_train_01_subset,
                      classifier = lrgd)
plt.tight_layout()
plt.show()
```

Code

- > Perform the logistic regression



Reference

- > CS229 Lecture Notes, Stanford
- > Generalized linear model
 - <https://www.dailydoseofds.com/generalized-linear-models-glms-the-supercharged-linear-regression/>
 - Heather Turner lecture note, University of Warwick
- > Linear regression code: 핸즈온 머신러닝 Ch.4
- > logistic regression code: 머신러닝 교과서 파이토치편 Ch.3

- > Further reading
 - <https://bookdown.org/robback/bookdown-BeyondMLR/ch-poissonreg.html>