

SME3006 Machine Learning – 2025 Fall

# Optimization and gradient descent



INHA UNIVERSITY

# Overview

- > Optimization
  - gradient
  - hessian
  - constraint
  - newton's method
- > Convex
  - LP, QP
- > Gradient descent
  - momentum
  - 2<sup>nd</sup> order model
  - stochastic optimization

# Optimization problem

- > Finding the minimizer of a function subject to constraints:

$$\underset{x}{\text{minimize}} J(x)$$

objective function


$$\begin{aligned} s.t. \quad & f_i(x) \leq 0, \quad i = \{1, \dots, k\} \\ & h_j(x) = 0, \quad j = \{1, \dots, l\} \end{aligned}$$

constraints


- Linear regression:  $\underset{w}{\text{minimize}} \|Xw - y\|^2$
- Classification (logistic regression):  $\underset{w}{\text{minimize}} \sum_{i=1}^n \log(1 + \exp(-y_i x_i^\top w))$
- Maximum likelihood estimation:  $\underset{\theta}{\text{maximize}} \sum_{i=1}^n \log p_\theta(x_i)$
- K-means:  $\underset{w}{\text{minimize}} J(\mu) = \sum_{j=1}^k \sum_i \|x_i - \mu_j\|^2$
- ...

# Optimization problem


> Constraints = defining feasible sets



$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad f(x)$$



$$\begin{aligned} &\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad f(x) \\ &\text{subject to} \quad h_i(x) = 0, \quad i = 1, \dots, n_h \end{aligned}$$




$$\begin{aligned} &\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad f(x) \\ &\text{subject to} \quad g_i(x) \geq 0, \quad i = 1, \dots, n_g \end{aligned}$$

# Gradient

- > Real-valued scalar function  $f(x): \mathbb{R}^n \rightarrow \mathbb{R}$

$$> \text{Gradient vector: } \nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \dots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$


- > Tangent line (hyperplane) at point  $(x^*, f(x^*))$ 
  - $y = f(x^*) + \nabla f(x^*)^\top (x - x^*)$
  - $\nabla f(x^*) = 0 \rightarrow x^*$  : maximum or minimum



# Hessian



$$f(x) \rightarrow \nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix} \rightarrow \nabla^2 f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

- Hessian matrix describes the local curvature of the function



# Optimality conditions

- > Unconstrained optimization
  - minimize  $f(x)$
  - What are the conditions for  $x$  to be a minimum point?
  - $\nabla f(x^*) = 0$
  - $\nabla^2 f(x^*) \geq 0$
  - $\nabla^2 f(x^*) > 0$



# Optimality conditions

## > Unconstrained optimization


- minimize  $f(x)$
- Necessary condition:  $\nabla f(x) = 0$ ,  $\nabla^2 f(x^*) \geq 0$
- Sufficient condition:  $\nabla f(x) = 0$ ,  $\nabla^2 f(x^*) > 0$
- Find the roots of  $\nabla f(x)$
- If obtained root satisfies  $\nabla^2 f(x^*) > 0$ , point  $x^*$  is local/global minimum

# Optimality conditions

## > Constrained optimization

- minimize  $f(x)$


*subject to*  $g(x) \geq 0$



- Optimality conditions of unconstrained optimization are not valid for constrained optimization
- It will lead to KKT (Karush-Kuhn-Tucker) necessary optimality condition (graduate level)

# Optimality conditions


- > Even in an unconstrained problem, we don't know whether the minimum point is local or global
- > How can we find a global minimum?



# Optimality conditions


## Nonconvex Optimization

- Multiple local minima ⭐
- Sensitive to initial point ⚡



## Convex Optimization

- Unique minimum: global/local




## Convex sets


> Definition:

A set  $C \subseteq \mathbb{R}^n$  is convex if for  $x, y \in C$  and any  $\alpha \in [0,1]$ ,

$$\alpha x + (1 - \alpha)y \in C$$



convex






non-convex

# Convex sets

## > Examples

- norm ball  $\{x: \|x\| \leq r\}$
- hyperplane  $\{x: a^T x = b\}$
- halfspace  $\{x: a^T x \leq b\}$
- affine space  $\{x: Ax = b\}$
- polyhedron  $\{x: Ax \leq b\}$




> Arbitrary intersection of convex sets is a convex set

## Convex function

> Definition:

A function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is convex if for  $x, y \in \text{dom } f$  and any  $\alpha \in [0,1]$ ,

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$




# Convex function

## > First order convexity condition:

- Suppose  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is differentiable  
 $f$  is convex iff for all  $x, y \in \text{dom } f$


$$f(y) \geq f(x) + \nabla f(x)^T(y - x)$$



## > Second order convexity condition:

- Suppose  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is twice differentiable  
 $f$  is convex iff for all  $x \in \text{dom } f$

$$\nabla^2 f(x) \succeq 0.$$




## Convex sets and convex functions

- > Epigraph of a function  $f$  is the set of points

$$\text{epi } f = \{(x, t) : f(x) \leq t\}$$

- >  $\text{epi } f$  is convex if and only if  $f$  is convex
- > Sublevel sets,  $\{x : f(x) \leq a\}$  are convex for convex  $f$



# Convex optimization

> An optimization problem is convex


- if its objective is a convex function
- inequality constraints  $f_i(x)$  are convex
- equality constraints  $h_j(x)$  are affine ( $h_j(x) = b^\top x + c$ )

$$\underset{x}{\text{minimize}} J(x) \quad \text{convex function}$$

$$\begin{aligned} s.t. f_i(x) \leq 0, \quad i = \{1, \dots, k\} \\ h_j(x) = 0, \quad j = \{1, \dots, l\} \end{aligned} \quad \begin{array}{l} \text{convex sets} \\ \text{affine} \end{array}$$

## Convex optimization

- > If  $x^*$  is a local minimizer of a convex optimization problem, then it is a global minimizer
- >  $\nabla f(x) = 0$  if and only if  $x$  is a global minimizer of  $f(x)$



## Convex optimization

- > If  $x^*$  is a local minimizer of a convex optimization problem, then it is a global minimizer
- >  $\nabla f(x) = 0$  if and only if  $x$  is a global minimizer of  $f(x)$ 
  - Proof.
  - $\nabla f(x) = 0$ 
    - we have  $f(y) \geq f(x) + \nabla f(x)^\top (y - x) = f(x)$
  - $\nabla f(x) \neq 0$ 
    - There is a direction of descent

# Class of convex optimizations

- > Linear program (LP)

- $\underset{x}{\text{minimize}} \quad c^T x$

$$\begin{aligned} s.t. \quad & x \geq 0, \\ & Ax = b \end{aligned}$$

- > Quadratic program (QP)

- $\underset{x}{\text{minimize}} \frac{1}{2} x^T P_0 x + q_0^T x + c_0$

$$\begin{aligned} s.t. \quad & Ax = b \\ & x \geq 0 \end{aligned}$$

- > Second-order cone program (SOCP)
- > Semidefinite program (SDP)
- > Cone program
- > ...

## Gradient descent

- > The simplest optimization algorithm

Optimization goal:  $\underset{x}{\text{minimize}} J(x)$


- > Just iterate

$$x_{t+1} = x_t - \alpha_t \nabla f(x_t)$$

where  $\alpha_t$  is step size (or learning rate)

# Gradient descent

- >  $x_{t+1} = x_t - \alpha_t \nabla f(x_t)$
- > Gradient descent is steepest descent method
  - direction  $\nabla f(x)$  gives greatest reduction in  $f(x_t)$  per unit change in  $x$
  - learning rate determines how far we move in that direction



## Gradient descent

- > 1<sup>st</sup> order Taylor series for  $f(x)$  around current  $x$  is

$$f(x + d) \approx f(x) + \nabla f(x)^T d$$

- for small enough  $d$ , this will be a reasonable approximation
- GD update computed by minimizing this within a sphere of radius  $r$


$$-\alpha_t \nabla f(x_t) = \arg \min_{d: \|d\| \leq r} f(x) + \nabla f(x)^T d$$

- > If  $f(x)$  is sufficiently smooth, and learning rate is small, gradient will keep pointing down-hill over the region




# Challenges

- > Choosing a proper learning rate can be difficult
  - a small learning rate lead slow convergence, while a large learning rate hinder convergence
  - a small learning rate may get stuck in local minima



# Challenges

- > Choosing a proper learning rate can be difficult
  - simple solution is scheduling the learning rate
  - $\alpha_t = \alpha_0 \beta^t, \beta \in [0,1]$
  - reducing the learning rate according to a predefined schedule  
(have to be defined in advance, thus unable to adapt to a dataset)






## Challenges

- > Choosing a proper learning rate can be difficult
  - the same learning rate applies to all parameter updates
  - if our data is sparse and our features have very different frequencies, we might not want to update all of them to the same extent, but perform a larger update for rarely occurring features
  - $x \leftarrow x - \alpha \nabla f(x) \Rightarrow x^i \leftarrow x^i - \alpha^i \nabla f(x^i)$

# Challenges


- > Gradient direction is not optimal
  - Steepest gradient is not the most efficient way to the minimum
  - e.g., 2D narrow valley example



# Challenges

- > Gradient direction is not optimal
  - GD minimizes the following primitive local 2<sup>nd</sup>-order approximation to  $h(\theta)$
  - smooth convex function

$$f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|x - y\|^2$$



## Challenges

- > Gradient direction is not optimal
  - GD minimizes the following primitive local 2<sup>nd</sup>-order approximation to  $h(\theta)$

$$\begin{aligned}f(\theta + d) &\approx f(\theta) + \nabla f(\theta)^T d + \frac{1}{2} d^T F(\theta) d \\&\approx f(\theta) + \nabla f(\theta)^T d + \frac{1}{2} d^T L I d = f(\theta) + \nabla f(\theta)^T d + \frac{L}{2} \|d\|^2\end{aligned}$$

- $\arg \min_d h(\theta) + \nabla h(\theta)^T d + \frac{L}{2} \|d\|^2 = -\frac{1}{L} \nabla f(\theta)$
- LI is a very conservative / pessimistic approximation to  $F(\theta)$  that treats all directions as having the same (very high) curvature
  - this explains why it struggles on problems where curvature varies a lot

# Challenges


- > Gradient direction is not optimal

- GD minimizes the following primitive local 2<sup>nd</sup>-order approximation to  $h(\theta)$
  - smooth convex function

$$f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|x - y\|^2$$

- strongly convex function

$$f(y) \geq f(x) + \nabla f(x)^\top (y - x) + \frac{\mu}{2} \|x - y\|^2$$




## Convergence theory

- >  $f(x)$  has Lipschitz continuous derivatives (Lipschitz smooth)
  - $\|\nabla f(x) - \nabla f(x')\| \leq L\|x - x'\|$  (upper bound on the curvature)
- >  $f(x)$  is strongly convex (perhaps only near minimum)
  - $f(x + d) \geq f(x) + \nabla f(x)^\top d + \frac{\mu}{2}\|d\|^2$  (lower bound on the curvature)
- > We take  $\alpha = \frac{2}{L+\mu}$ ,  $f(x_t) - f(x^*) \leq \frac{L}{2} \left(\frac{\kappa-1}{\kappa+1}\right)^{2t} \|x_0 - x^*\|^2$  where  $\kappa = \frac{L}{\mu}$
- > Number of iterations to achieve  $f(x_t) - f(x^*) \leq \epsilon$  is

$$t \in \mathcal{O}\left(\kappa \log \frac{1}{\epsilon}\right)$$

## Momentum method

- > Motivation: gradient has a tendency to flip back and forth as we take steps when the learning rate is large




- > Key idea: accelerate movement along directions that point consistently down-hill across many consecutive iterations

$$x_{t+1} = x_t + \alpha_t v_{t+1}$$


$\gamma$  is momentum constant  
(usually 0.9)

$$v_{t+1} = \gamma v_t - \nabla f(x_t), \quad v_0 = 0$$


# Momentum method




Gradient descent with large learning rate



Gradient descent with small learning rate




Momentum method




# Nesterov's momentum

> Intuition: look ahead before you leap!




## Simple convex function performance


>  $f(x, y) = x^2 + 10y^2$



(a) SGD with learning rate 0.08




(b) SGD with learning rate 0.09




(c) SGD with learning rate 0.10

# Simple convex function performance


>  $f(x, y) = x^2 + 10y^2$




(a) Momentum with momentum factor 0.7



(b) Momentum with momentum factor 0.9



(c) NAG with momentum factor 0.7



(d) NAG with momentum factor 0.9

## Convergence theory


- > A first-order method is one where updates are linear combinations of observed gradients.
  - i.e.,  $x_{k+1} - x_k = d \in \text{span}\{\nabla f(x_0), \nabla f(x_1), \dots, \nabla f(x_k)\}$
  - gradient descent
  - momentum methods
  - conjugate gradients
- > To achieve  $f(x_k) - f(x^*) \leq \epsilon$ , the number of iterations  $k$  satisfies:
  - (worst-case) lower bound for 1<sup>st</sup>-order methods:  $k \in \Omega\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$
  - upper bound for gradient descent:  $k \in \mathcal{O}\left(\kappa \log \frac{1}{\epsilon}\right)$
  - upper bound for GD w/ Nesterov's momentum:  $k \in \mathcal{O}\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$

## 2<sup>nd</sup>-order method

- > For any 1<sup>st</sup>-order method, the number of steps needed to converge grows with condition number

$$\kappa = \frac{L}{\mu}$$

Max curvature  
Min curvature





- > 2<sup>nd</sup> –order methods can improve this dependency
  - approximate  $f(x)$  by its 2<sup>nd</sup> order Taylor series around current  $x$

$$f(x + d) \approx f(x) + \nabla f(x)^\top d + \frac{1}{2} d^\top H(x) d$$

## 2<sup>nd</sup>-order method


- > Minimize the 2<sup>nd</sup>-order approximation
  - $f(x + d) \approx f(x) + \nabla f(x)^\top d + \frac{1}{2} d^\top H(x) d$
  - $d = -H(x)^{-1} \nabla f(x)$
- > Update current iterate with this:
  - $x_{t+1} = x_t - H(x)^{-1} \nabla f(x_t)$



## Newton's method

- > Newton's method: finding a zero of a function (numerical optimization)

- To find  $f(x) = 0$ ,  $x \leftarrow x - \frac{f(x)}{f'(x)}$




- The maxima of  $l$  corresponds to points where its first derivate  $l'(\theta)$  is 0
- To find  $l'(\theta) = 0$ ,  $\theta \leftarrow \theta - \frac{l'(\theta)}{l''(\theta)}$
- In multi-dimensional setting  $\theta \leftarrow \theta - H^{-1}\nabla_{\theta} l(\theta)$ ,

## Comparison to gradient descent

- > GD:  $x_{t+1} = x_t - \alpha \nabla f(x_t)$ 
  - To avoid divergence,  $\alpha \leq 1/L$   
 $L$  is maximum curvature (Lipschitz constant)
- > Newton's method:  $x_{t+1} = x_t - H(x)^{-1} \nabla f(x_t)$
- >  $H(x)$  corresponds to  $LI$ 
  - $f(x + d) \approx f(x) + \nabla f(x)^\top d + \frac{1}{2} d^\top H(x) d$   
$$\approx f(x) + \nabla f(x)^\top d + \frac{1}{2} d^\top L I d$$
  - GD implicitly minimizes a bad approximation of 2<sup>nd</sup>-order Taylor series
  - GD treats all directions as having max curvature

## 2<sup>nd</sup>-order method



- > Quadratic approximation of objective is only trustworthy in a local region around current  $x$ 
  - Constrain update  $d$  to lie in a trust region  $R$  around where approximation remains good enough
- > Or we use a better quadratic approximation instead of using  $H(x)$ 
  - Generalized Gauss-Newton matrix (GCN), Fisher information matrix, ...



## 2<sup>nd</sup>-order method

### > Disadvantages

- For neural networks,  $x \in \mathbb{R}^n$  can have millions or billions of dimensions
- We can't compute and store  $n \times n$  matrix
- To use 2<sup>nd</sup>-order methods, we must simplify the curvature matrix's
  - computation
  - storage
  - and inversion
- This can be done by approximating the matrix with a simpler form
  - such as using only the diagonal entries
  - or using only the diagonal blocks



## 2<sup>nd</sup>-order method

- > Many optimizers retain some benefits of curvature-aware optimization without the full computational burden
  - We may call it quasi-2<sup>nd</sup> methods
  - use first-order information (gradients) only, but implicitly incorporate curvature information (typically through the averages)
- > Most modern optimizers (used in deep learning) are quasi-2<sup>nd</sup> order
  - Adagrad, Adadelta, RMSProp, Adam, AdamW, Nadam, AMSGrad

## Optimizer: Adagrad

- > It adapts the learning rate to the parameters
  - smaller update to frequently occurring features
  - larger update to infrequent features
  - well-suited for dealing with sparse data
  - gradient:  $g_{t,i} = \nabla J(x_{t,i})$
  - update:  $x_{t+1,i} = x_{t,i} - \frac{\alpha}{\sqrt{G_{t,ii} + \epsilon}} g_{t,i}$
  - $G_t$  is a diagonal matrix where its diagonal entry  $i$  is the sum of the squares of the gradients w.r.t.  $x_i$  up to time step  $t$   
(contains the past gradients)
  - Learning rate is autonomously reduced for each parameter
  - Weakness:  $G_t$  is keeps accumulating, causing the learning rate to keep shrinking


## Optimizer: RMSProp

- > To resolve Adagrad's radically diminishing learning rates
  - restricts the window of accumulated past gradients
- > Instead of inefficiently storing  $w$  previous squared gradients, the sum of gradients is recursively defined as a decaying average
  - $E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma)g_t^2$       (exponentially moving average)
  - $x_{t+1,i} = x_{t,i} - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} g_{t,i}$
- > Hinton suggests  $\gamma$  to be set to 0.9, while a good default value for  $\alpha = 0.001$

## Optimizer: Adam

- > Adaptive momentum estimation (Adam)
  - In addition to storing an EMA of past squared gradients, Adam keeps an EMA of past gradients, similar to momentum
  - $E[g^2]_t = \gamma_1 E[g^2]_{t-1} + (1 - \gamma_1)g_t^2$
  - $m_t = \gamma_2 m_{t-1} + (1 - \gamma_2)g_t$
  - $m_t$  and  $E[g^2]_t$  are biased towards zero initially, thus correct it
  - $\hat{m}_t = \frac{m_t}{1 - \gamma_2^t}$ ,  $\hat{E}[g^2]_t = \frac{E[g^2]_t}{1 - \gamma_1^t}$
  - $x_{t+1,i} = x_{t,i} - \frac{\alpha}{\sqrt{\hat{E}[g^2]_t + \epsilon}} \hat{m}_t$
  - Authors propose default values of  $(0.999, 0.9)$  for  $(\gamma_1, \gamma_2)$



# Optimizers



출처 : <https://www.slideshare.net/yongho/ss-79607172>

# Optimizers


## > Convergence performance



# Optimizers

## > Which optimizer to use?

- Adam / AdamW are dominant and default choices
- In cases where another optimizer did better than Adam, it was usually RMSProp or NAG
- Adafactor, Lion, Signum are promising for LLMs
- SGD (+momentum) is good for vision tasks, but requires careful tuning



# Reference

- > An overview of gradient descent optimization algorithms
  - <https://www.ruder.io/optimizing-gradient-descent/>
- > Simple performance example
  - <https://tiddler.github.io/optimizers/>
- > Lecture notes
  - Theoretical derivation and mathematics:
    - optimization in machine learning – Simon Weissmann
    - optimization for machine learning – Bernd Gartner and Martin Jaggi