

SME3006 Machine Learning – 2025 Fall

Dimensionality reduction and PCA




INHA UNIVERSITY

Why dimension reduction?

- > high-resolution images
 - 1080p: $1920 * 1080$ (pixels) * 3 (RGB)
= 6,200,800 dims
 - Raw data ~ 6Mb
 - image compression:
PNG: 1~3Mb, JPEG: 0.1~0.5Mb

- > language tokens
 - One-hot encoding
 - Dog = [1, 0, 0, ...], Cat = [0, 1, 0, ...],
Rabbit = [0, 0, 1, ...] $x \in \mathbb{R}^{\dim(V)}$
 - Oxford English Dictionary $\dim(V) \approx 5\sim 600,000$
 - 표준국어대사전 $\dim(V) \approx 4\sim 500,000$




Why dimension reduction?

> Good things

- memory friendly, computationally efficient
- extract only meaningful data → performance improvement
- avoid the curse of dimensionality
- generalizability (prevent overfitting)
- enable visualization (2D, 3D)

> Bad things


- might lose some (important) information
- might reduce interpretability
- require extra steps to process the data



Why dimension reduction?

> Application


- Communication



- Robot control / autonomous driving




- Bio informatics




Feature extraction

- > Ex) distinguishing a red ball and a blue ball
 - image inputs ($32 \times 32 \times 3 = 3072$ dims)



- features: position (x, y), radius (r), color (a.k.a latent vector)
 - [13, 23, 3]
 - [14, 8, 5]
 - [14, 18, 5]
 - [18, 21, 7]

- Auto-encoder structure (embedding and reconstruction)



Feature extraction

- > A feature can be task-specific (e.g., classification)
 - an individual measurable property or characteristics of a data sample


- car : wheels, windows, doors, headlights, ...



- dog : legs, tail, fur texture, ears, nose, ...






- Are they still cars?



Feature extraction


> Auto-encoder: minimize $_{\theta} \|x - g(f(x))\|^2$



- There is no labels
- It is an unsupervised learning
- objective: while preserving the essential structure and information of the data, learn low-dimensional representations
- dimension reduction → feature extraction → representation learning

Linear dimension reduction




- > Map high-dim. data into a low-dim. space using linear transformations
 - $z = W^T x$, where $W \in \mathbb{R}^{D \times d}$, $d < D$
 - W is called a projection matrix
 - recall linear projection (from linear algebra)



$$Q = [\mathbf{u}_1 \quad \dots \quad \mathbf{u}_k]$$
$$\mathbf{v}_{\parallel} = Q(Q^T Q)^{-1} Q^T \mathbf{v}$$



Linear dimension reduction

- > linear projection onto 1-dimension






Linear dimension reduction

- > linear projection onto 1-dimension




Linear dimension reduction

- > linear projection onto 2-dimension



Linear dimension reduction

- > linear projection
 - if I had to choose one of those vectors, which do I prefer?



- > Principal component analysis (PCA)
 - project d -dimensional data into k -dimensional space using linear transformation while preserving as much as possible
 - choose projection with minimum reconstruction error

PCA


> 1) preprocessing data

- data is centered, so that every variable has mean 0
- (optional) we can normalize the data to have variance 1
not usual (variance can be meaningful)
yes, when we have different units of measurement

- $x_i \leftarrow x_i - \mu$
- covariance matrix $\Sigma = \frac{x^T x}{n}$

$$Var(X) = \frac{\sum (X_i - \bar{X})^2}{N} = \frac{\sum x_i^2}{N}$$



$$Cov(X, Y) = \frac{\sum (X_i - \bar{X})(Y_i - \bar{Y})}{N} = \frac{\sum x_i y_i}{N}$$



PCA

> 2) project to 1-dim vector

- project x_i on to a line through the origin with direction w (unit vector)
- distance of the projection from the origin: $(x_i \cdot w)w$



- residual of the projection: $\|x_i - (x_i \cdot w)w\|^2 = x_i \cdot x_i - (x_i \cdot w)^2$

PCA

> 2) project to 1-dim vector

- residual of the projection: $\|x_i - (x_i \cdot w)w\|^2 = x_i \cdot x_i - (x_i \cdot w)^2$
- Add those residuals up across all the vectors


$$MSE(w) = \frac{1}{n} \sum_{i=1}^n [x_i \cdot x_i - (x_i \cdot w)^2]$$

$$= \frac{1}{n} \sum_{i=1}^n \|x_i\|^2 - \left(\frac{1}{n} \sum_{i=1}^n x_i \cdot w \right)^2 - Var[x_i \cdot w]$$

$$= \frac{1}{n} \sum_{i=1}^n \|x_i\|^2 - Var[x_i \cdot w] \quad (\text{the mean of the projection is zero})$$

Find space z_k which maximize variance of projected data

- minimizing the residual sum of squares
= maximizing the variance of the projections
- In general, we project on to multiple principal components those are orthogonal, having the unit vectors w_1, w_2, \dots, w_k



PCA

> 3) maximizing the variance

- let $x = [x_1 \ x_2 \ \dots \ x_n]^T \in \mathbb{R}^{n \times p}$, $w \in \mathbb{R}^{p \times 1}$
- variance $\sigma_w^2 = \frac{1}{n} \sum_{i=1}^n (x_i \cdot w)^2 = \frac{1}{n} w^T x^T x w = w^T v w$, $v = \text{cov}(x)$
- we want to choose w so as to maximize σ_w^2
- Eigenvalue decomposition (diagonalization)
 - since covariance matrix is symmetric and positive semi-definite
 - $v = Q \Lambda Q^T$ where Q is an orthogonal matrix, composed of eigenvectors
 $\Lambda = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_p)$, $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_p$
- Any vector can be represented by eigenvectors (basis)
 - $w = Qa$, $\|w\|^2 = 1 \rightarrow \|a\|^2 = 1$
- $w^T v w = a^T Q^T (Q \Lambda Q^T) Q a = a^T \Lambda a$
- maximizing $a^T \Lambda a$ s.t. $\|a\|^2 = 1 \rightarrow a = [1, 0, 0, \dots, 0]^T$
 - $w = Qa = q_1$: normalized eigenvector corresponding to the maximum eigenvalue






PCA

> Basic PCA algorithm

- 1) re-center: subtract mean from each row of x , $x_i \leftarrow x_i - \mu$
- 2) compute covariance matrix: $\nu = \frac{1}{n} x^T x$
- 3) find eigenvectors q_i and eigenvalues λ_i of ν
- 4) principal components: k eigenvectors with the highest eigenvalues
$$\hat{x}_i = \sum_{j=1}^k (q_j \cdot x_i) q_j = x_i^T Q_k \quad \text{where } Q_k = [q_1 \ q_2 \ \dots \ q_k]^T$$

PCA

> 3D example



PCA

> Application: eigenface



$$\approx w_1 * \text{PC 1} + w_2 * \text{PC 2} + w_3 * \text{PC 3} + \dots$$



PCA


> Application: eigenface



PCA

> Application: eigenface

- reconstruction error occurs when we have minor features



Best reconstructions




Worst reconstructions

PCA


> Working examples

- the first eigenvalue accounts for most variable, so the dimensionality is 1

Example: $\lambda_1 = 0.0938, \lambda_2 = 0.0007$




Example: $\lambda_1 = 0.1260, \lambda_2 = 0.0054$



> Difficult examples

- true non-linear dimension of the data is 1 (using polar coordinates)


Example: $\lambda_1 = 0.0884, \lambda_2 = 0.0725$



Kernel PCA

> Making PCA non-linear

- instead of using x_i , we use some different feature space $\phi(x_i) \in \mathbb{R}^p$
- e.g. using polar coordinates instead of cartesian coordinates
- covariance matrix $\nu = \frac{1}{n} \sum_{i=1}^n \phi(x_i)\phi(x_i)^\top = \frac{1}{n} \boldsymbol{\phi}(x)^\top \boldsymbol{\phi}(x)$
- Kernel: $K(x_i, x_j) = \phi(x_i)^\top \phi(x_j)$
 - we don't have to know the $\phi(x_i)$, we only need $K(x_i, x_j)$
 - given any algorithm that can be expressed solely in terms of dot products, this kernel trick allows us to construct difficult nonlinear versions of it



Kernel PCA

> 3-c) maximizing the variance

- let $\phi(x) = [\phi(x_1) \ \phi(x_2) \ \dots \ \phi(x_n)]^\top \in \mathbb{R}^{n \times p}$, $w \in \mathbb{R}^{p \times 1}$
- $\sigma_w^2 = \frac{1}{n} \sum_{i=1}^n (\phi(x_i) \cdot w)^2 = \frac{1}{n} w^\top \phi(x)^\top \phi(x) w = w^\top v_\phi w$, $v_\phi = \text{cov}(\phi(x))$
- we want to choose w so as to maximize σ_w^2
- Since $\phi(x)$ has (infinite dimensional) basis
 - Any vector can be represented by eigenvectors (basis)
 - $w = \phi(x)^\top a$, $\|w\|^2 = 1 \rightarrow a^\top \phi(x) \phi(x)^\top a = 1 \rightarrow a^\top K a = 1$
- $w^\top v_\phi w = a^\top \phi(x) \phi(x)^\top \phi(x) \phi(x)^\top a = a^\top K^2 a$
- maximizing $a^\top K a$ s.t. $a^\top K a = 1$
 - $K a = n \lambda a$, pick the largest eigenvalue and corresponding eigenvector a and normalize it to satisfy $a^\top K a = 1$

Kernel PCA

> Popular kernels

- Polynomial: $K(x_i, x_j) = (x_i^\top x_j + c)^d$
- RBF (Gaussian): $K(x_i, x_j) = \exp(-\frac{\|x_i - x_j\|^2}{2\sigma^2})$
- Sigmoid: $K(x_i, x_j) = \tanh(\alpha x_i^\top x_j + c)$
- Fisher, Neural tangent, Laplacian, Bessel, ...

> Normalizing the feature space

- $\tilde{\phi}(x_i) = \phi(x_i) - \frac{1}{n} \sum_{k=1}^n \phi(x_k)$
- $\tilde{K}(x_i, x_j) = \tilde{\phi}(x_i)^\top \tilde{\phi}(x_j)$
- $$\begin{aligned}\tilde{K} &= K - \frac{2}{n} \sum_{k=1}^n K(x_i, x_k) + \frac{1}{n^2} \sum_{j,k=1}^n K(x_j, x_k) \\ &= K - 2 \mathbf{1}_{1/n} K + \mathbf{1}_{1/n} K \mathbf{1}_{1/n}\end{aligned}$$

where $\mathbf{1}_{1/n}$ is a matrix with all elements 1/n

Kernel PCA

> Algorithm



- 1) pick a kernel
- 2) construct the normalized kernel matrix \tilde{K} of the data
- 3) find eigenvectors q_i and eigenvalues λ_i of \tilde{K}
- 4) principal components: k eigenvectors q_1, \dots, q_k with the highest eigenvalues $\hat{z}_i(x) = \sum_{j=1}^n q_{ij} K(x_j, x)$

> Comparison to basic PCA

- can capture complex nonlinear structure
- works well for clustering and visualization
- computationally expensive
- requires storing full kernel matrix
- difficult to interpret components
- can't easily reconstruct original data



Kernel PCA

> Comparison to basic PCA / depending on the choice of kernel functions




Nonlinear dimension reduction

- > Assuming data are supported on smooth nonlinear low dimensional geometric objects
- > Manifold: a generalization of curves and surfaces
 - a subset of points in the high-dimensional space that locally looks like a low-dimensional space



(a) Swiss Roll




(b) Torus

- idea is that once the manifold is unfolded, the analysis becomes easy



Nonlinear dimension reduction

- > Key assumption: high dimensional data actually resides in a lower dimensional space that is locally Euclidean
 - find a lower dimensional representation of data that preserves distances between points (MDS)



Geodesic distance


- > Multidimensional scaling (MDS)
 - find a lower dimensional representation of data that preserves distances between points
 - it doesn't have to be Euclidean distance



- we are interested in preserving distances along the manifold (geodesic distances)

Geodesic distance


- > Geodesic distances can be approximated using a graph in which vertices represent data points
 - option1: define a local radius and connect vertices within the radius
 - option2: connect each point to its k nearest neighbors



Isomap

> Algorithm

- construct the similarity graph
- compute geodesic distances by finding shortest paths on the graph
- construct the geodesic distance matrix
- apply MDS



Manifold learning


- > Starting from neighborhood graph and distance matrix
- > Locally-linear embedding (LLE)
 - describe each point as a linear combination of its neighbors
 - finds a low-dimensional embedding that preserves these local relationships via eigenvalue decomposition
- > Local tangent space alignment (LTSA)
 - estimate the tangent space at each point by performing PCA
 - produce an embedding in which the tangent spaces are aligned
- > Laplacian eigenmaps (LE)
 - compute the Laplacian of the graph
 - solve the eigenvalue problem for the graph Laplacian

Manifold learning


> Pros and cons

- Isomap – best for flat manifolds with no holes, small data
- LLE – simple and fast
- LTSA – correct at boundaries
- LE – local smoothness, but only topology


Original data
(Swiss Roll with hole)




Laplacian Eigenmaps (LE)



Isomap




Hessian Eigenmaps (HE)



Local Linear Embedding (LLE)




Local Tangent Space Alignment (LTSA)





Manifold learning

- > Blur the line between dimension reduction and data visualization
 - axes don't correspond to anything in the input space
 - often can't transform new data
 - t-SNE, UMAP, LargeVis
- > How to use t-SNE effectively
 - <https://distill.pub/2016/misread-tsne/>



Manifold learning

- > MNIST dataset visualization



Reference

- > Dimension reduction
 - <https://amueller.github.io/COMS4995-s20/slides/aml-13-dimensionality-reduction>
- > PCA
 - <https://cs229.stanford.edu/notes2022fall/pca.pdf>
 - <https://setosa.io/ev/principal-component-analysis/>
- > Manifold learning
 - <https://arxiv.org/pdf/2311.03757>
 - <https://sites.stat.washington.edu/mmp/Talks/mani-fields22-notes.pdf>
 - https://trevorcohn.github.io/comp90051-2017/slides/16_manifold_learning.pdf