

SME3006 Machine Learning – 2025 Fall


Support Vector Machine



INHA UNIVERSITY


Support vector machines (SVM)

- > Supervised method for binary classification (two classes)
 - maximal margin classifier: only applicable to linearly separable data
 - support vector classifier: can be applied to data that is not linearly separable. Decision boundary still linear
 - SVM: generalization of the above two. Non-linear decision boundary




Hyperplanes

- > A hyperplane is a flat, affine subspace of one dimension less than its ambient space
 - i.e., $(n - 1)$ -dimensional affine subspace in \mathbb{R}^n
 - recall: A subspace of a vector space is a subset that is closed under addition and scalar multiplication.
 - Affine subspace: not required to include zero vector, not closed





Hyperplanes

- > Defined as set of points $x \in \mathbb{R}^n$ satisfying $w^T x + b = 0$
 - In 2D, $w_x x + w_y y + b = 0$ (line)
 - In 3D, $w_x x + w_y y + w_z z + b = 0$ (plane)
 - Any point satisfying the equation lies on the hyperplane
- > Separating with a hyperplane
 - $w^T x + b > 0$ or $w^T x + b < 0$
 - point lies on either side of the hyperplane
 - it divides the n -dimensional space into two halves




Hyperplanes

- > Hyperplane classifier
 - use a separating hyperplane for binary classification
 - key assumption: classes can be separated by a linear decision boundary
- > recall: logistic regression
 - it effectively finds a separating hyperplane



Hyperplanes

- > Note that for a linearly separable dataset, there are many possible separating hyperplanes (in fact, an infinite number)
 - they perform equally well on the training set
 - but show different results on the test set (generalization)




Maximal margin hyperplane

> Which hyperplanes should we choose?


- maximal margin hyperplane: separating hyperplane that is the farthest from the training samples
- margin: smallest distance between point x_i and the hyperplane

- For hyperplane, $w^T x = 0$ (unbiased)
 - w is orthogonal to the hyperplane
 - the unit direction is $w/\|w\|$
 - compute the projection of x_i
 - x_i has distance $|w^T x_i|/\|w\|$



Maximal margin hyperplane


- > Which hyperplanes should we choose?
 - maximal margin hyperplane: separating hyperplane that is the farthest from the training samples
 - margin: smallest distance between point x_i and the hyperplane
- For hyperplane, $w^T x + b = 0$ (biased)
 - assume unbiased hyperplane HP_o (translated)
 - vectors $d(HP, x_i) + d(HP, HP_o) = d(x_i, HP_o)$
 - w is orthogonal to the hyperplane
 - $d(x_i, HP_o) = w^T x_i / \|w\|$
 - $d(HP, HP_o) = b / \|w\|$
 - distance $|d(x_i, HP_o)| = |w^T x_i + b| / \|w\|$



Maximal margin hyperplane

> Which hyperplanes should we choose?

- maximal margin hyperplane: separating hyperplane that is the farthest from the training samples
 - margin: smallest distance between any point x_i and the hyperplane
 - support vectors: data points that have the distance equal to the margin
-
- we want to maximize the margin
 - we have data points (x_i, y_i) , $y_i \in \{-1, 1\}$
 - $\max M$
 - $s.t.$ $\sum_{j=0}^n w_j^2 = 1$
 - $y_i(w^\top x_i + w_0) \geq M, \forall i$
 - constraint necessary for well-defined optimization problem




Maximal margin hyperplane

> Optimization problem

$$\begin{aligned} & \max_w |w^\top x + w_0| / \|w\| \\ & \text{s.t. } \sum_{j=0}^n w_j^2 = 1 \end{aligned}$$

$$y_i(w^\top x_i + w_0) \geq M, \forall i$$




$$\begin{aligned} & \min_w \|w\|^2 \\ & \text{s.t. } y_i(w^\top x_i + w_0) \geq 1 \end{aligned}$$

- quadratic cost function and linear constraints: quadratic program (convex)
- Quadratic program (QP)

$$\begin{aligned} & \underset{x}{\text{minimize}} \frac{1}{2} x^\top P_0 x + q_0^\top x + c_0 \\ & \text{s.t. } Ax = b \\ & \quad x \geq 0 \end{aligned}$$


Convex Optimization

> Unique minimum: global/local




Maximal margin hyperplane

- > What if there is no separating hyperplane?
- > In addition,
 - it can be sensitive to individual data points
 - it may overfit training data
 - one outlier could ruin the algorithm



Support vector classifier

- > Like a maximal classifier, it looks for a hyperplane to perform classification
- > Training samples are allowed to be on the ‘wrong side’
- > It separates the classes using a soft margin
 - maximal margin hyperplane is a hard-margin SVM
 - now, margin is not the closest distance, it is determined by the parameter
 - support vectors are points within the margin or on the wrong side



Support vector classifier

> We want to maximize the margin

- we have m data points (x_i, y_i) , $y_i \in \{-1, 1\}$

- $\max_w M$

$$s.t. \quad \sum_{j=0}^n w_j^2 = 1 \quad \longrightarrow$$

constraint necessary for well-defined optimization problem

$$\begin{aligned} y_i(w^\top x_i + b) &\geq M(1 - \epsilon_i), \quad \forall i \\ \sum_{i=1}^m \epsilon_i &\leq C, \quad \epsilon_i \geq 0, \quad \forall i \end{aligned}$$


- points must be at least distance M from hyperplane, or pay a penalty ϵ_i
- there is a limit on the total penalties C
- i.e., you can violate the margin, but only by a total amount C

Support vector classifier

- > We want to maximize the margin
 - we have m data points (x_i, y_i) , $y_i \in \{-1, 1\}$
 - $\max_w M$
 - s.t. $\sum_{j=0}^n w_j^2 = 1$

$$y_i(w^\top x_i + b) \geq M(1 - \epsilon_i), \forall i$$
$$\sum_{i=1}^m \epsilon_i \leq C, \epsilon_i \geq 0, \forall i$$


- Penalties
 - $\epsilon_i = 0$: training point is on correct side of margin
 - $\epsilon_i > 0$: training point violates the margin
 - $\epsilon_i > 1$: training point is misclassified (wrong side of hyperplane)






Support vector classifier

- > C determines the total budget for violations
 - it is a hyperparameter that we tune using cross-validation
 - if $C = 0$, we recover the maximal margin classifier
 - As C goes from small to large, there is bias-variance tradeoff

- Large C
high bias, low variance
large margin
many support vectors




- Small C
low bias, high variance
small margin
few support vectors



Support vector classifier


- > We are still using a linear decision boundary
 - Some datasets are not linearly separable, but linearly separable when transformed into a higher dimensional space

Original feature space



variables x_1, x_2

New feature space



variables x_1, x_2, x_1x_2

Expanding feature space

- > We are lifting the data into a new feature space

- original data

$$x = (x_1, x_2, \dots, x_n) \in \mathbb{R}^n$$



- lifted data

$$\text{e.g., } \tilde{x} = (x_1, x_2, \dots, x_n, x_1^2, x_2^2, \dots, x_n^2) \in \mathbb{R}^{2n}$$

- support vector classifier will find a hyperplane in $2n$ dimensions


- hyperplane will be nonlinear in the original space

- in this case, it is an ellipse




Expanding feature space

- > Can imagine adding higher order terms and to expand feature set
- > Large number of features becomes computationally challenging
- > We need an efficient way to work with large number of features




SVM

- > Extends the support vector classifier by using kernel functions to achieve non-linear decision boundaries



- recall: kernel function $K(x_i, x_j) = \phi(x_i)^\top \phi(x_j)$
- they implicitly map data into higher-dimensional space



SVM

- > Support vector classifier (linear SVM)
 - hyperplane: $w^T x + b$,
 - If we solve the optimization problem,
 $w = \sum_i \alpha_i y_i x_i$ (from the optimality condition of Lagrangian)
 - classifier: $f(x) = w^T x + b - M$ $\rightarrow f(x) > 0 \text{ or } f(x) < 0$
 $= (\sum_i \alpha_i y_i x_i^T x) + b - M$
inner product
 - > General SVM
 - we replace the inner product with some kernel function
 - classifier: $f(x) = (\sum_i \alpha_i y_i K(x_i, x)) + b - M$


> Properties of kernels

- generalization of inner product without explicitly computing the mapping
 - $\phi: X \rightarrow X^\phi$
 - $x \mapsto \phi(x)$
 - $K(x, x') = \phi(x)^\top \phi(x')$
- Symmetric: $K(x, x') = K(x', x)$
- Kernel matrix is positive semidefinite: $K \geq 0$

SVM

> Popular kernels

- Polynomial: $K(x_i, x_j) = (x_i^\top x_j + c)^d$
- RBF (Gaussian): $K(x_i, x_j) = \exp(-\frac{\|x_i - x_j\|^2}{2\sigma^2})$
- Sigmoid: $K(x_i, x_j) = \tanh(\alpha x_i^\top x_j + c)$
- Fisher, Neural tangent, Laplacian, Bessel, ...



> Creating more complicated kernels

- $K(x_i, x_j) = K_1(x_i, x_j) + K_2(x_i, x_j)$
- $K(x_i, x_j) = \alpha K_1(x_i, x_j)$
- $K(x_i, x_j) = K_1(x_i, x_j)K_2(x_i, x_j)$





More about RBF kernel

- > RBF (Gaussian): $K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$
 - for $d = 1$, $\phi(x) = \exp\left(-\frac{x^2}{2\sigma^2}\right) \left[1, \frac{x}{\sigma\sqrt{1!}}, \frac{x^2}{\sigma\sqrt{2!}}, \frac{x^3}{\sigma\sqrt{3!}}, \dots\right]^\top$
 - this is an infinite vector, nobody actually uses this value
- Very popular in practice
 - gives very smooth hypothesis
 - behaves somewhat like k-nearest neighbors, but smoother
 - oscillates less than polynomials
- choose σ by validation
 - σ trades off bias vs. variance
 - larger $\sigma \rightarrow$ wider Gaussians & smoother $h \rightarrow$ more bias & less variance

SVM

> Decision boundaries in Kernel SVMs

- linear: best for linearly separable data; simplest and most interpretable
- polynomial: captures curved decision boundary, flexibility increases with deg.
- RBF: highly flexible, local decision boundary
- sigmoid: inspired by NN, but rarely used due to instability



SVM with 3+ classes



> Adapt SVMs to perform classification for more than 2 classes

- One vs. one

- construct an SVM for each pair of classes
- for k classes, this requires training and testing $k(k - 1)/2$ SVMs
- computationally expensive for large k

- One vs. all

- construct an SVM for each class against the $k - 1$ other classes pooled together
- for k classes, this requires training and testing k SVMs
- may exacerbate class imbalances, distance to hyperplane may not correspond well to confidence



SVM

> Summary


- regularization parameter C helps avoid overfitting
- use of kernel gives flexibility in shape of decision boundary
- optimization problem is convex – unique solution

- must tune hyperparameters (e.g., C , kernel function)
- must formulate as binary classification
- difficult to interpret

RKHS (deeper look into kernel methods)

> Introduction to Reproducing kernel Hilbert space (RKHS)

- In kernel trick, we don't know the feature map explicitly
- We're still doing optimization, regularization, and learning in this unknown space
 - Is it okay? Is this optimization mathematically sound?
- A RKHS is a special Hilbert space of functions satisfying:
 - the kernel function acts as a feature
 - every function in this space satisfies the reproducing property



Hilbert space

- > Inner product space containing Cauchy sequence limits (complete)
 - it is like a generalization of Euclidean space to functions
- > Inner product: A function $\langle \cdot, \cdot \rangle_{\mathcal{H}} : \mathcal{H} \times \mathcal{H} \rightarrow \mathbb{R}$
 - linear: $\langle \alpha_1 f_1 + \alpha_2 f_2, g \rangle_{\mathcal{H}} = \alpha_1 \langle f_1, g \rangle_{\mathcal{H}} + \alpha_2 \langle f_2, g \rangle_{\mathcal{H}}$
 - symmetric: $\langle f, g \rangle_{\mathcal{H}} = \langle g, f \rangle_{\mathcal{H}}$
 - $\langle f, f \rangle_{\mathcal{H}} \geq 0$ and $\langle f, f \rangle_{\mathcal{H}} = 0$ if and only if $f = 0$
- > Cauchy sequence
 - a sequence where the elements become arbitrarily close to each other as the sequence progresses
 - ex) 1, 1.4, 1.41, 1.414, 1.4142, ... (approaching $\sqrt{2}$) is Cauchy
In \mathbb{Q} (rational number), its limit $\sqrt{2} \notin \mathbb{Q}$, so not complete
In \mathbb{R} , $\sqrt{2} \in \mathbb{R}$, so \mathbb{R} is a complete space

Reproducing property

- > Obtain the value of a function $f(x)$ by taking an inner product between the function and the kernel
 - $f(x) = \langle f(\cdot), k(\cdot, x) \rangle_{\mathcal{H}}$ (reproducing property)
 - Imagine $f(\cdot)$ is hidden, but you want to know $f(x)$
 - RKHS is a Hilbert space full of functions
 - A positive semidefinite kernel function leads to the reproducing kernels (Moore-Aronszajn Theorem)
- > Then we estimate $\hat{f}(\cdot) = \sum_{i=1}^n \alpha_i k(x_i, \cdot)$ (Representer Theorem)
 - In regularized empirical risk minimization problems over an RKHS, the optimal function \hat{f} can always be expressed as a linear combination of kernel functions centered at the training data points
 - $\hat{f}(x) = \langle \sum_{i=1}^n \alpha_i k(x_i, \cdot), k(\cdot, x) \rangle_{\mathcal{H}} = \sum_{i=1}^n \alpha_i k(x_i, x)$


RKHS

> To sum up,

- Any symmetric, PSD kernel defines an RKHS (Moore-Aronszajn theorem)
(We don't have to define a function space – the kernel defines it)
- In learning problems, we don't need to search over infinite-dimensional functions – we need a weighted sum of kernels at the data points
(Representer theorem)

> RKHS norm $\|f\|_{\mathcal{H}}^2 = \sqrt{\langle f(\cdot), f(\cdot) \rangle_{\mathcal{H}}}$

- the norm acts as a built-in measure of how complex a function is
- small norm results in smooth functions -> better generalization



Reference

- > SVM
 - https://web.stanford.edu/class/cme250/files/cme250_lecture5.pdf
 - https://www.cs.princeton.edu/courses/archive/spring16/cos495/slides/ML_basics_lecture4_SVM_I.pdf
 - https://www.cs.princeton.edu/courses/archive/spring16/cos495/slides/ML_basics_lecture5_SVM_II.pdf
- > RKHS
 - <https://www.gatsby.ucl.ac.uk/~gretton/coursefiles/rkhscourse.html>