

SME3006 Machine Learning – 2025 Fall

Clustering: K-means and more



INHA UNIVERSITY

Overview

- > Introduction to clustering problem
- > How close is the data? distance/similarity
- > K-means and extensions
- > Other clustering methods
 - DBSCAN
 - Hierarchical Clustering
 - (Spectral Clustering)

Types of machine learning

> Supervised learning

- training a model from input data and its corresponding targets

Training data

X	y
😺	CAT
😺	CAT
...	...
🐶	DOG
🐺	DOG

Learning algorithm



ML model

Learned
function f

Unseen test data

X	y
😺	?
🐶	?

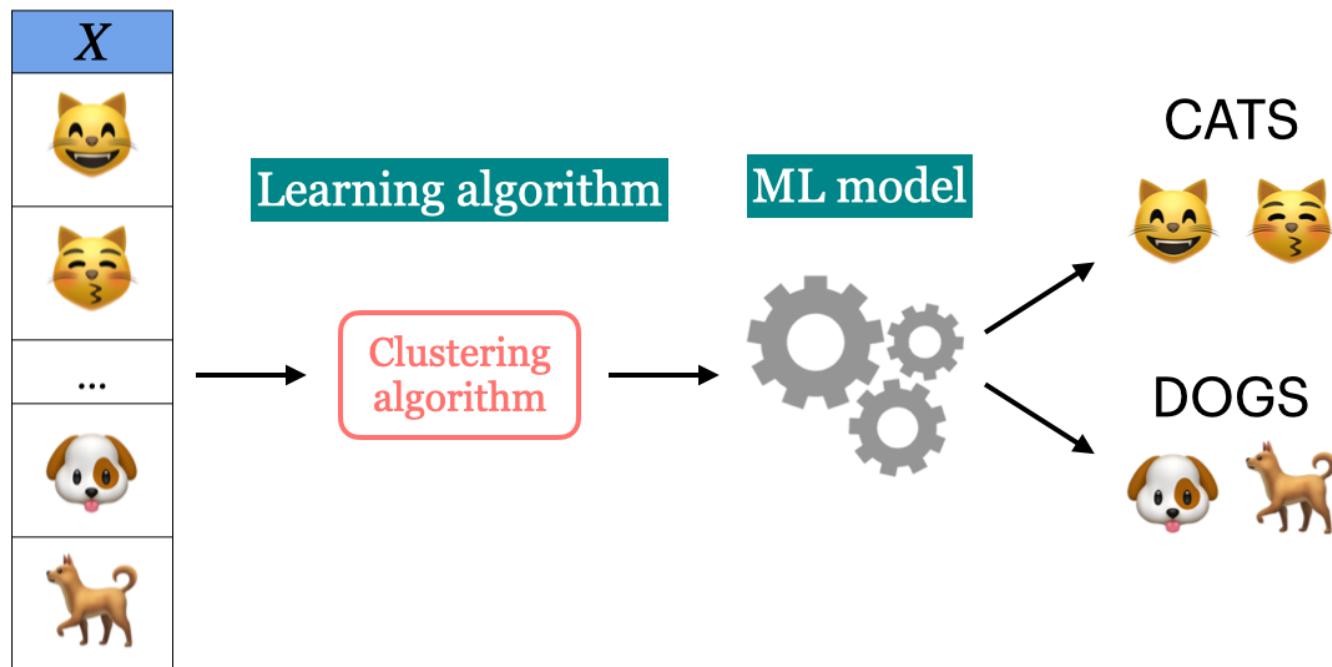
Predictions

\hat{y}
CAT
DOG

Types of machine learning

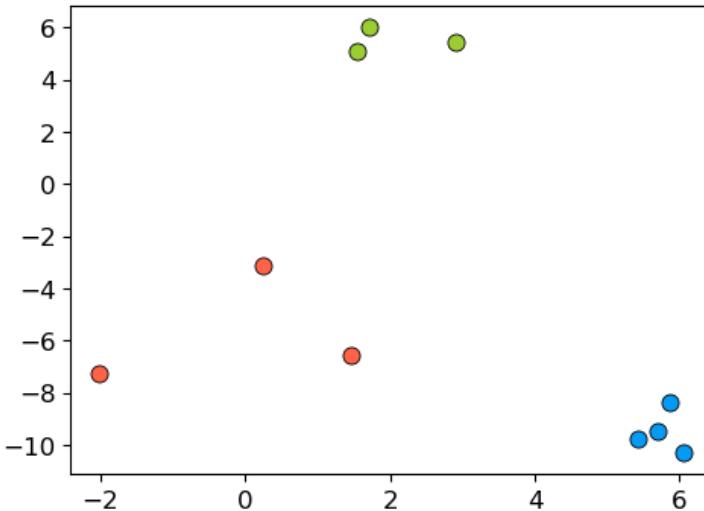
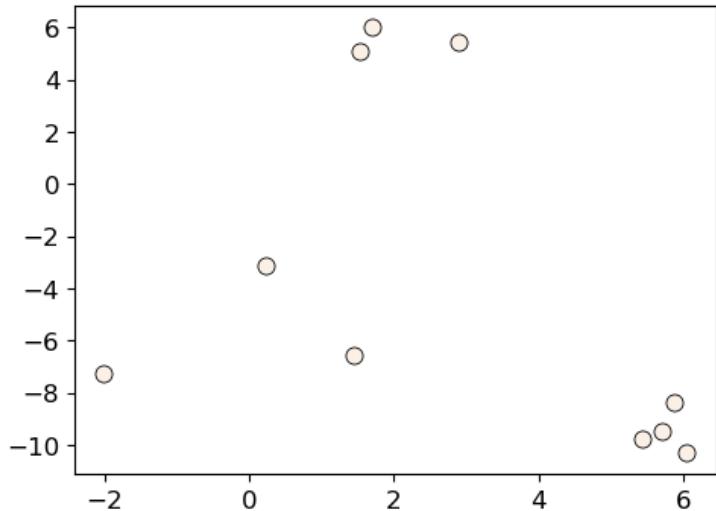
- > Unsupervised learning
 - only training input data (since labeling is expensive)
 - learning will be focused on finding the underlying structures of the inputs

Training data



Clustering

- > Group similar examples together to get some insight into the data



- clusters are identified by a cluster label
- label is only for separating the clusters (knowing the different group)
- In real-world data, we often do not know how many clusters are there

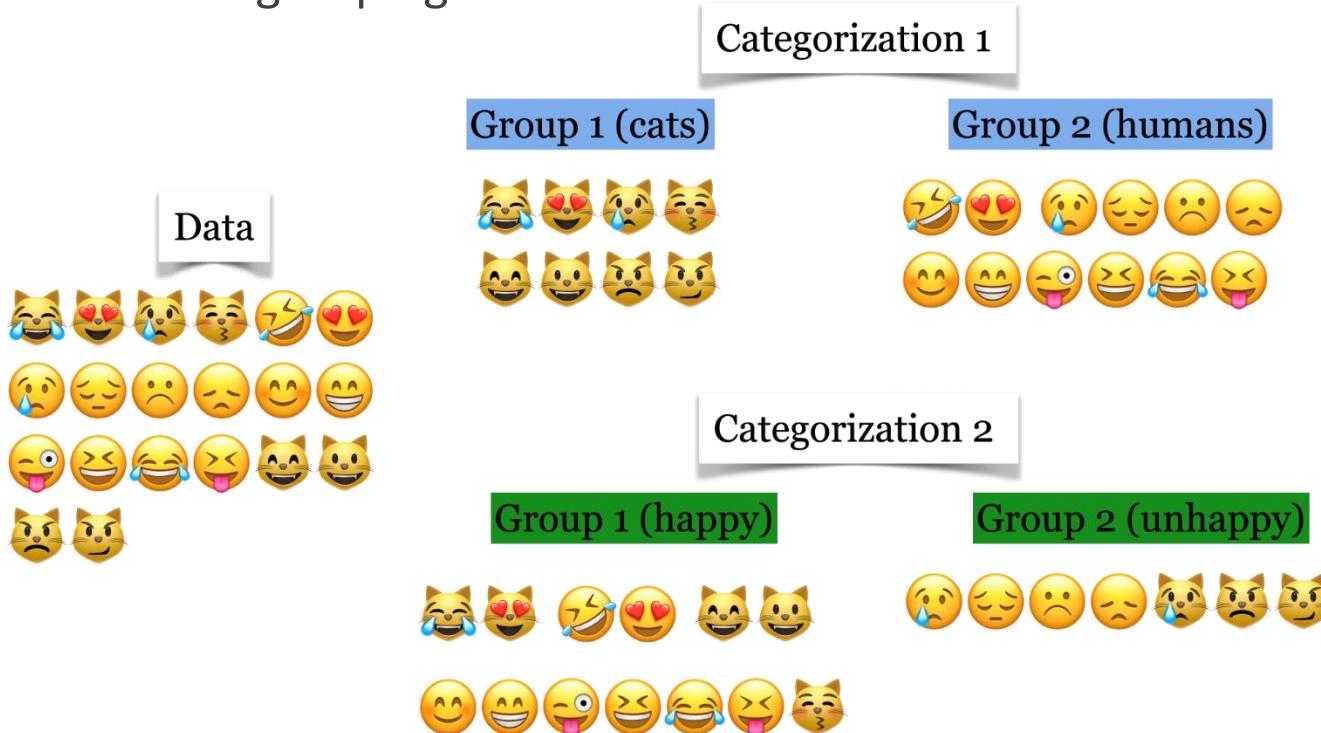
Clustering

- > What is correct grouping?



Clustering

- > What is correct grouping?



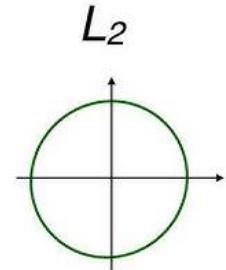
- Both seem reasonable
- This makes it hard to for us to measure the quality of a clustering algorithm

Clustering

- > How can we define the similarity (or closeness) between data points?
 - Euclidean distance (L2 norm)

$$d(x, y) = \sqrt{\sum_i (x_i - y_i)^2}$$

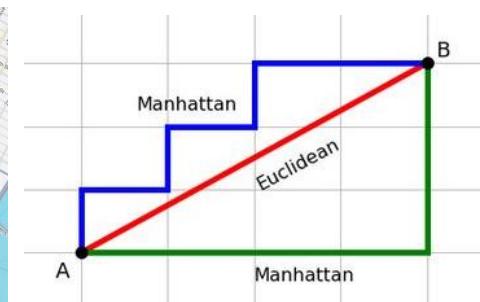
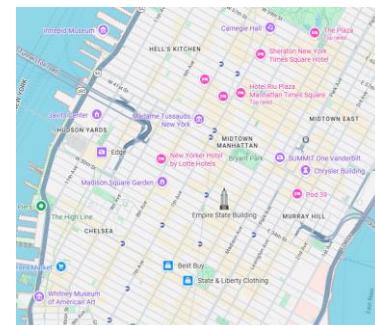
- measure straight-line distance in continuous feature space
- assumes all features are equally scaled and uncorrelated



- Manhattan distance (L1 norm)

$$d(x, y) = |\sum_i (x_i - y_i)|$$

- less sensitive to outliers
- intuitively, we can't go diagonally



Clustering

- > How can we define the similarity (or closeness) between data points?
 - Cosine similarity

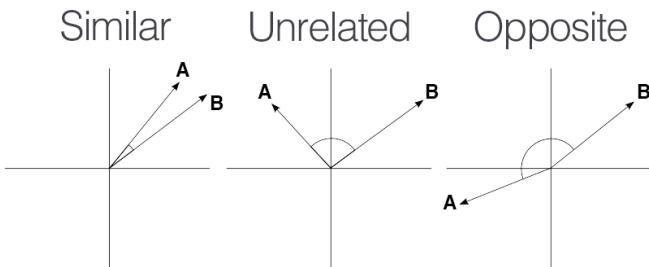
$$sim(x, y) = \frac{x \cdot y}{\|x\| \|y\|}$$

- measures only the angle between two vectors (ignoring the magnitude)
- effective in high-dimensional, sparse data

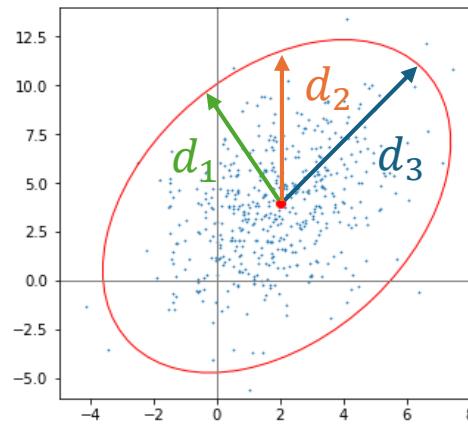
- Mahalanobis distance

$$d(x, y) = \sqrt{(x - y)^\top \Sigma (x - y)}$$

- accounts for correlations between features using the covariance matrix Σ
- useful when variables have different scales



$$d_1 = d_2 = d_3$$



Clustering

- > How can we define the similarity (or closeness) between data points?

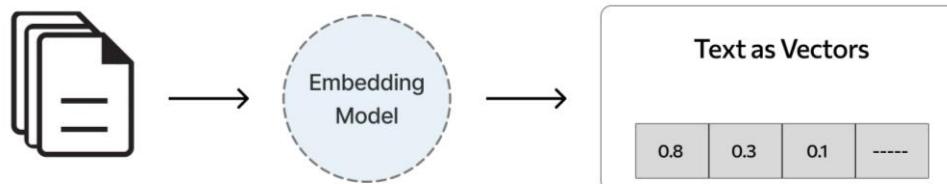
- Categorical data

- Hamming distance:
count the number of mismatch
- Jaccard similarity:
measure the size of intersection

$$sim(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

	Blood Type	Hair Color
student 1	A	Black
student 2	B	Brown
student 3	B	Black
student 4	AB	Red

- In practice, text and image data
 - embed to high-dimensional vector (dimension reduction)
 - compute the geometric distance in the feature space



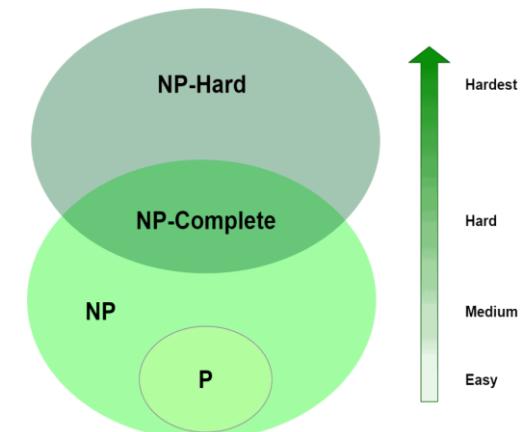
K-means

> Assumptions

- the data lives in a Euclidean space
- the data belongs to K classes
- the data points from same class are similar (close in Euclidean distance)
- K-means assumes there are k clusters and each point is close to cluster center

> Chicken and egg problem

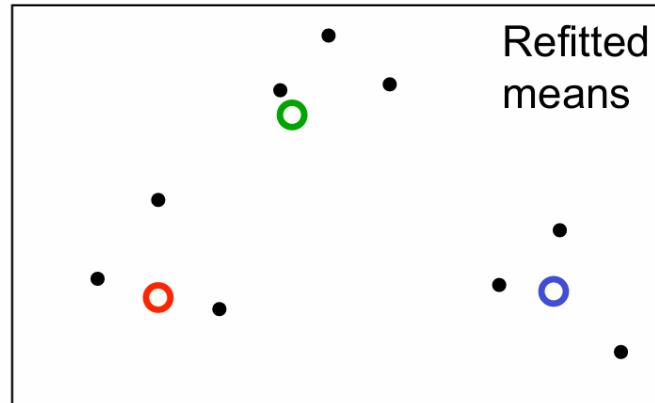
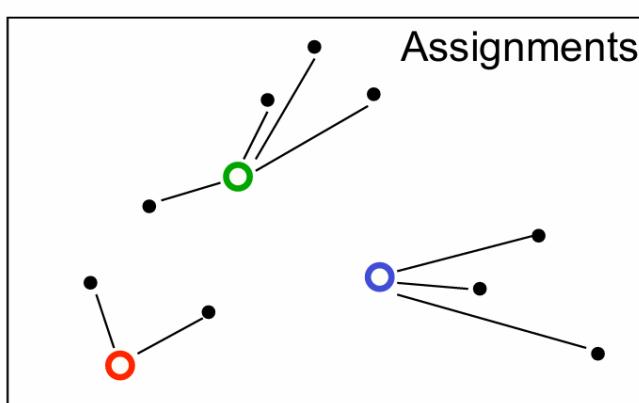
- if we knew the cluster assignment, we could easily compute means
- if we know the means, we could easily compute cluster assignment
- it is a NP hard



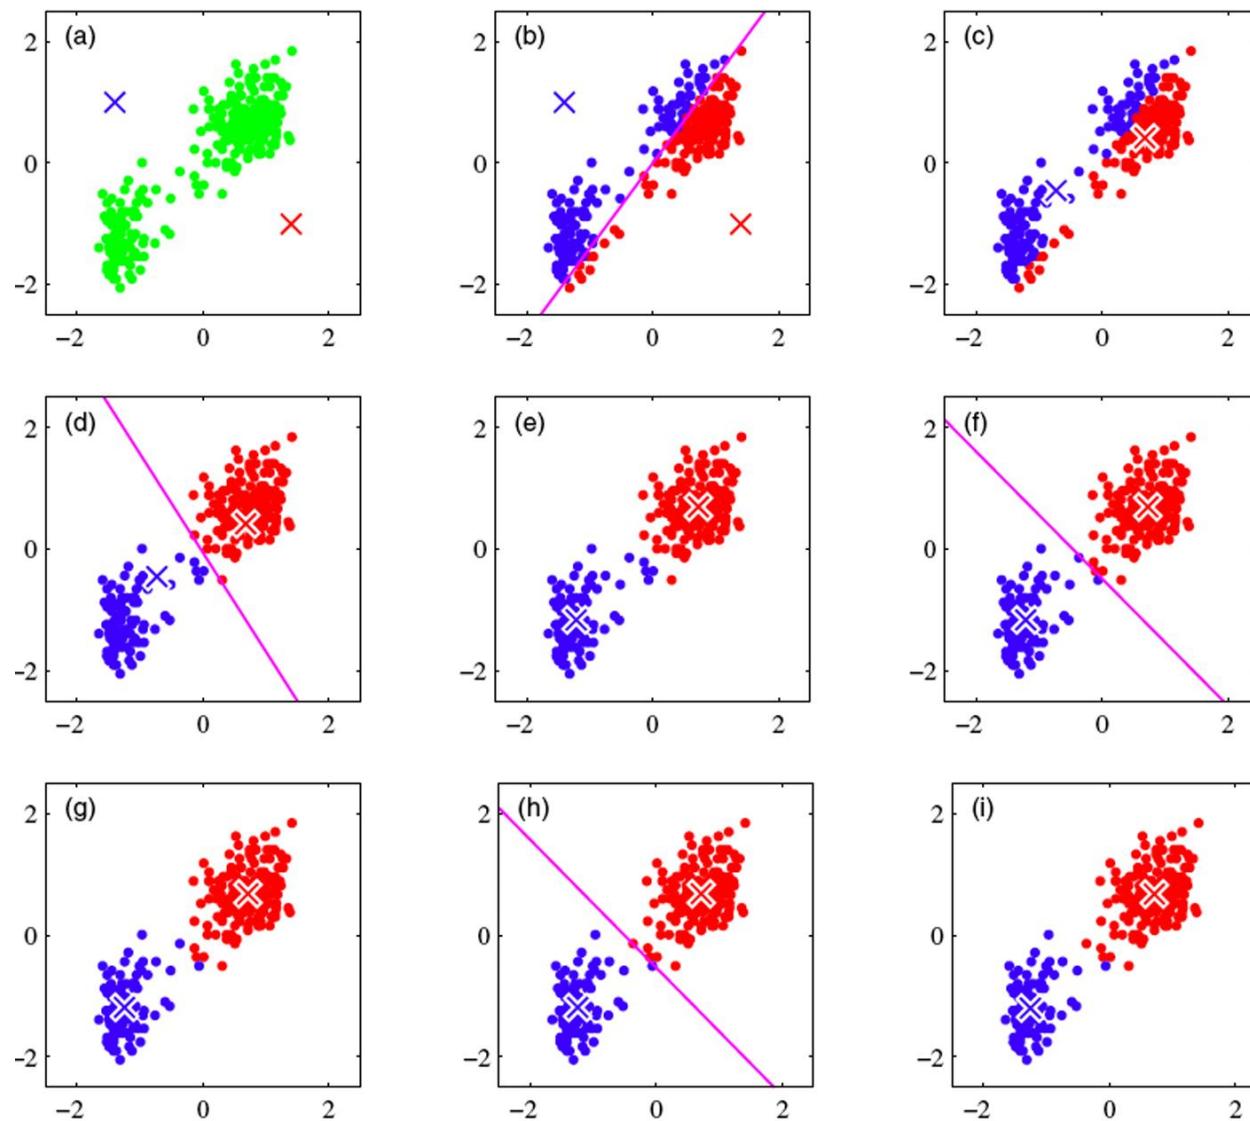
> Simple heuristics: start randomly and alternate between the two

K-means

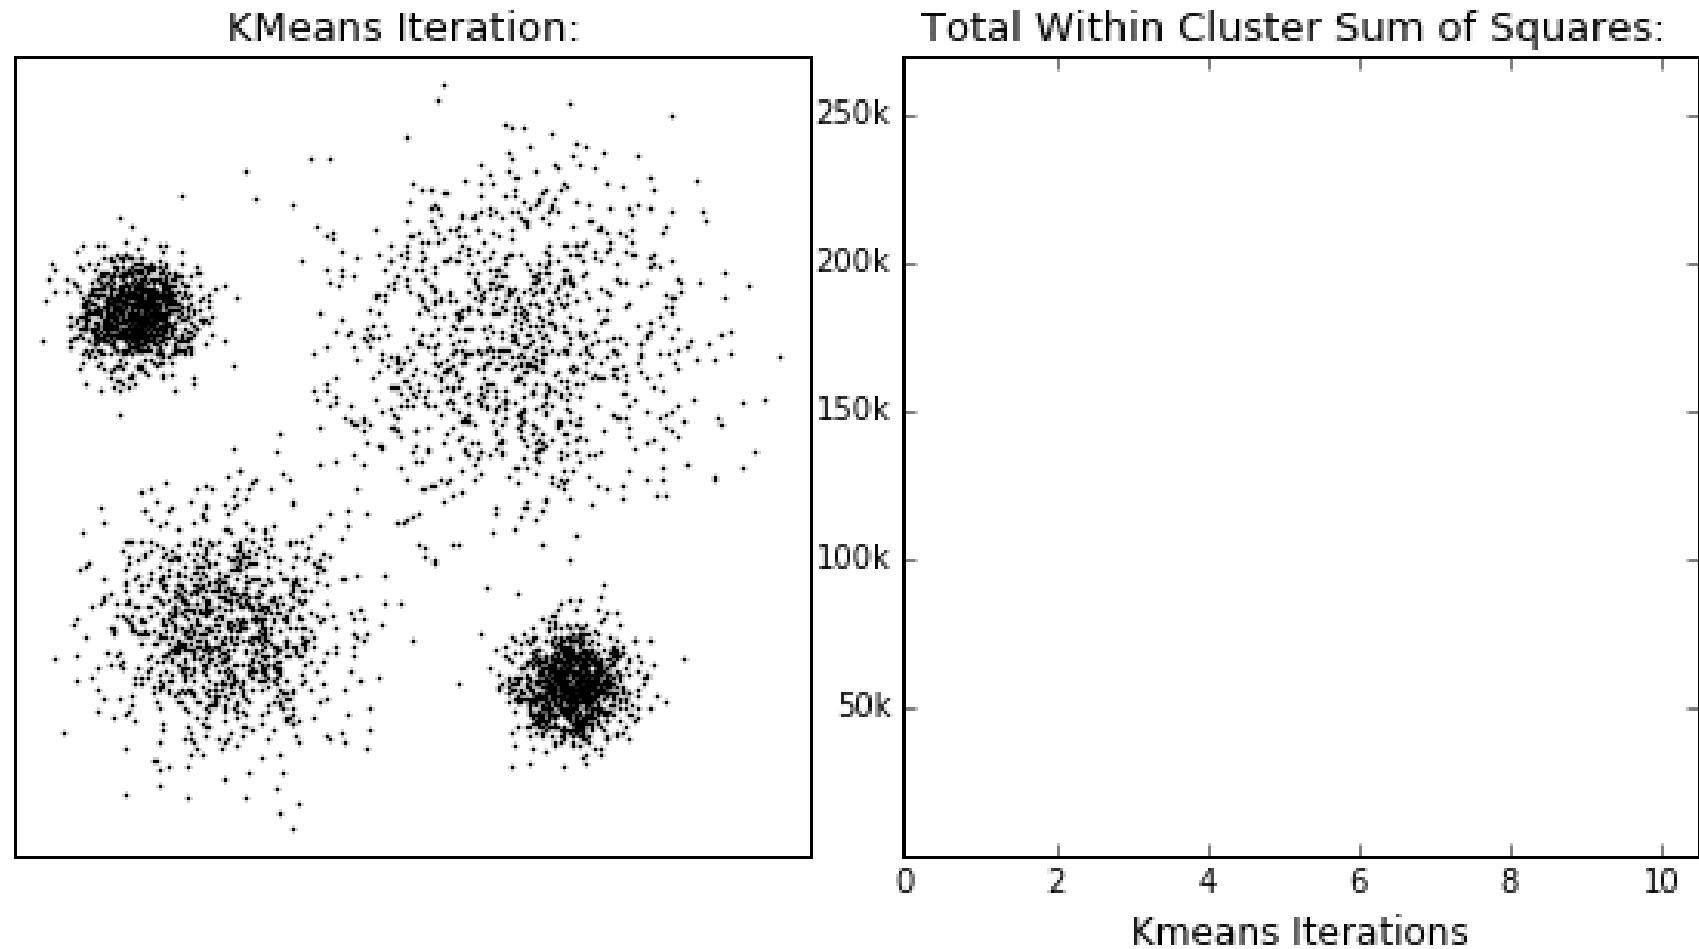
- > Initialization: randomly initialize cluster centers
- > Iteratively alternates between two steps:
 - assignment step: assign each data point to the closest cluster
 - refitting step: move each cluster center to the center of gravity of the data assigned to it



K-means



K-means

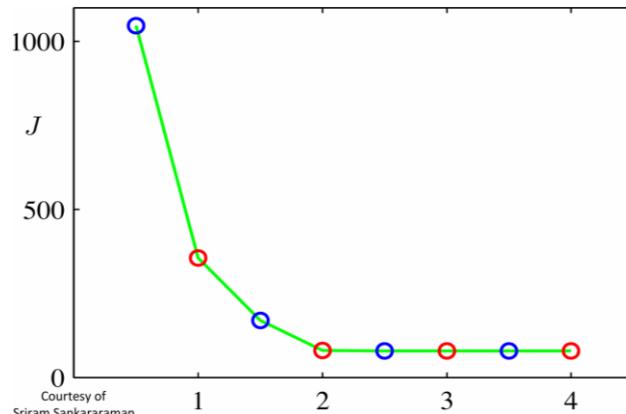


K-means

> Objective: find cluster center μ_i that minimizes the following cost

$$J = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2 \text{ where } C_i \text{ is the set of points assigned to cluster } i$$

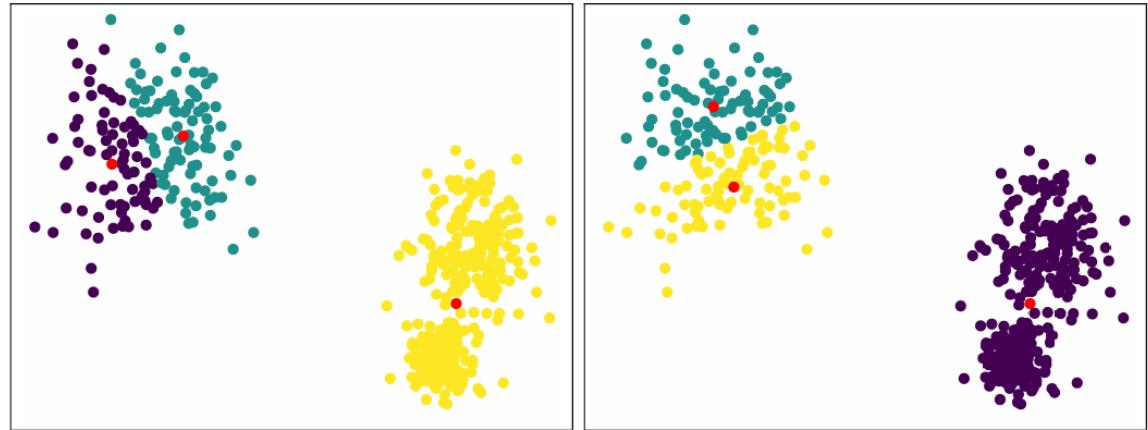
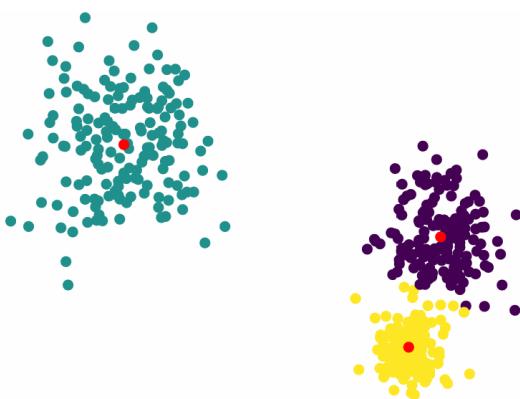
- assignment step: assign each data point to the closest cluster
- refitting step : move each cluster center to the center of gravity of the data assigned to it
- cost J cannot increase
 - assignment step: for each point, distance decreases
 - refitting step: for fixed cluster, the best centroid is the mean of the points



K-means

> Poor clustering results

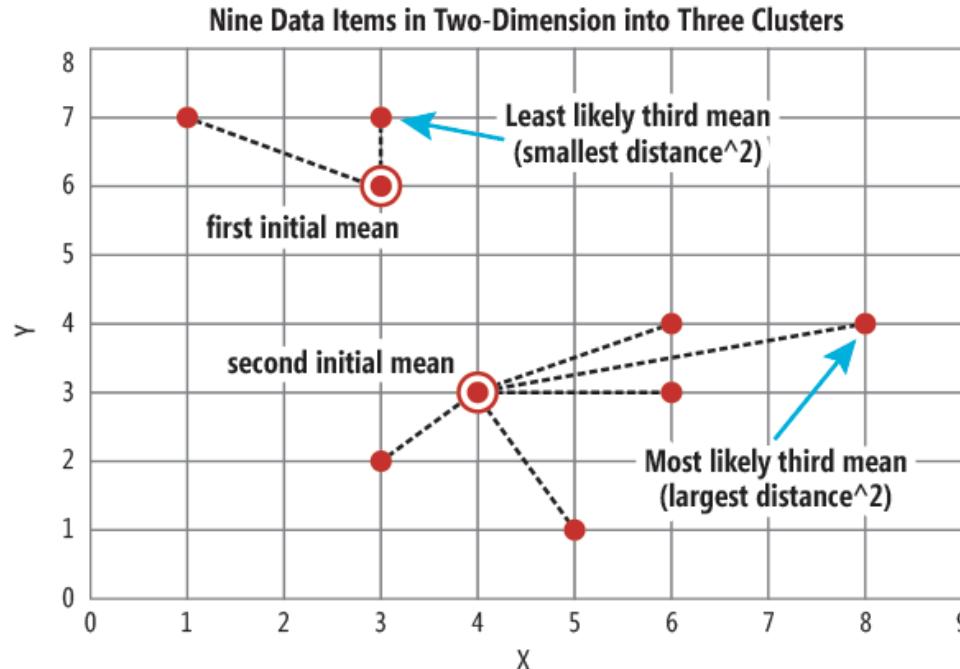
- it depends on the random initialization (since it a non-convex problem)
- there is nothing to prevent k-means getting stuck at local minima



- try many random starting points
- try non-local split-and-merge moves
 - simultaneously merge two nearby clusters
 - and split a big cluster into two

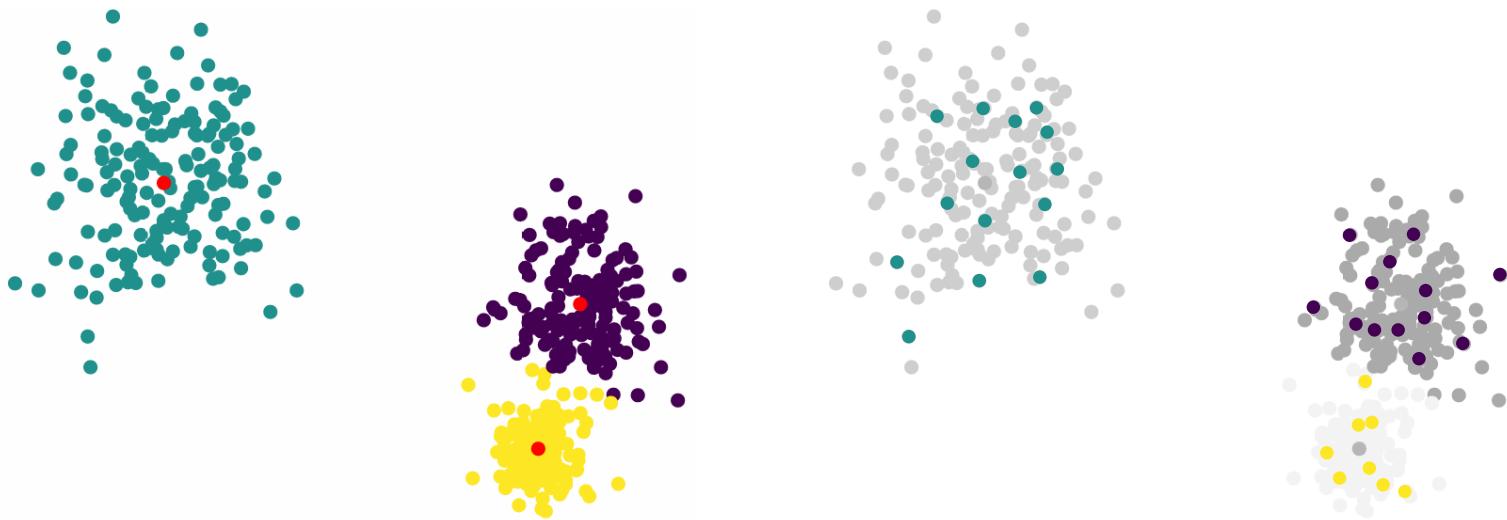
K-means extensions

- > K-means++: improves initialization
 - choose the first center uniformly at random from data points
 - For each new center, pick a point with probability proportional to the square of its distance from the nearest chosen center
 - once all centers are selected, run K-means
- > In practice: more accurate and faster than k-means



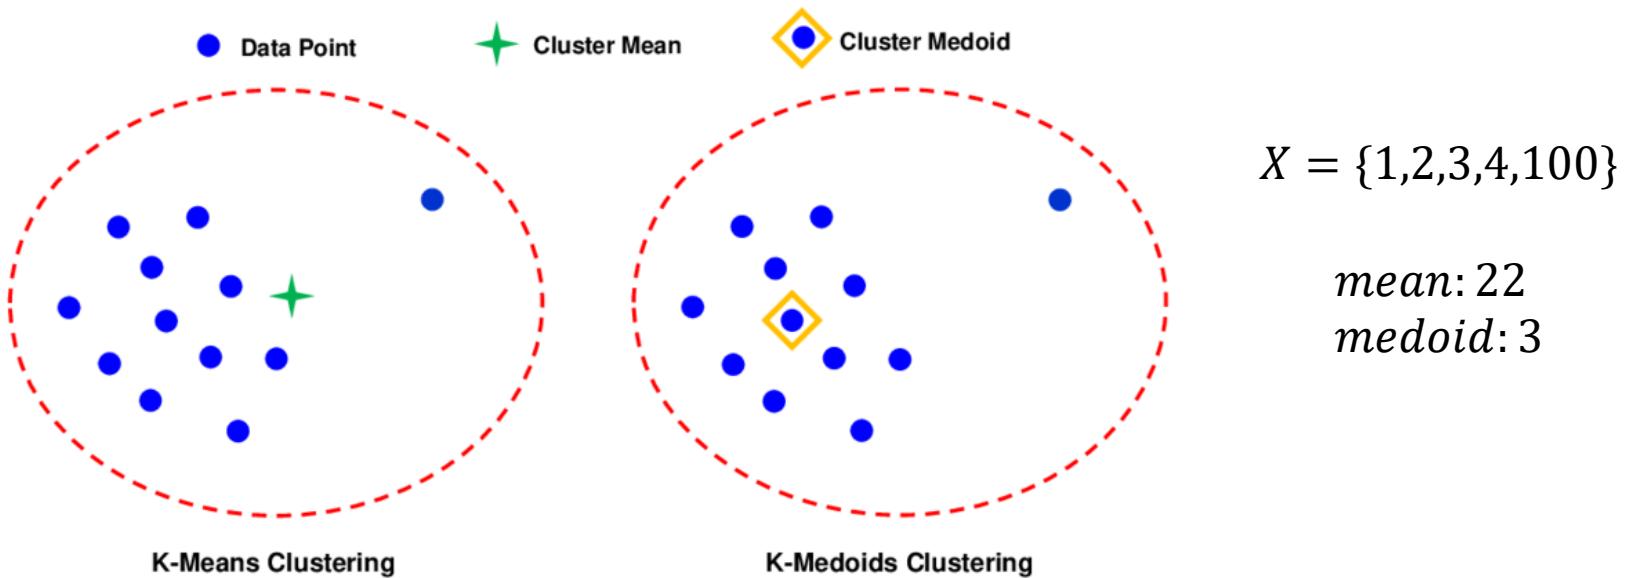
K-means extensions

- > mini-batch K-means: improves the memory and computation
 - randomly sample a mini-batch of data points
 - update cluster centroids using only the sampled data, then resample (like gradient descent → stochastic gradient)



K-means extensions

- > K-medoids: improves robustness
 - instead of using means, we chose the central data point (medoid)
 - we can use other distance (dissimilarity) measures (e.g., categorical)
 - it is more robust to outliers than K-means



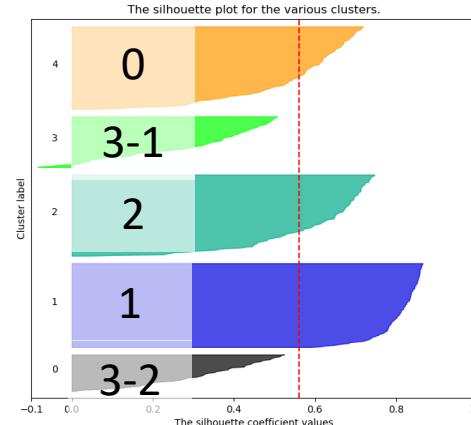
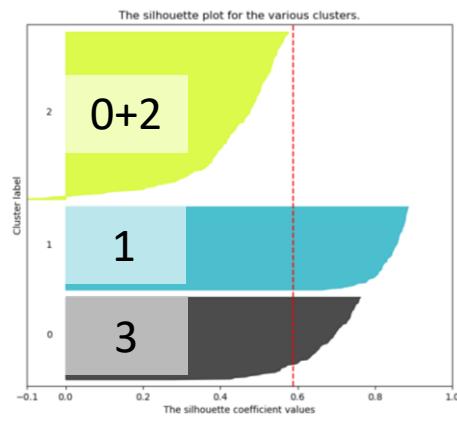
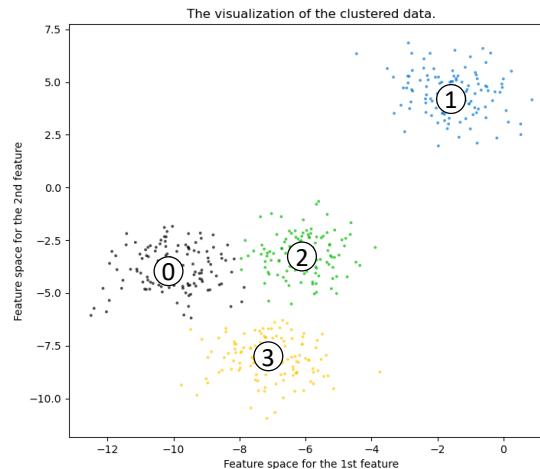
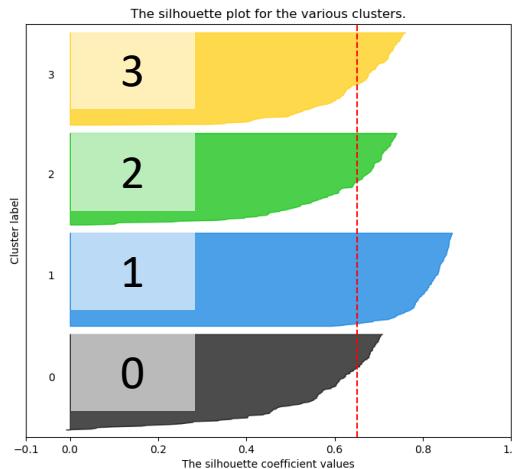
K-means extensions

- > how to choose the number of clusters k
 - elbow method
 - for each k , compute the cost
 - plot the cost vs. k and look for the elbow point
 - cost is monotonically decreasing (why?)



K-means extensions

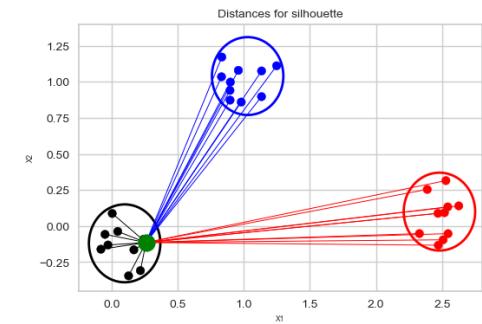
- > how to choose the number of clusters k
 - Silhouette analysis



$$S_i = \frac{b_i - a_i}{\max(a_i, b_i)}$$

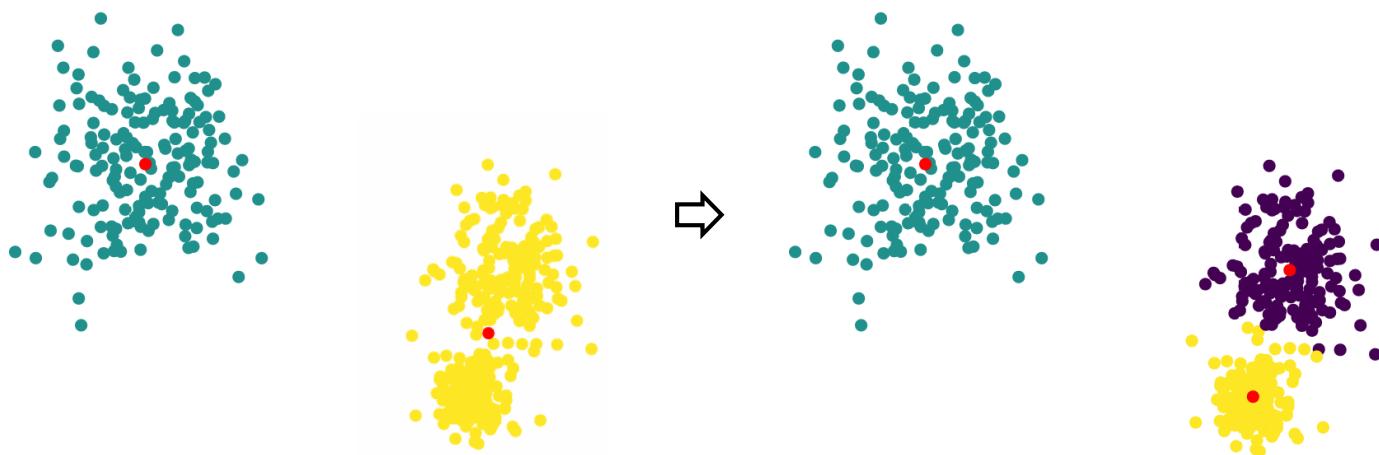
a_i : avg. distance to other points in the same cluster

b_i : avg. distance to points in the nearest different cluster



K-means extensions

- > how to choose the number of clusters k
 - X-means: if splitting a cluster into two improves the model, then split it
 - based on Bayesian information criterion
 - G-means: if the data within a cluster doesn't look Gaussian, then split it
 - assumes each true cluster should follow a Gaussian distribution
 - Both start with small number of clusters (e.g., 1 or 2)



K-means applications

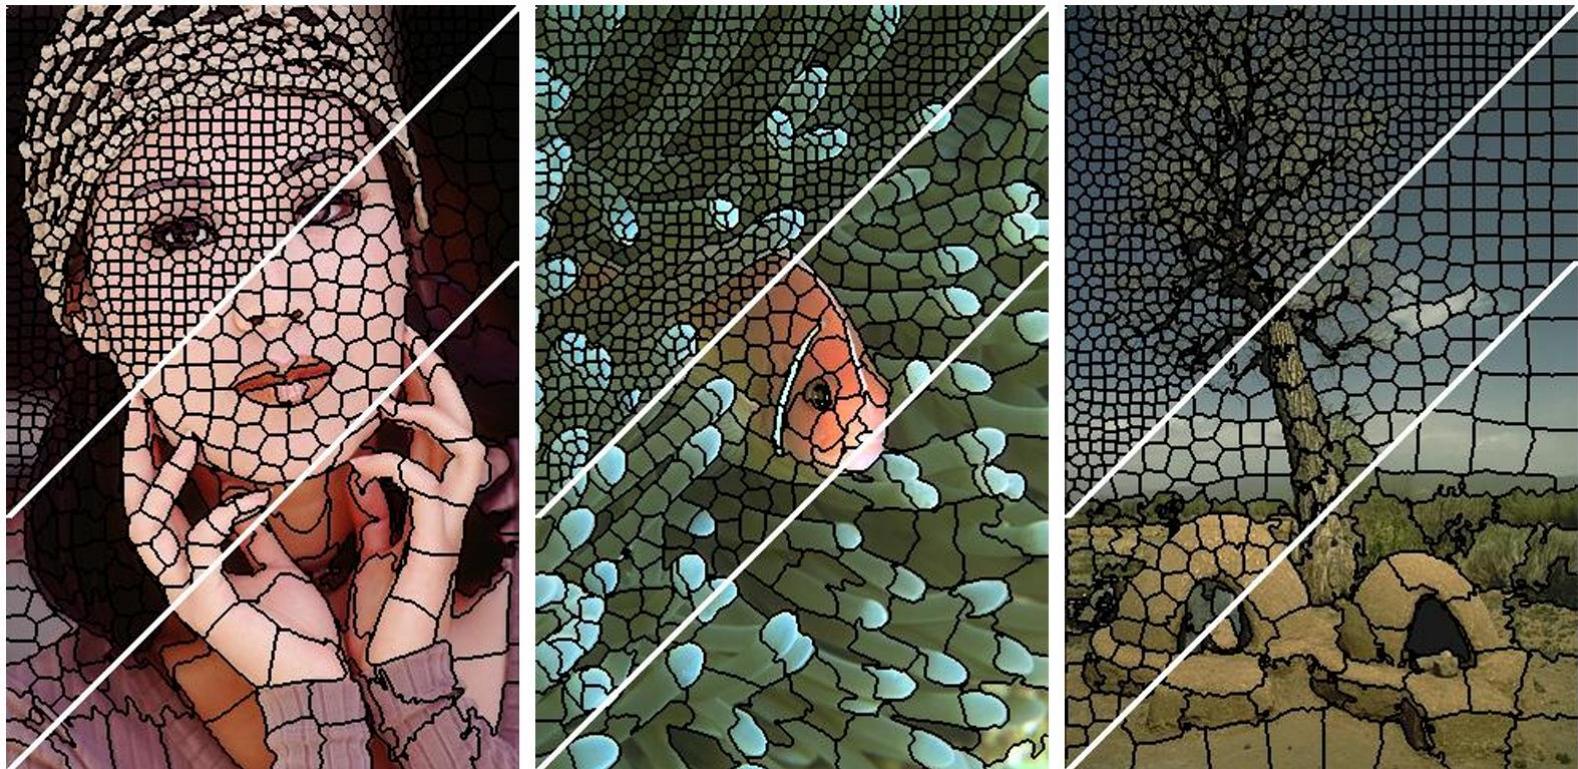
> Vector quantization

- replace each image pixel with the closest representative color



K-means applications

- > Image segmentation

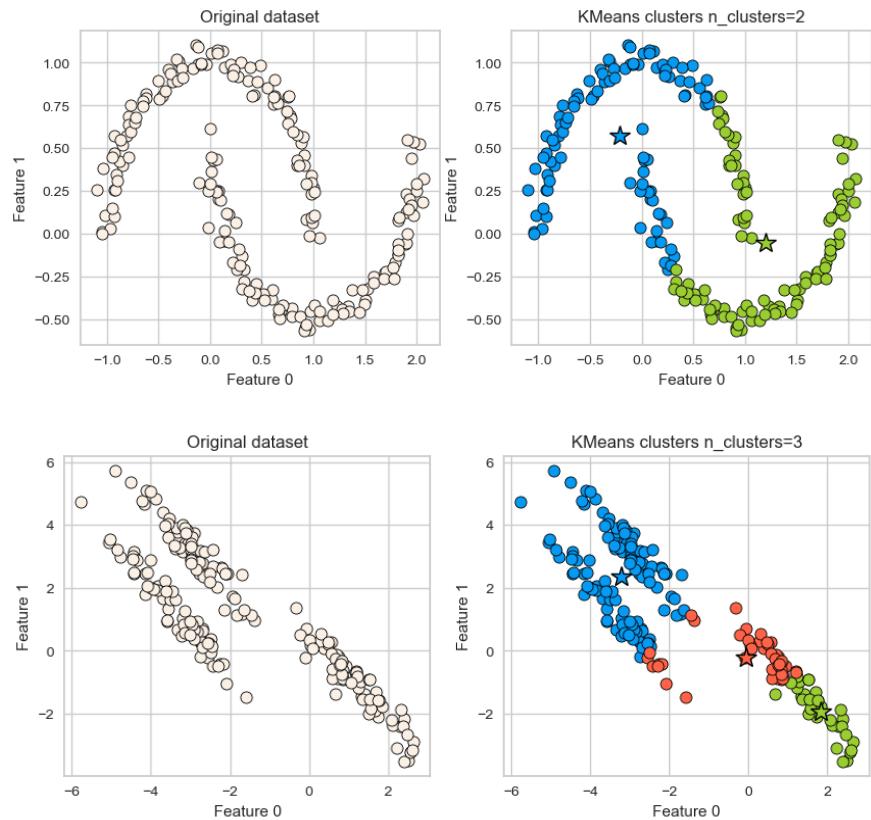
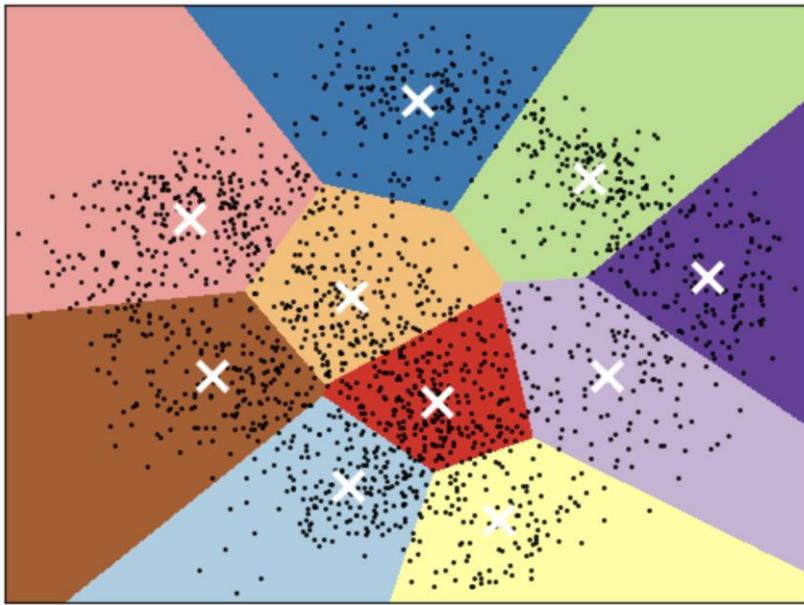


K-means summary

- > To sum up,
 - clustering goal: segment datapoints into similar groups
 - k-means: iteratively optimize cluster assignments and centroids
- > Issues
 - squared Euclidean objective restrictive → K-medoids
 - non-convex optimization problem → K-means++
 - running time → mini-batch K-means
 - must choose k → elbow method, Silhouette analysis, X-means, G-means
- > More issues
 - hard assignments are unstable under small perturbations of data
 - gives equal weight to each coordinates and clusters
 - clusters change arbitrarily with different k
 - works poorly on non-convex / non-spherical clusters
 - different cluster density and size

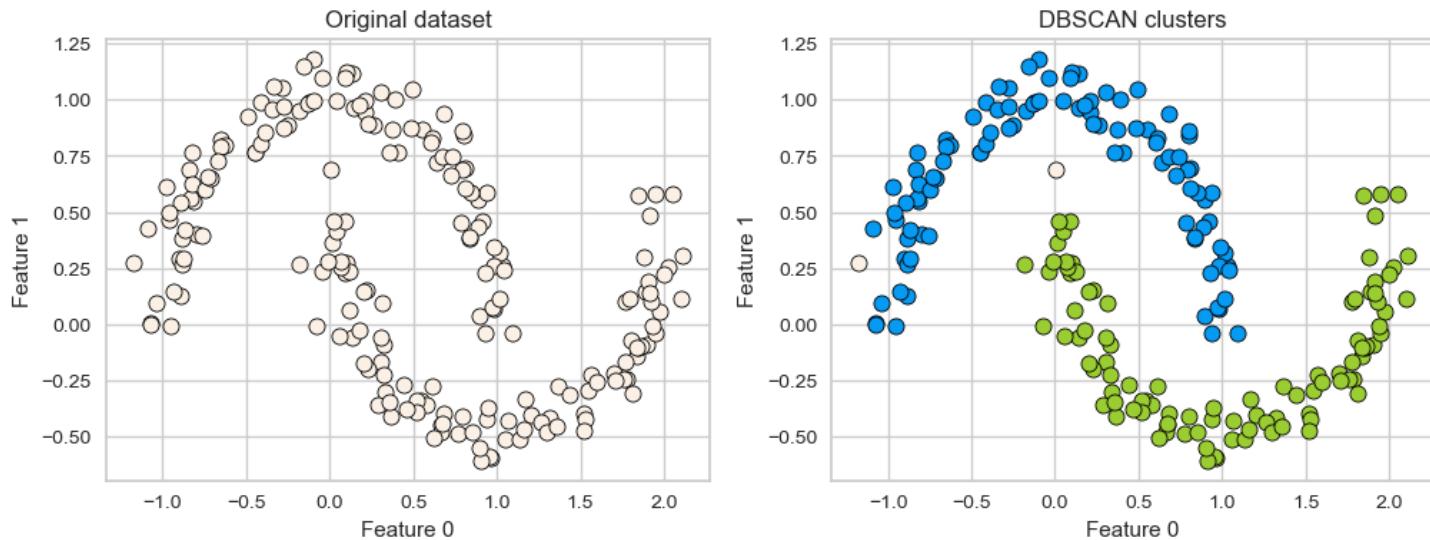
K-means limitations

- > Shape of K-means clusters
 - boundaries between clusters are linear



DBSCAN

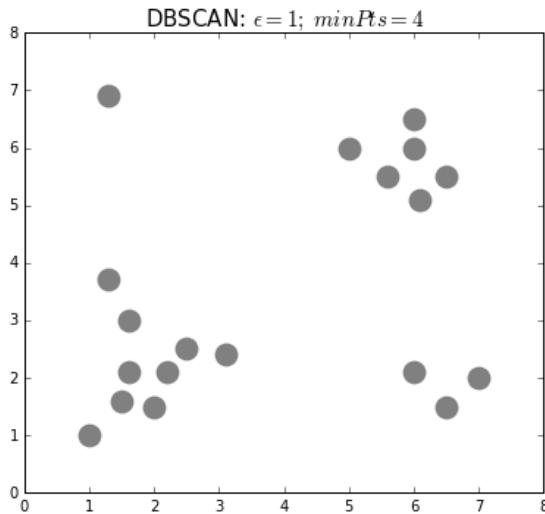
- > Density-based spatial clustering of applications with noise
 - idea: clusters are dense regions in the feature space, so identify them
 - it does not require to specify the number of clusters
 - it can identify points that are not part of any clusters
 - it can capture clusters of complex shapes



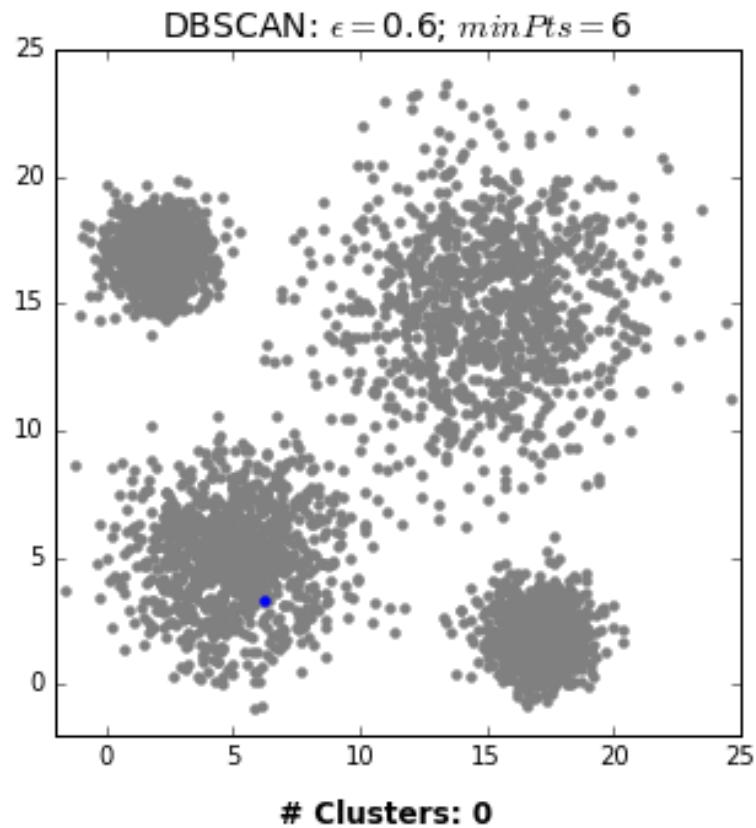
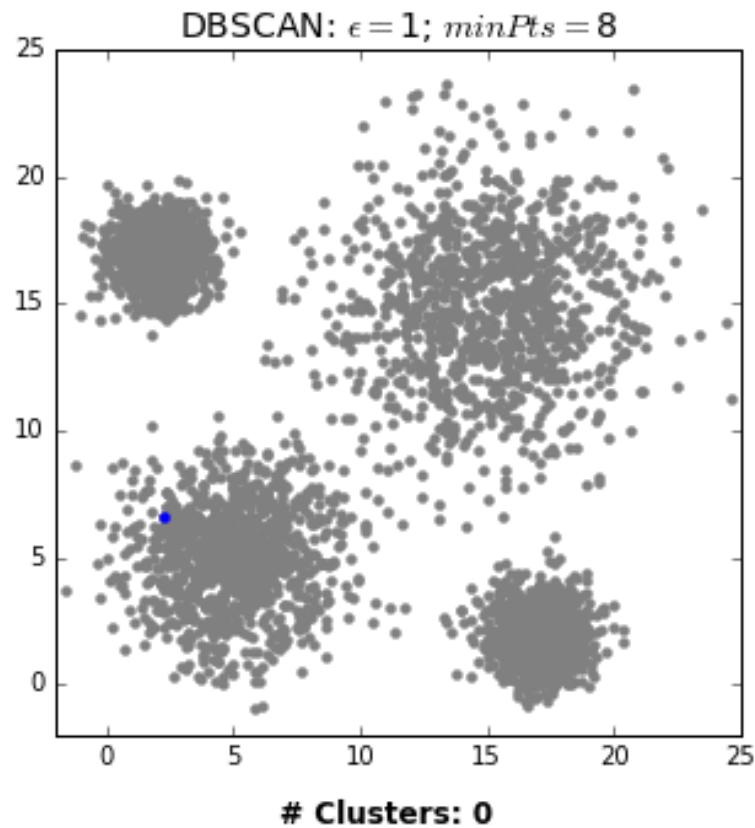
DBSCAN

> Algorithm

- pick a point at random
- check whether the point is a core point
 - core point: that have at least n points within a distance of d
- if the point is a core point, give it a color (label)
- spread the color to all of its neighbors
- check if any of the neighbors is a core point, if yes, spread the color
- once there is no more core point left to spread, pick a new unlabeled point



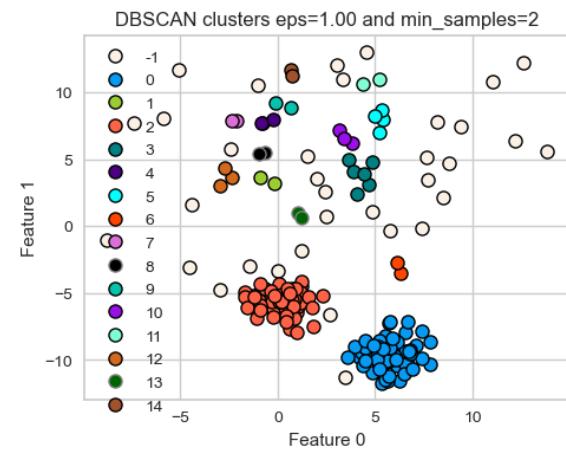
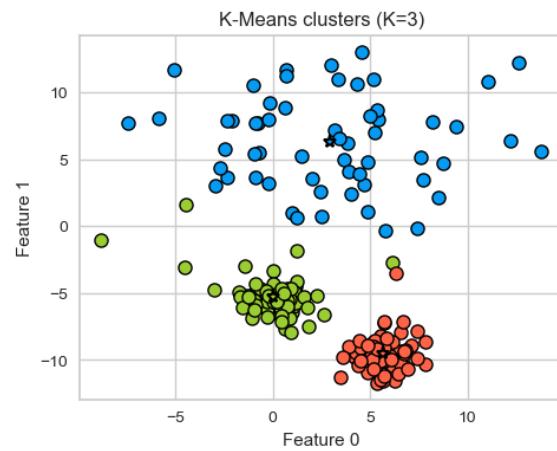
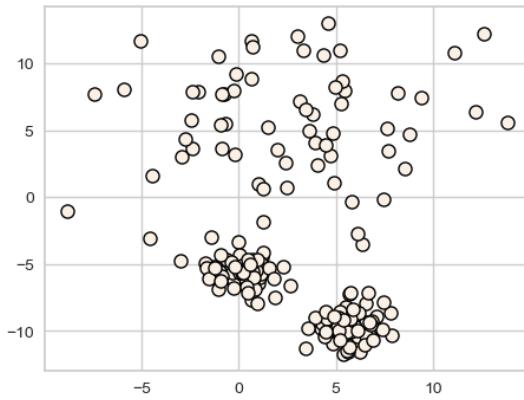
DBSCAN



DBSCAN

> Pros and cons

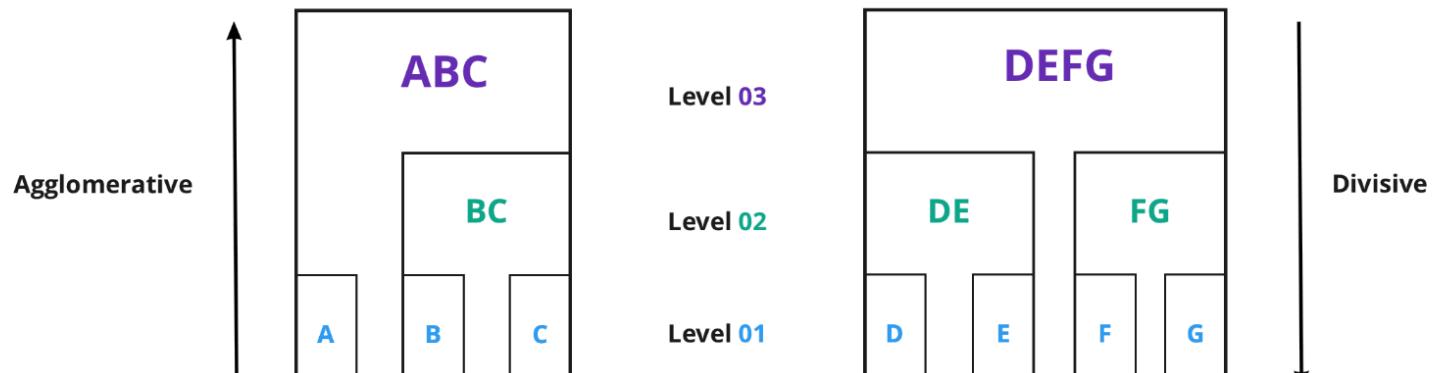
- can learn arbitrary cluster shapes
- can detect outliers
- cannot predict on new examples
- needs tuning of two non-obvious hyperparameters (n, d)
- doesn't do well when we have clusters with different densities



Hierarchical clustering

- > Construct a hierarchy of clusters based on proximity
 - Divisive: start with the one cluster (entire dataset) and split
 - Agglomerative: start with the all individual points (all clusters) and combine

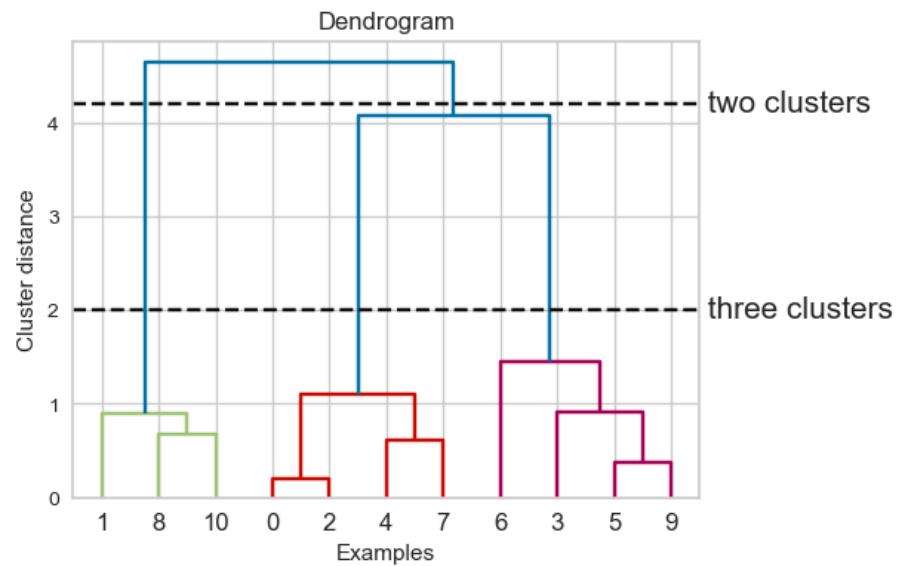
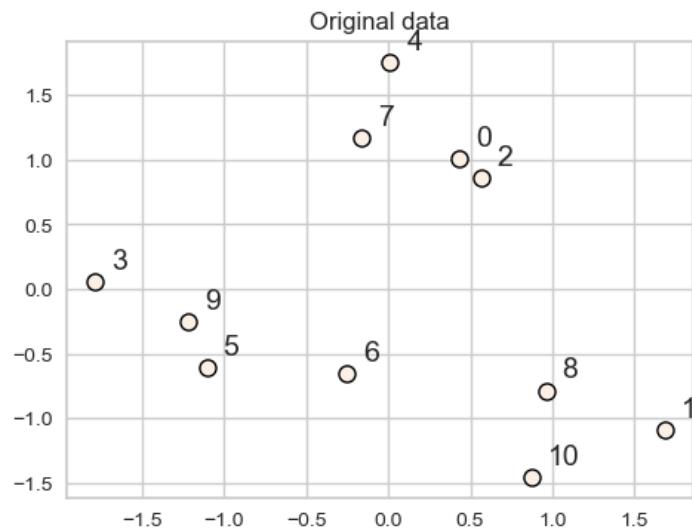
Hierarchical Clustering Algorithm Types



Hierarchical clustering

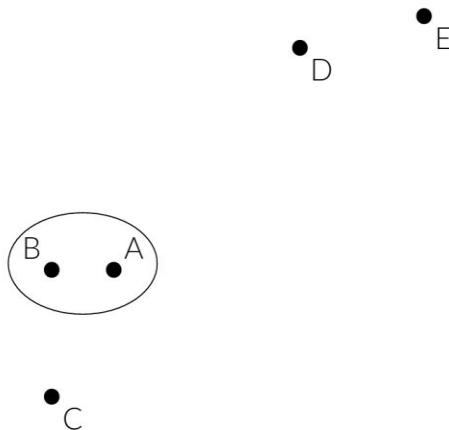
> Dendrogram

- a tree that shows how clusters are merged/split hierarchically
- a clustering is obtained by cutting the dendrogram at the desired level



Hierarchical clustering

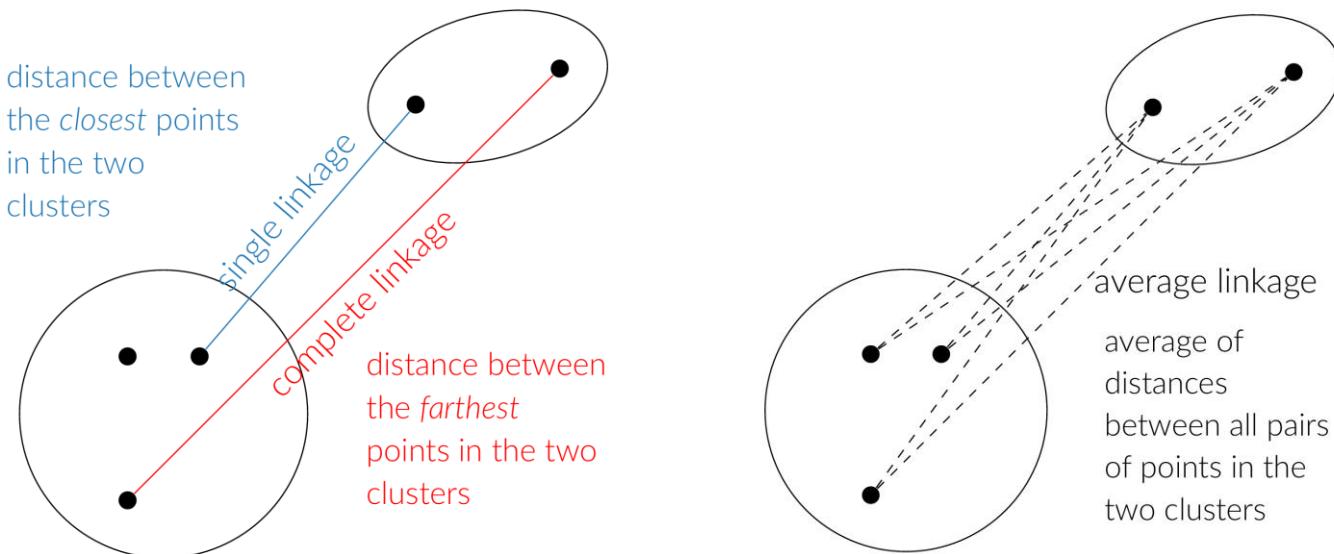
- > Agglomerative
 - first, we merge the two closest points into a cluster
 - next, we want to merge the next closest into a cluster



- > How do we measure the distance between a cluster and a point?
- > How do we measure distance between two clusters?

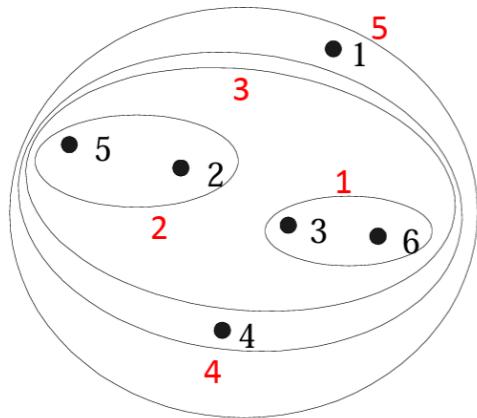
Hierarchical clustering

- > How to measure distances between clusters is called the linkage
 - simple linkage: minimal distance
 - complete linkage: maximum distance
 - average linkage: average distance between all pairs
 - ward linkage: increase in within-cluster variance

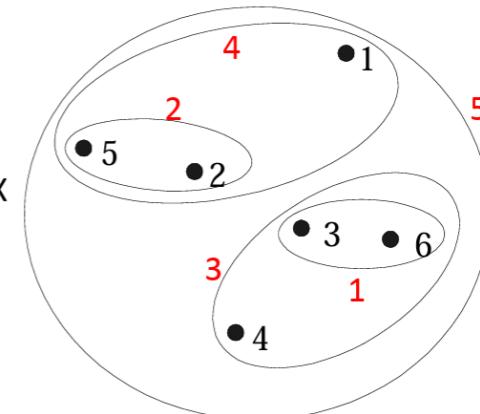


Hierarchical clustering

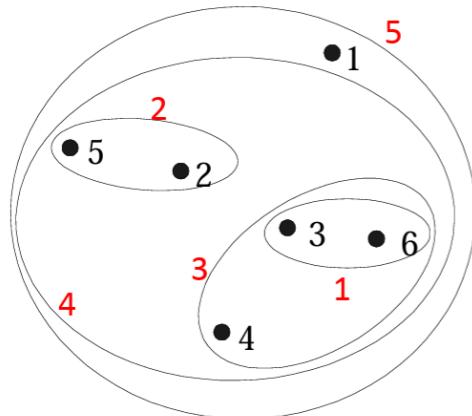
> Example



MIN

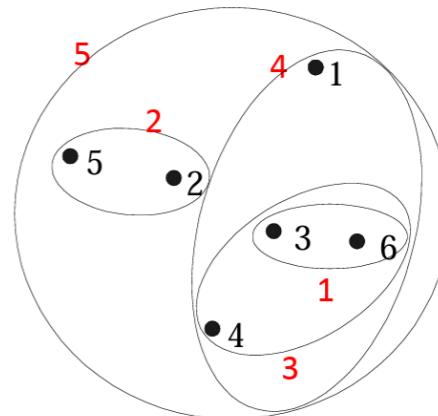


MAX



Group Average

Ward's Method



Hierarchical clustering

- > Pros and cons
 - do not have to assume any particular number of clusters
 - easy to decide the number of clusters by merely looking at the Dendrogram
 - they may correspond to meaningful taxonomies
 - once a decision is made to combine two clusters, it cannot be undone
 - no objective function is directly minimized
 - does not work well on vast amounts of data
 - different measures have problems with one or more:
 - sensitive to noise and outliers
 - breaking large clusters
 - difficulty handling different sized clusters and irregular shapes

Reference

> K-means

- https://web.stanford.edu/~lmackey/stats306b/doc/stats306b-spring14-lecture1_slides.pdf
- <https://www-users.cse.umn.edu/~jwcalder/Clustering.pdf>
- https://www.cs.toronto.edu/~rgrosse/courses/csc411_f18/slides/lec15-slides.pdf

> Other clustering

- http://pajarito.materials.cmu.edu/Data_Analytics-lectures/27737-Clustering-L11-24Mar21.pdf
- https://ubc-cs.github.io/cpsc330-2024W2/lectures/notes/16_DBSCAN-hierarchical.html
- <https://dashee87.github.io/data%20science/general/Clustering-with-Scikit-with-GIFs/>
- https://cse.buffalo.edu/~jing/cse601/fa12/materials/clustering_hierarchical.pdf
- https://dlsun.github.io/stats112/slides/hierarchical_clustering.pdf