

SME3006 Machine Learning – 2025 Fall

# Decision trees and ensemble methods



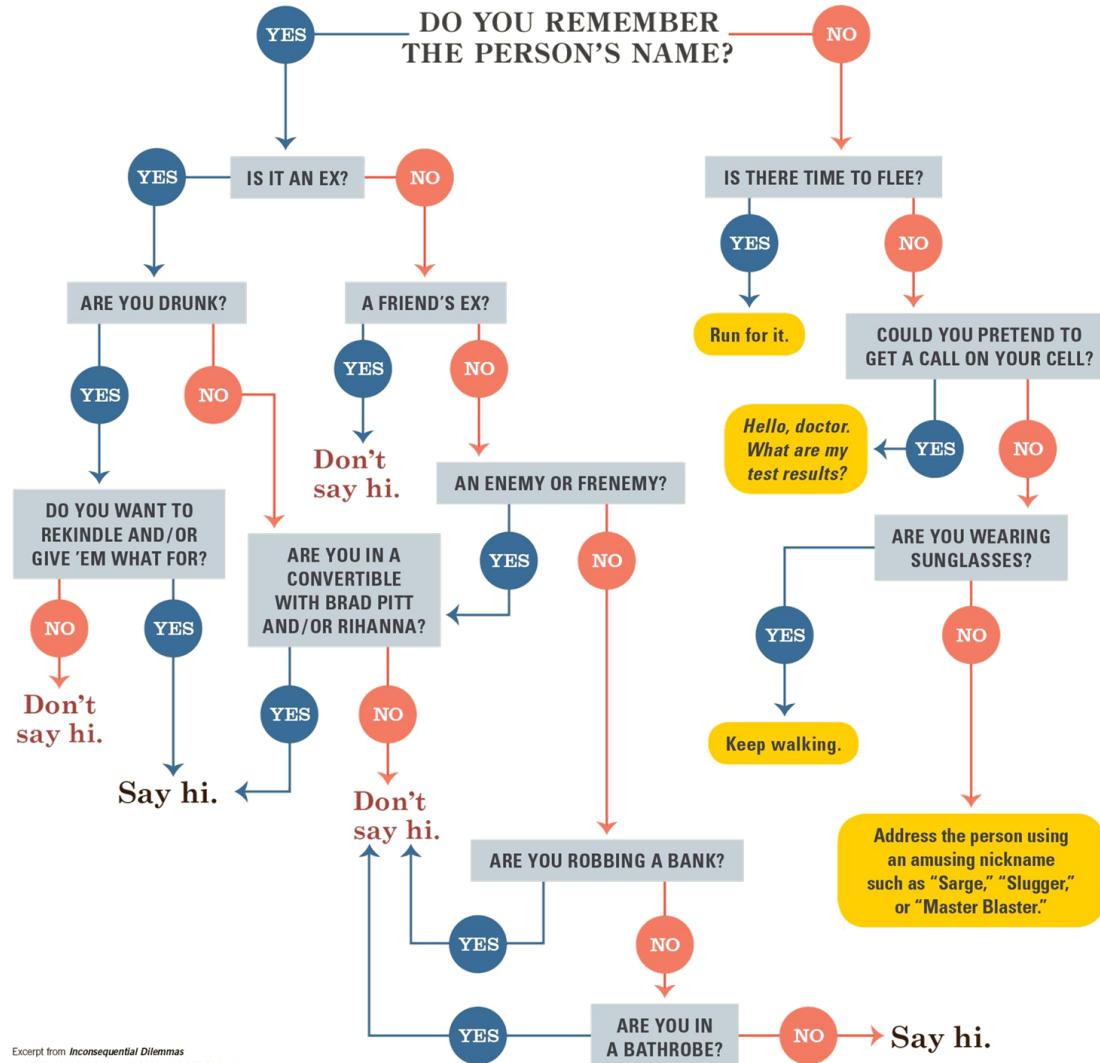
INHA UNIVERSITY

## Overview

- > Why decision trees?
- > limitation of simple models
- > from single tree to ensemble
- > bagging and random forest
- > boosting methods

# Decision tree

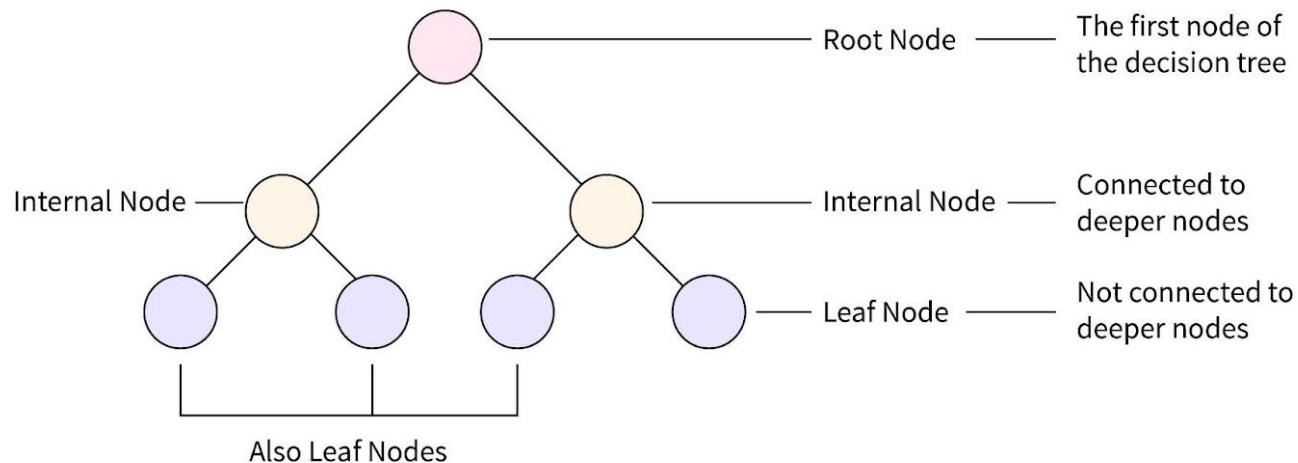
> Do I say hi?



Excerpt from *Inconsequential Dilemmas*  
KNOCKKNOCKSTUFF.COM • © 2013 WHO'S THERE INC.

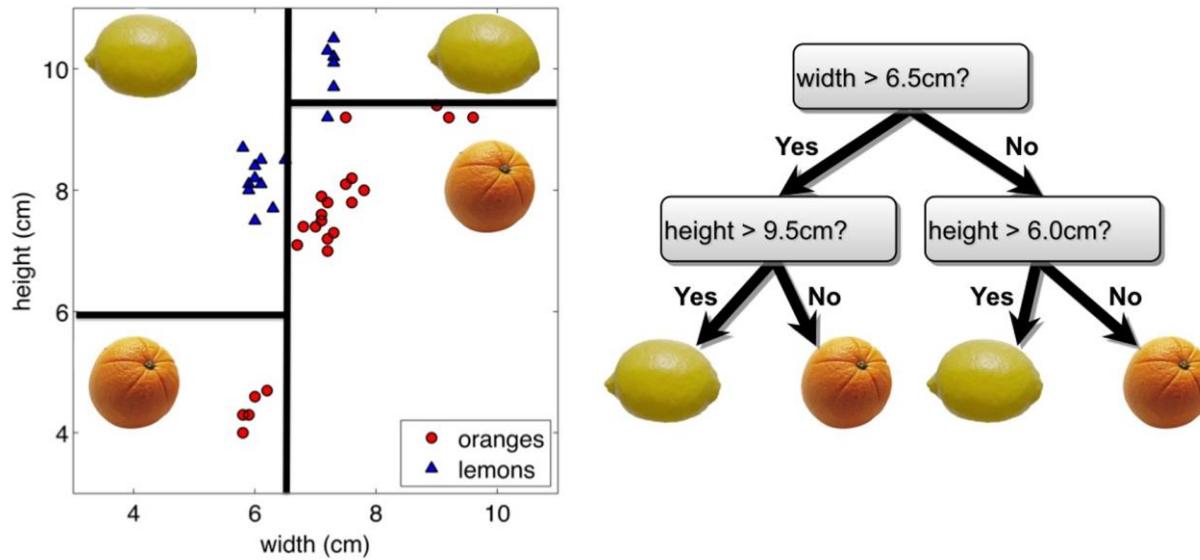
# Decision tree

- > A decision tree (DT) is a tree that
  - at each inner node has a decision rule that assigns instances uniquely child nodes of the actual node
  - at each leaf node has a class label



## Decision tree

- > DT make predictions by recursively splitting on different attributes according to a tree structure
  - decision boundaries are rectangular
- > Example) classifying a lemon and an orange

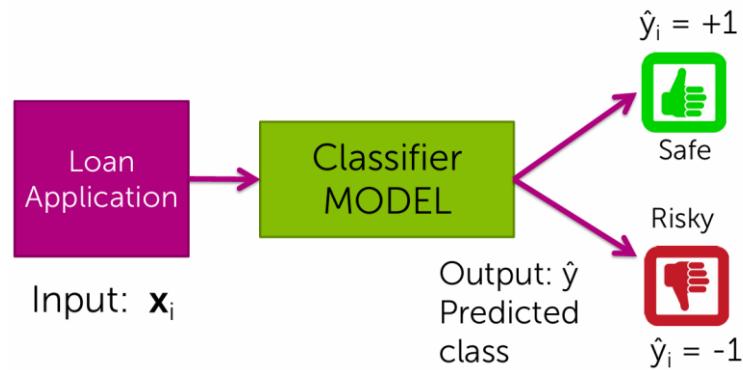


## Decision tree

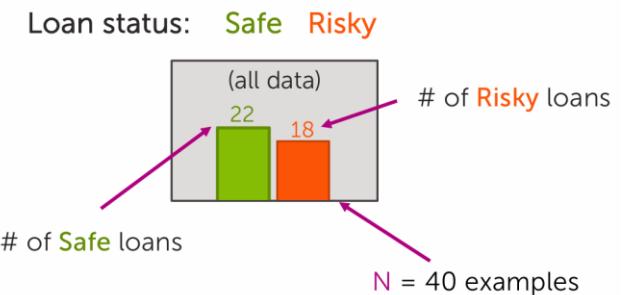
- > How do we construct a useful decision tree?
  - learning the simplest (smallest) decision tree is an NP complete problem
  - number of possible trees increases exponentially
- > Resort to a greedy heuristic:
  - start from an empty decision tree
  - split on the best attribute
  - recurse
- > Which attribute is the best?

# Decision tree

- > Example) buying a new house

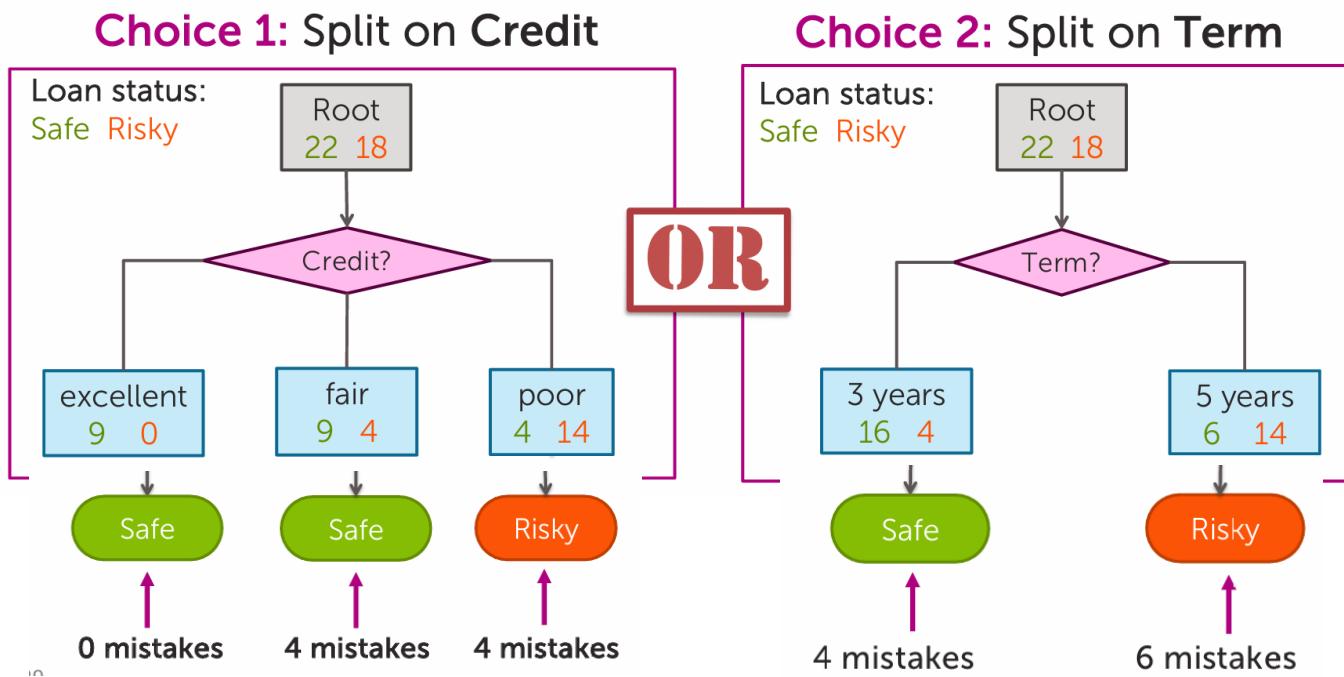


Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe



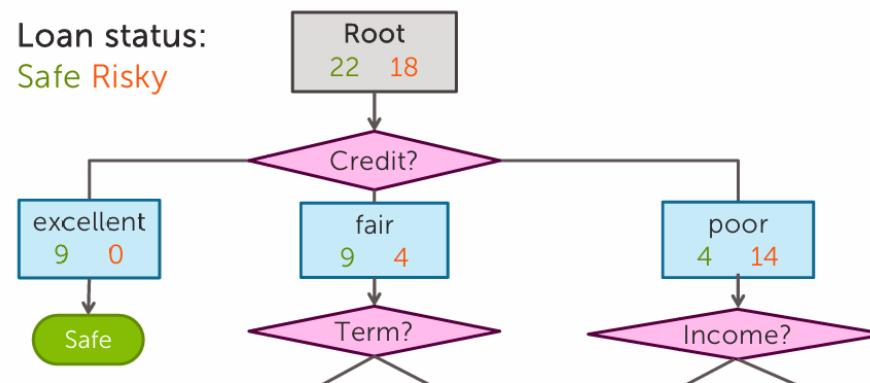
# Decision tree

- > Quality metric:
  - classification error  $Err = \frac{\# \text{ incorrect predictions}}{\# \text{ examples}} \in [0, 1]$
- > What is the best feature to split on



# Decision tree

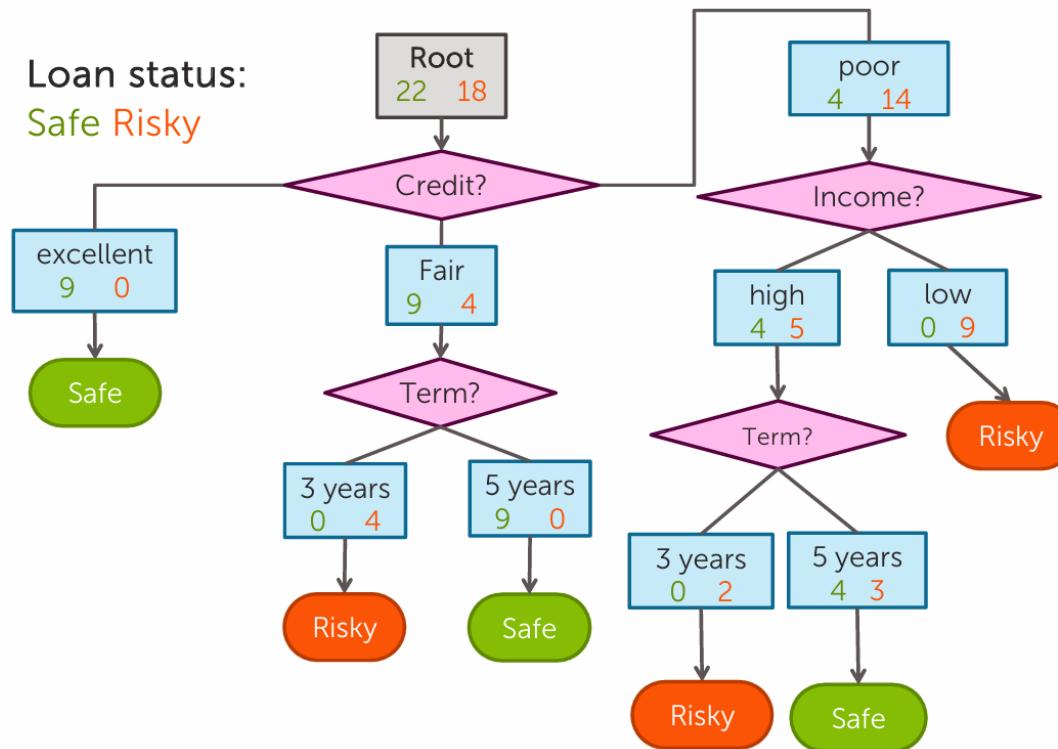
- > Feature split selection algorithm
  - given a subset of data  $M$
  - For each feature  $h_i(x)$ :
    - split data of  $M$  according to feature  $h_i(x)$
    - compute classification error of split
  - Choose feature  $h^*(x)$  with lowest classification error
- > Recursion



# Decision tree

> When does the recursion stop?

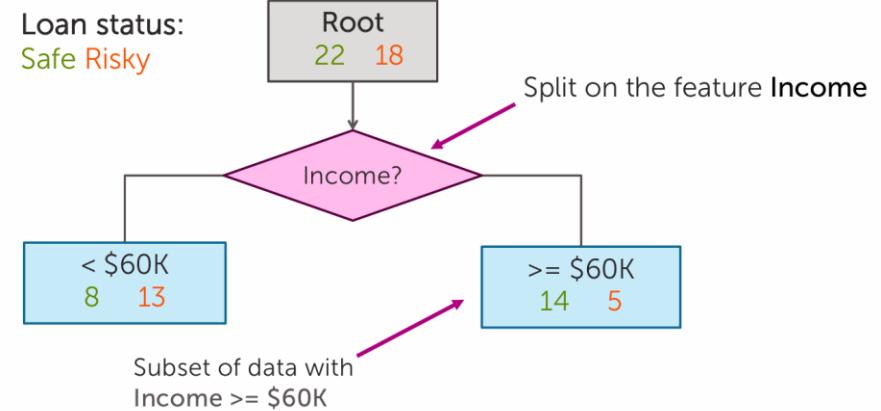
- 1) all data agrees on *output*
- 2) already split on all features
- 3) if no split reduces the classification error



# Decision tree

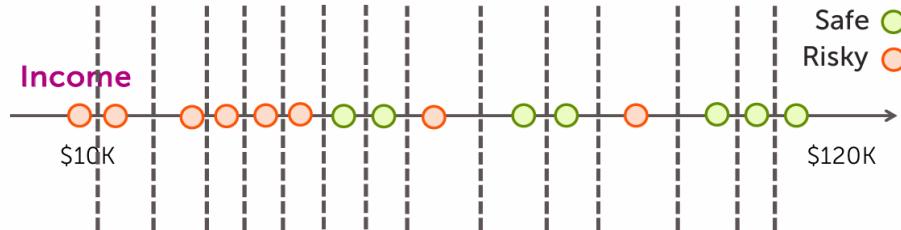
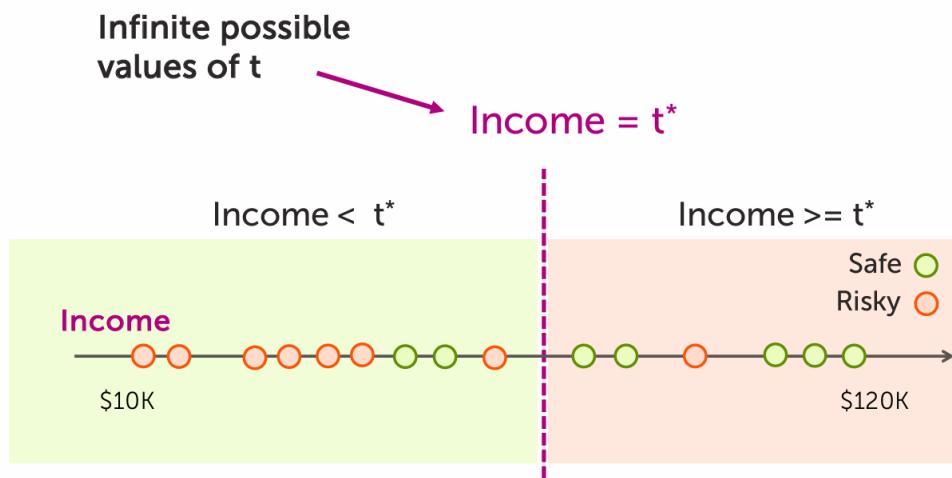
- > What if we have continuous values?
  - threshold split

Income	Credit	Term	y
\$105 K	excellent	3 yrs	Safe
\$112 K	good	5 yrs	Risky
\$73 K	fair	3 yrs	Safe
\$69 K	excellent	5 yrs	Safe
\$217 K	excellent	3 yrs	Risky
\$120 K	good	5 yrs	Safe
\$64 K	fair	3 yrs	Risky
\$340 K	excellent	5 yrs	Safe
\$60 K	good	3 yrs	Risky



# Decision tree

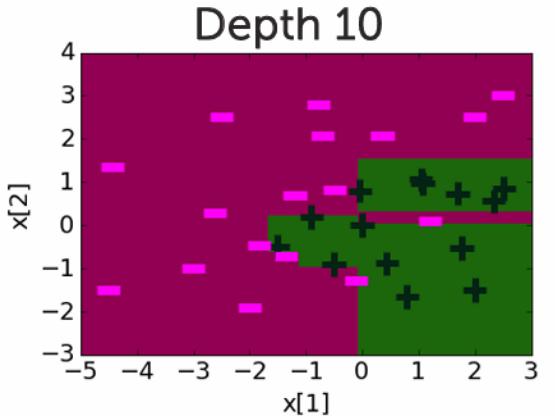
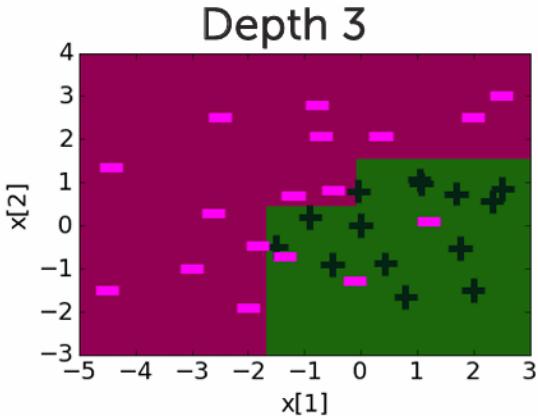
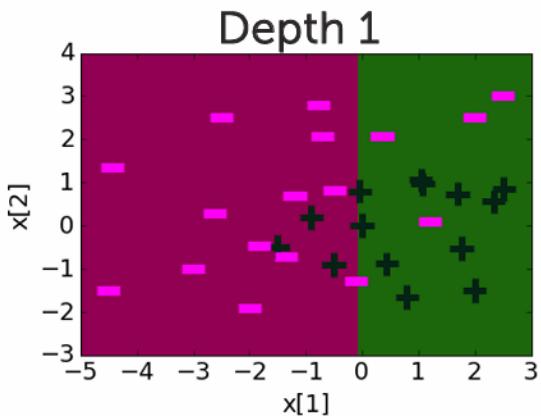
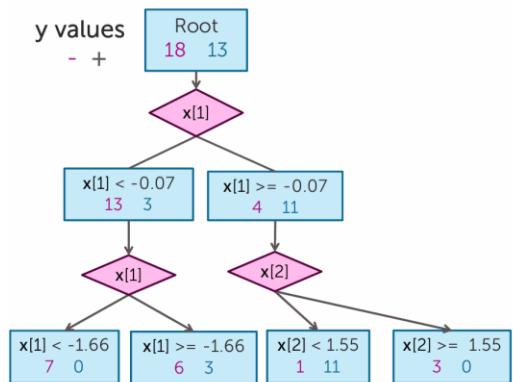
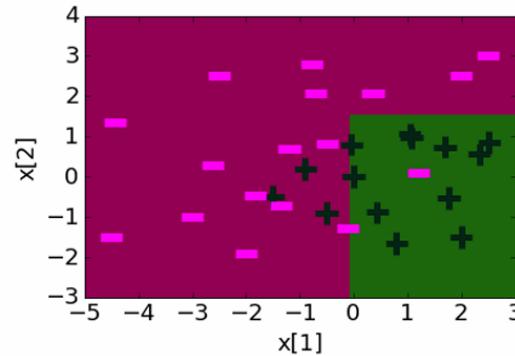
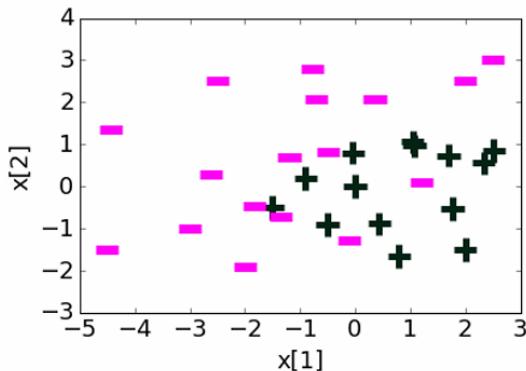
- > What if we have continuous values?
  - threshold split



# Decision tree

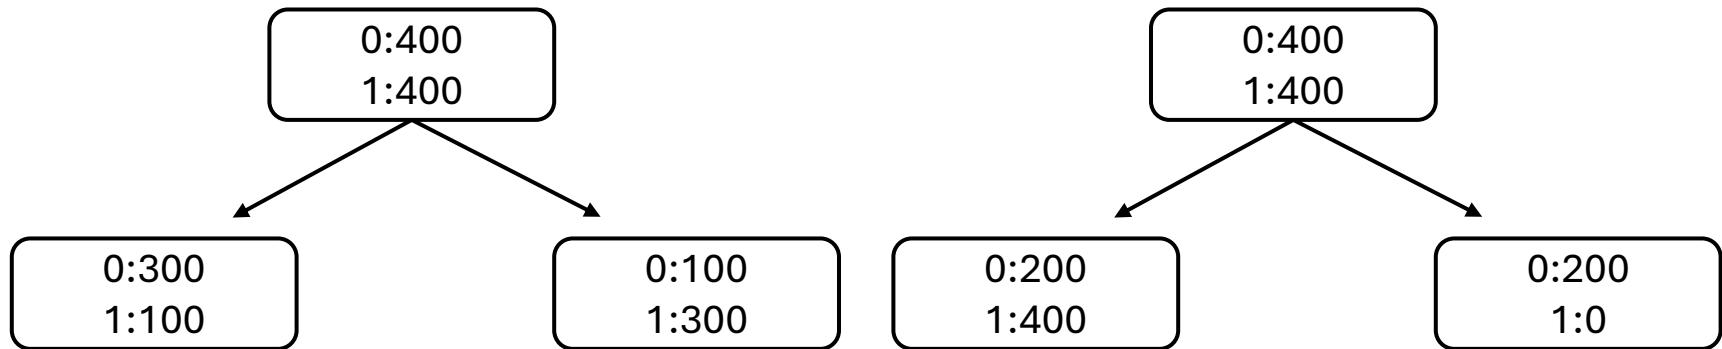
## > Decision boundaries

- for threshold splits, same feature can be used multiple times



## Decision tree

- > It is possible to use classification error as quality criterion, it usually is not a good idea
- > Imagine a dataset with a binary target variable (0/1)
  - we have 400 cases per each( 400 / 400 )
  - assume the following split



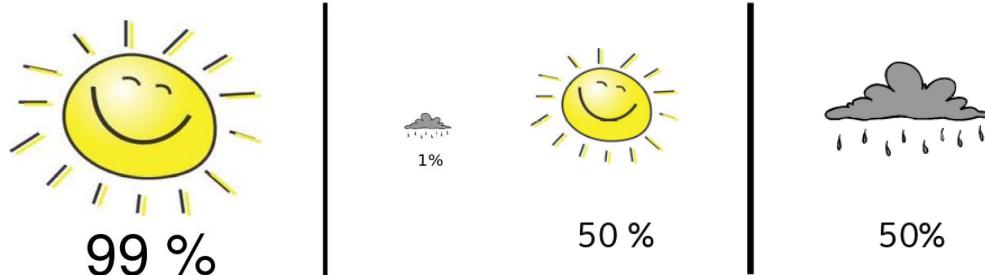
- both have 200 errors (same classification error)
- but the right split is preferred as it contains a pure node

# Information theory

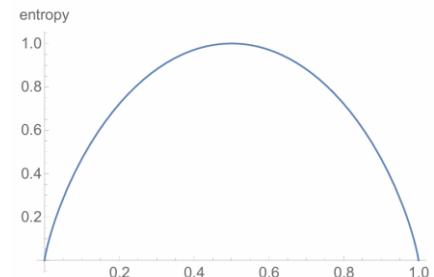
- > We will use techniques from information theory
  - define probability distributions to measure uncertainty
- > Which feature (attribute) is better to split on
  - deterministic (all are true or false) → good
  - uniform distribution (all classes in leaf equally probable) → bad
  - what about distributions in between?

# Information theory

- > Entropy
  - weather forecasts



- which of the two forecasts would give you the most information?
  - first one: you (certainly) know tomorrow's weather
  - second one: you have no idea (the highest uncertainty)
- entropy measures the lack of information or uncertainty
  - maximal if the random distribution is uniform
  - minimal if the random distribution is deterministic



# Information theory

- > Entropy measures the lack of information or uncertainty

$$H(X) = -\sum_{x \in X} p(x) \log_2 p(x)$$

- Why? (Shannon uniqueness theorem, math proof)
  - continuity: small changes in probabilities should cause small changes in  $H$
  - maximal for uniform distribution: when all outcomes are equally likely, uncertainty should be largest
  - additivity for independent events: the total uncertainty of two independent processes should be the sum of their individual uncertainties  
$$H(X, Y) = H(X) + H(Y) \text{ if } X \text{ and } Y \text{ are independent}$$
- high entropy  $\rightarrow$  outcome is very unpredictable  $\rightarrow$  you need information
- zero entropy  $\rightarrow$  outcome is fully predictable  $\rightarrow$  no need information

# Information theory

## > Entropy of a joint distribution

$$H(X, Y) = -\sum_{x,y} p(x, y) \log_2 p(x, y)$$

- we have two random variables and measures the total uncertainty
- weather forecasts
  - predicting Seoul and New York's weather (assume they are independent)  
→ need information for both Seoul and New York,  $H(x, y) = H(x) + H(y)$
  - predicting Seoul and Incheon's weather (they are somewhat dependent)  
→ need information for Seoul and rest for Incheon,  $H(x, y) = H(x) + H(y|x)$

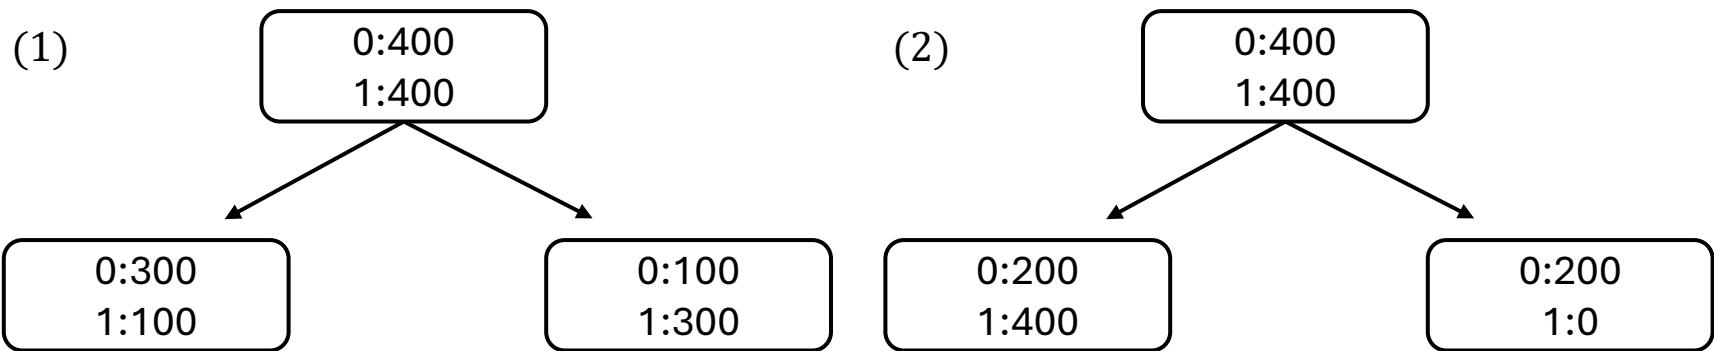
# Information theory

- > Useful properties:
  - $H$  is always non-negative
  - chain rule:  $H(X, Y) = H(X|Y) + H(Y) = H(Y|X) + H(X)$
  - If  $X$  and  $Y$  independent,  $X$  doesn't tell us anything about  $Y$ :  $H(Y|X) = H(Y)$
  - but  $y$  tells us everything about  $y$ :  $H(Y|Y) = 0$
  - by knowing  $x$ , we can only decrease about  $y$ :  $H(Y|X) \leq H(Y)$
- > Information gain
  - how much information do we get by discovering  $x$
  - $IG(Y|X) = H(Y) - H(Y|X)$
  - this is called the information gain or the mutual information
    - $x$  is completely uninformative about  $y$ :  $IG(Y|X) = 0$
    - $x$  is completely informative about  $y$ :  $IG(Y|X) = H(Y)$

# Information theory

$$H(X) = - \sum_{x \in X} p(x) \log_2 p(x)$$

- > Information gain measures the informativeness of a variable
  - what is the information gain of the splits?



- Root entropy:  $H(Y) = -\frac{400}{800} \log_2 \left(\frac{400}{800}\right) \times 2 = 1$
- Leaf entropy (1):  $H(Y|X_{(1)}) = \frac{1}{2} \times \left(-\frac{300}{400} \log_2 \frac{300}{400} - \frac{100}{400} \log_2 \frac{100}{400}\right) \times 2$
- Leaf entropy (2):  $H(Y|X_{(2)}) = \frac{3}{4} \times \left(-\frac{400}{600} \log_2 \frac{400}{600} - \frac{200}{600} \log_2 \frac{200}{600}\right) - \frac{1}{4} \times \left(\frac{200}{200} \log_2 \frac{200}{200}\right)$
- Information gain:  $IG_{(1)} = H(Y) - H(Y|X_{(1)}) = 0.1887$   
 $IG_{(2)} = H(Y) - H(Y|X_{(2)}) = 0.3113$

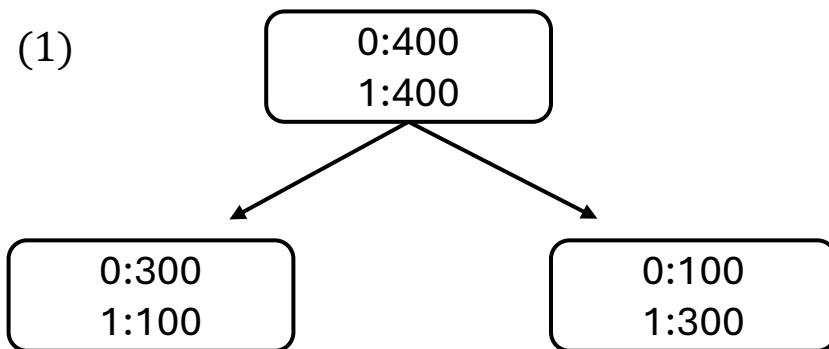
# Information theory

## > Quadratic entropy (Gini impurity)

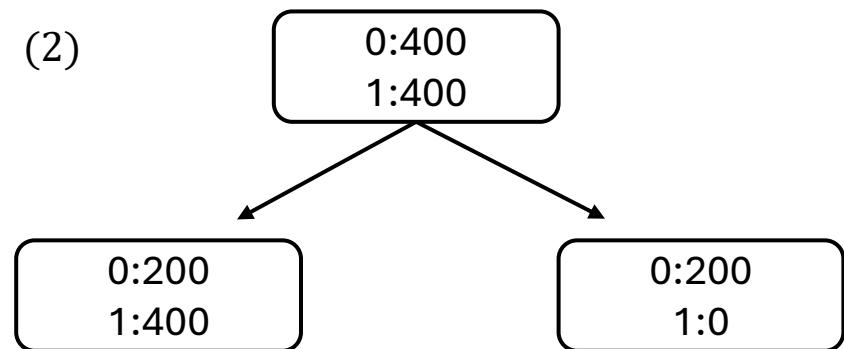
$$H(X) = \sum_{x \in X} p(x)(1 - p(x)) = 1 - \sum_{x \in X} p(x)^2$$

- measure probability of misclassification
- only multiplications and additions (no logarithm)

(1)



(2)



Gini-Impurity

$$\begin{aligned} &= \frac{1}{2}\left(\left(\frac{3}{4}\right)^2 + \left(\frac{1}{4}\right)^2\right) + \frac{1}{2}\left(\left(\frac{3}{4}\right)^2 + \left(\frac{1}{4}\right)^2\right) \\ &= 0.625 \end{aligned}$$

Gini-Impurity

$$\begin{aligned} &= \frac{3}{4}\left(\left(\frac{1}{3}\right)^2 + \left(\frac{2}{3}\right)^2\right) + \frac{1}{4}(1^2 + 0^2) \\ &\approx 0.667 \end{aligned}$$

# Decision tree

- > Concept: split data to reduce uncertainty about the target variable
  - what makes a good split? → good criteria
  - which split strategy to use (multiway vs. binary)
  - how to stop splitting? (pruning needed to avoid overfitting)
  - how to deal with incomplete / missing data?
- > Algorithms
  - ID3: multiway split, Information gain
  - C4.5: Gain ratio (fixed IG bias), continuous features, pruning, missing value handling
  - CART: binary splits only, Gini impurity for classification, variance reduction for regression, cost-complexity pruning, surrogate splits

# Ensemble methods

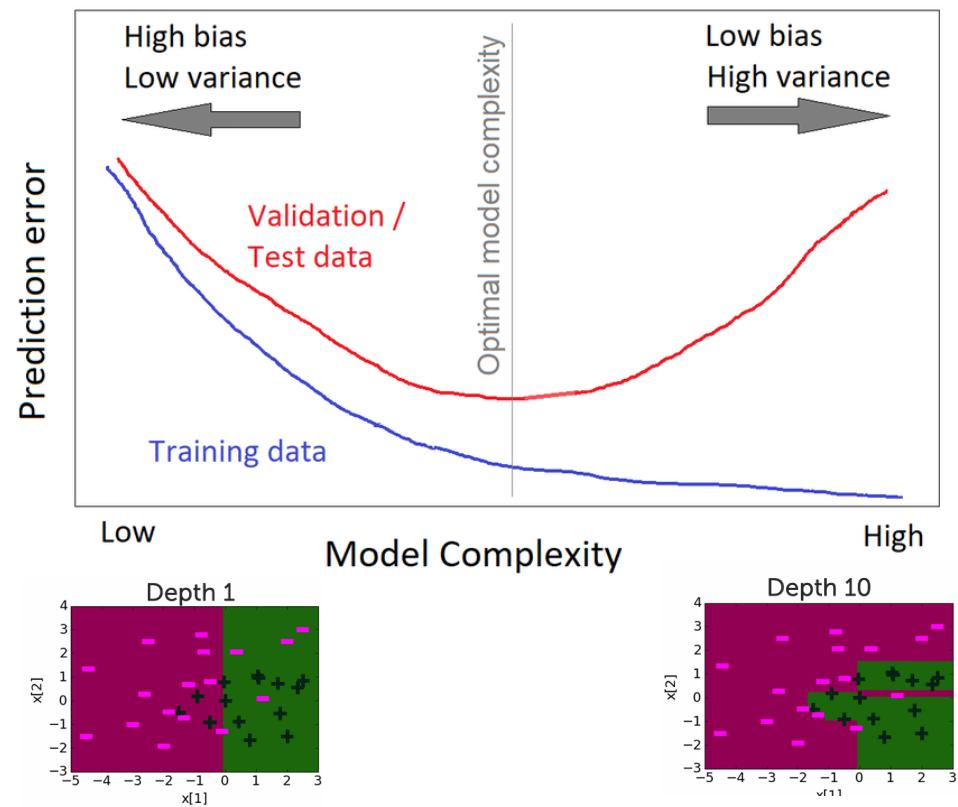
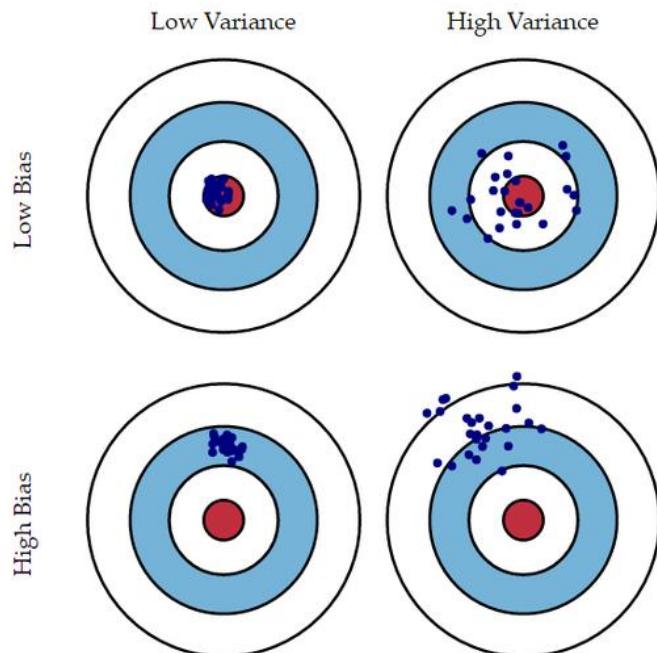
- > Tree-based models
  - interpretable
  - can capture non-linear relationships
  - don't require scaling of the data
  - but single decision trees are likely to overfit
- > Ensembles
  - key idea: combine multiple machine learning models to create more powerful methods
  - can be applied to almost any learning algorithms
  - particularly well suited to decision trees

## Ensemble methods

- > An ensemble of predictors is a set of predictors whose individual decisions are combined in some way to predict new examples
  - E.g., (weighted) majority vote
- > For this to be nontrivial, the learned hypotheses must differ somehow
  - different algorithm
  - different choice of hyperparameters
  - trained on different data
  - trained with different weight of the training examples
- > Bagging: train classifiers independently on random subsets of data
- > Boosting: train classifiers sequentially, focusing on examples that the previous ones got wrong

# Ensemble methods

- > Bias and variance
  - one way to reduce model variance is averaging

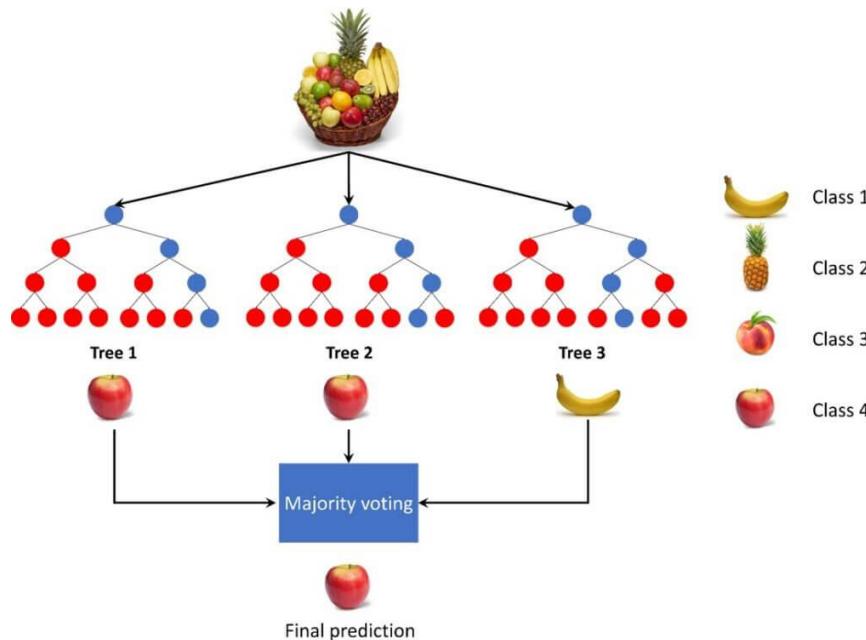


## Bagging

- > Suppose we could somehow sample  $m$  independent training sets
  - we could then compute the prediction  $y_i$  based on each one
  - bias: unchanged,  $\mathbb{E}[y] = \mathbb{E}\left[\frac{1}{m} \sum_i y_i\right] = \mathbb{E}[y_i]$
  - variance: reduced,  $Var[y] = Var\left[\frac{1}{m} \sum_i y_i\right] = \frac{1}{m^2} \sum_i Var[y_i] = \frac{1}{m} Var[y_i]$
- > In practice, we don't have access to the data generating distribution
  - solution: bootstrap aggregation or bagging
  - from a single dataset  $\mathcal{D}$ , generating  $m$  new datasets, each by sampling  $n$  training examples from  $\mathcal{D}$
  - average the predictions of models trained on each of these datasets

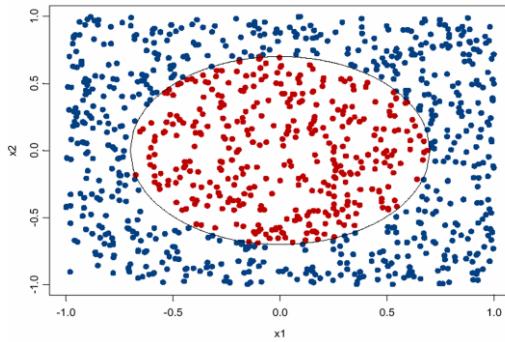
# Bagging

- > Random forests = bagged decision trees, with one extra trick to decorrelate the predictions
  - when choosing each node of the decision tree, choose a random set of  $d$  input features, and only consider splits on those features
  - random forests work well – one of the most widely used algorithms
  - final prediction: average (regression) or majority vote (classification)



# Bagging

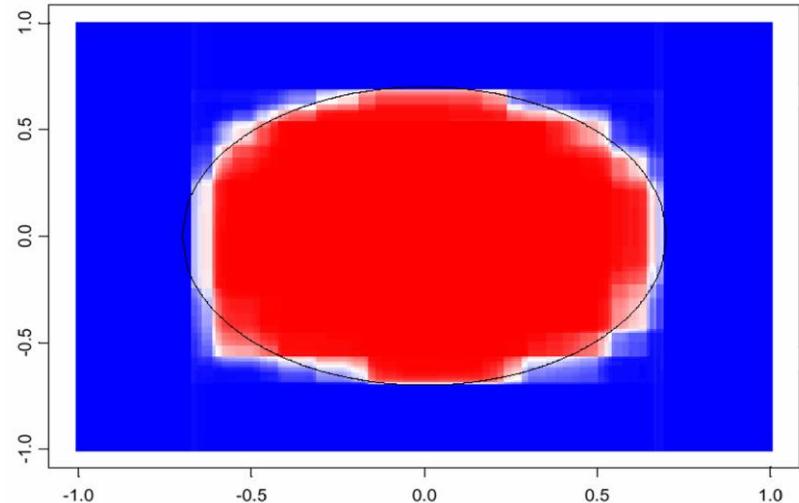
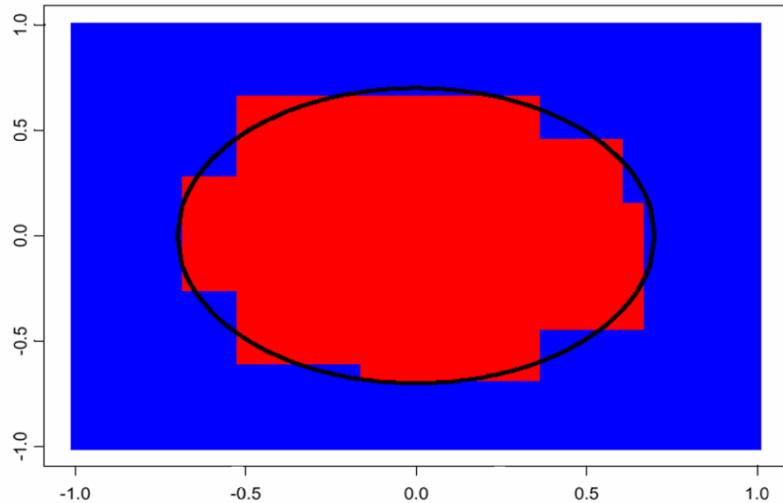
- > Example: classifying points as being inside a circle



CART decision boundary

vs.

100 bagged trees

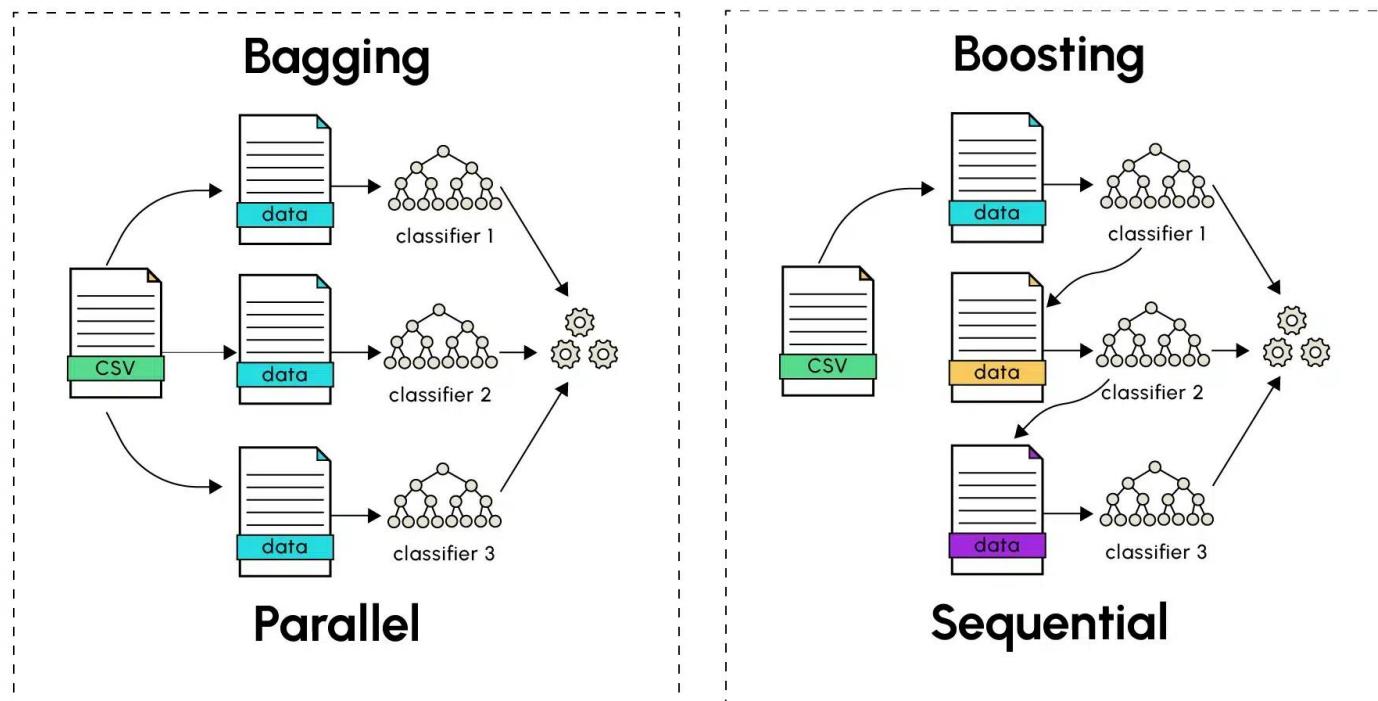


# Bagging

- > Random forest reduces overfitting by averaging predictions
  - each tree overfits on some part of the data, but we can reduce overfitting by averaging the results (can be shown mathematically)
  - work well without heavy tuning of hyperparameters
  - even if a single model is great, a small ensemble usually helps
- > Limitations
  - require more memory
  - hard to interpret
  - does not reduce bias
  - there is still correlation between classifiers

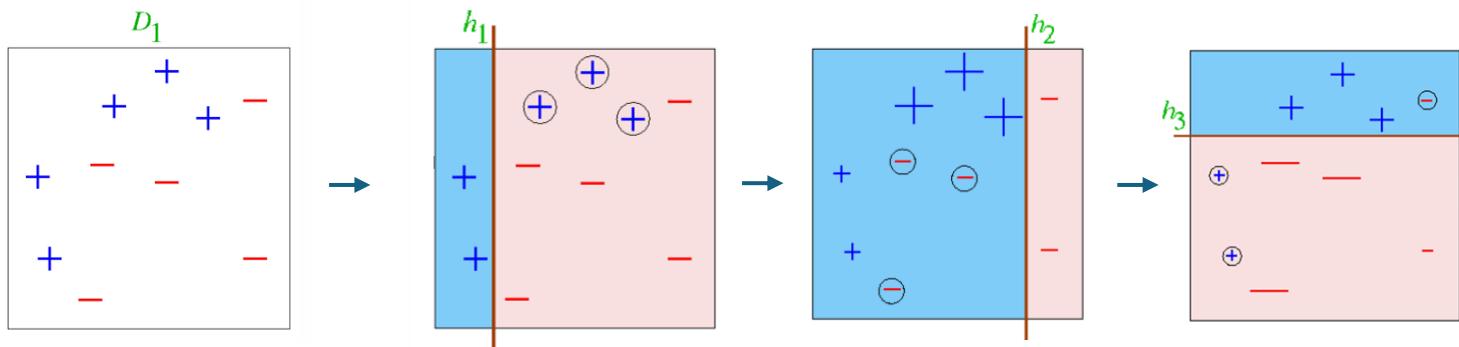
# Boosting

- > Weak learners are trained sequentially, each one focusing more on the mistakes of the previous ones
  - boosting is strong at reducing bias (use simple, weak learners)
  - can a set of weak learners to be combined to create a stronger learner?



# Boosting

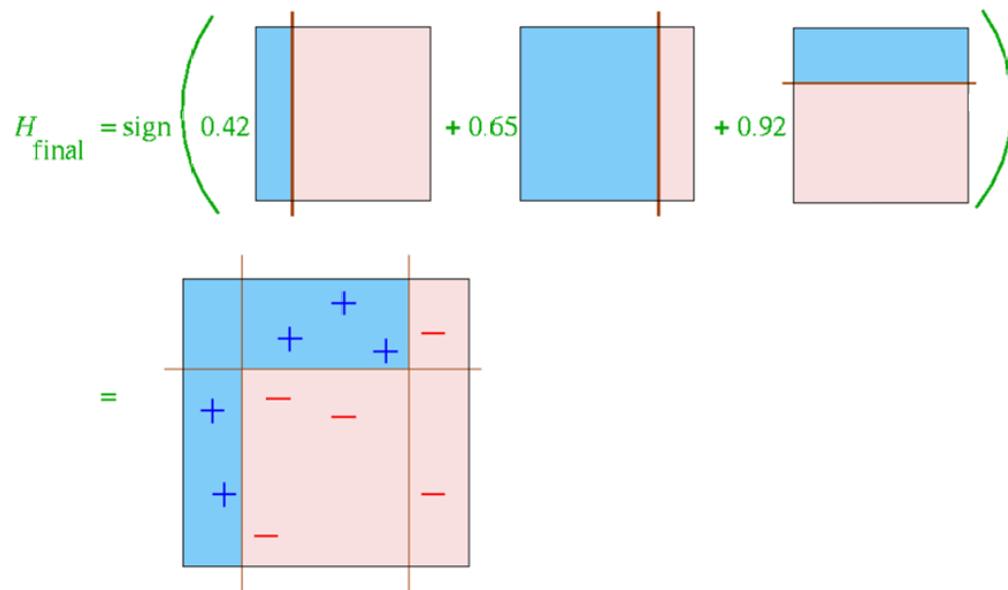
- > Classifier tries harder on examples with higher cost
  - misclassification error  $\sum_i \mathbb{I}[h(x^i) \neq t^i]$
  - weighted error  $\sum_i w^i \mathbb{I}[h(x^i) \neq t^i]$
- Iterate
  - at each iteration we re-weight the training samples by assigning larger weights to samples that were classified incorrectly
  - we train a new weak classifier and add to the ensemble



# Boosting

## > Final classifier

- weighted sum of each weak learners
- training error converges to zero
  - if each weak learner must have error rate under 0.5 on the weighted data



# Boosting

- > It is quite resilient to overfitting, though it can overfit
- > Types
  - AdaBoost (1996): assigns weights to data points and trains new models on updated weights
  - Gradient boosting (1999): fits new model to the residual errors of previous models
  - Stochastic gradient boosting (2002): uses random subsets of training data and features for each new model
  - XGBoost (2016): optimized gradient boosting with regularization and parallelization
  - LightBGM (2017): gradient boosting with leaf-wise growth and histogram-based splits

# Ensemble methods

- > Ensembles combine models to improve performance
- > Bagging
  - reduce variance (large ensemble can't cause overfitting)
  - bias is not changed (much)
  - parallel
  - need to minimize correlation between ensemble elements
- > Boosting
  - reduce bias
  - increase variance (large ensemble can cause overfitting)
  - sequential
  - high dependency between ensemble elements

# Reference

- > Decision trees
  - <https://cs229.stanford.edu/notes2022fall/decision-trees.pdf>
  - <https://people.csail.mit.edu/dsontag/courses/ml16/slides/lecture11.pdf>
  - <https://www.ismll.uni-hildesheim.de/lehre/ml-09w/script/ml-04-decisiontrees-2up.pdf>
  - [https://www.cs.toronto.edu/~mren/teach/csc411\\_19s/lec/lec03.pdf](https://www.cs.toronto.edu/~mren/teach/csc411_19s/lec/lec03.pdf)
- > Ensemble methods
  - [https://ubc-cs.github.io/cpsc330-2024W2/lectures/notes/12\\_ensembles.html](https://ubc-cs.github.io/cpsc330-2024W2/lectures/notes/12_ensembles.html)
  - <https://cs229.stanford.edu/notes2022fall/boosting.pdf>
  - [https://www.cs.toronto.edu/~mren/teach/csc411\\_19s/lec/lec04.pdf](https://www.cs.toronto.edu/~mren/teach/csc411_19s/lec/lec04.pdf)
  - [https://www.cs.toronto.edu/~mren/teach/csc411\\_19s/lec/lec05\\_matt.pdf](https://www.cs.toronto.edu/~mren/teach/csc411_19s/lec/lec05_matt.pdf)