

SME3006 Machine Learning – 2025 Fall

# Hyperparameter tuning and gradient-free optimization



INHA UNIVERSITY

# Optimization problem

- > Finding the minimizer of a function subject to constraints:

$$\underset{x}{\text{minimize}} J(x)$$

objective function

$$\begin{aligned}s.t. f_i(x) \leq 0, \quad i = \{1, \dots, k\} \\ h_j(x) = 0, \quad j = \{1, \dots, l\}\end{aligned}$$

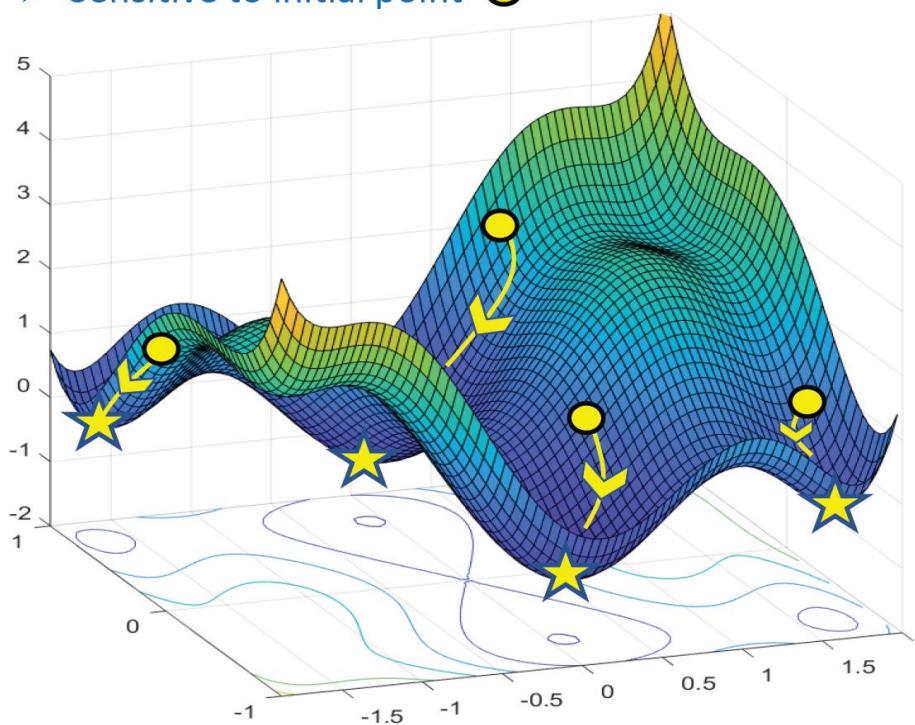
constraints

- Linear regression:  $\underset{w}{\text{minimize}} \|Xw - y\|^2$
  - Classification (logistic regression):  $\underset{w}{\text{minimize}} \sum_{i=1}^n \log(1 + \exp(-y_i x_i^\top w))$
  - Maximum likelihood estimation:  $\underset{\theta}{\text{maximize}} \sum_{i=1}^n \log p_\theta(x_i)$
  - K-means:  $\underset{\mu}{\text{minimize}} \sum_{j=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$
- ...

# Optimality conditions

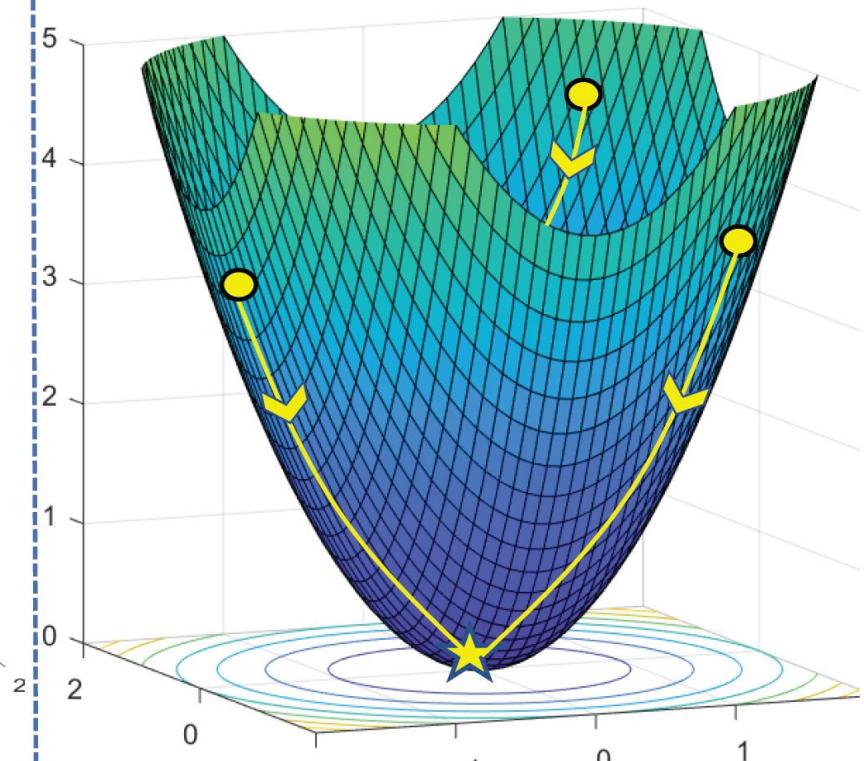
## Nonconvex Optimization

- Multiple local minima ⭐
- Sensitive to initial point ⚡



## Convex Optimization

- Unique minimum: global/local



# Hyperparameter tuning

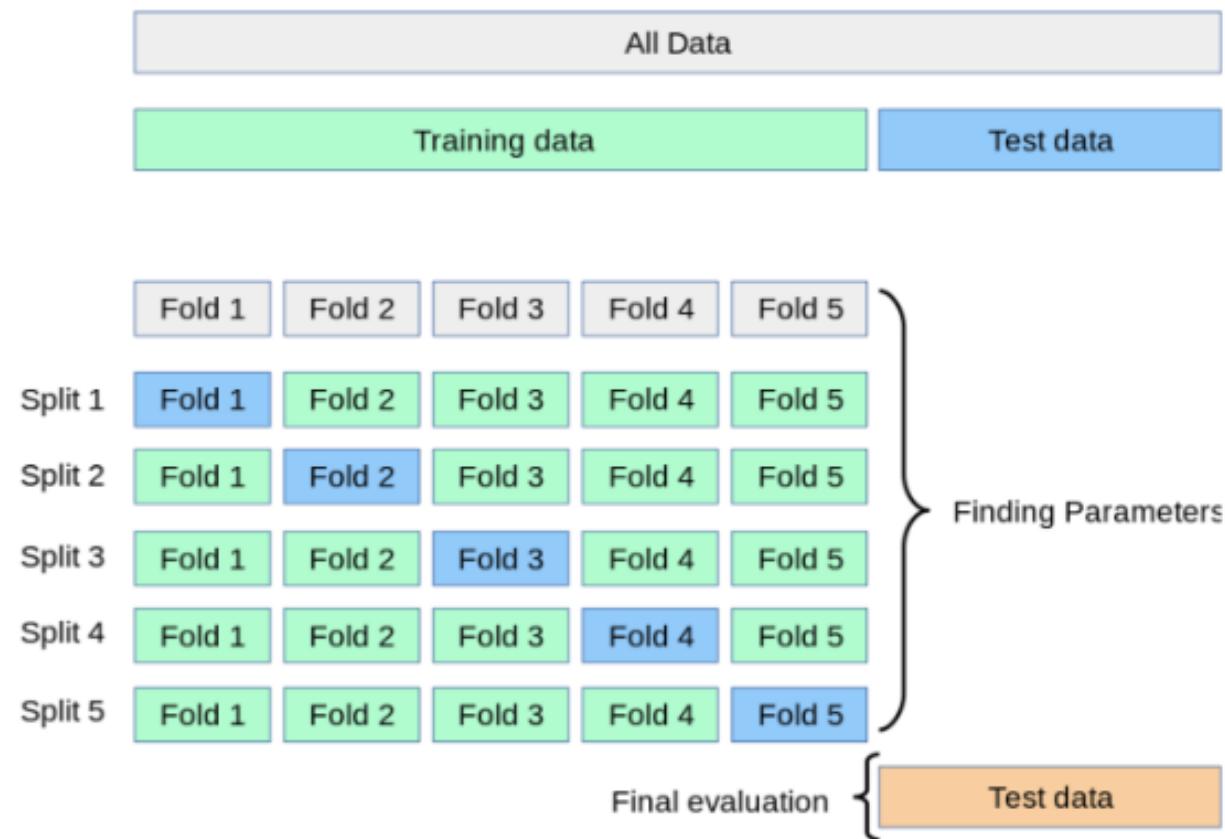
- > Gradient-free (derivative-free) optimization
  - When gradients cannot be computed
    - discontinuous or non-smooth design space (discrete or mixed-integer variables)
    - hyperparameters of machine learning models
      - e.g. learning rate, kernel type, regularization choice
    - black-box simulators with no analytic derivates
  - When following gradients is inefficient or unreliable
    - computing/approximating gradient is too costly relative to function evaluations
    - objectives/constraints are noisy or stochastic
    - complex landscapes with many local minima where gradients mislead search
  - Gradient-free methods have a higher chance of finding the global optimum

# Hyperparameter tuning

- > Model selection
  - which input features to include
  - what preprocessing to do
  - what machine learning method to use
- > Hyperparameter tuning
  - for k-nearest neighbors, hyperparameters include:
    - k
    - metric (e.g. Euclidean distance)
- > We use cross-validation and pick the model / hyperparameter with the smallest test error

# Hyperparameter tuning

- > Cross-validation
  - splitting the data into multiple train/validation sets



# Hyperparameter tuning

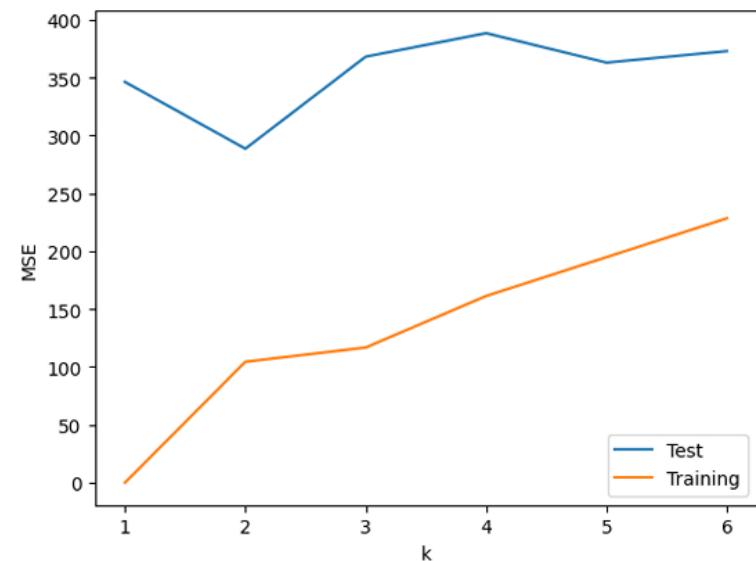
- > Wine quality prediction



# Hyperparameter tuning

## > Wine quality prediction

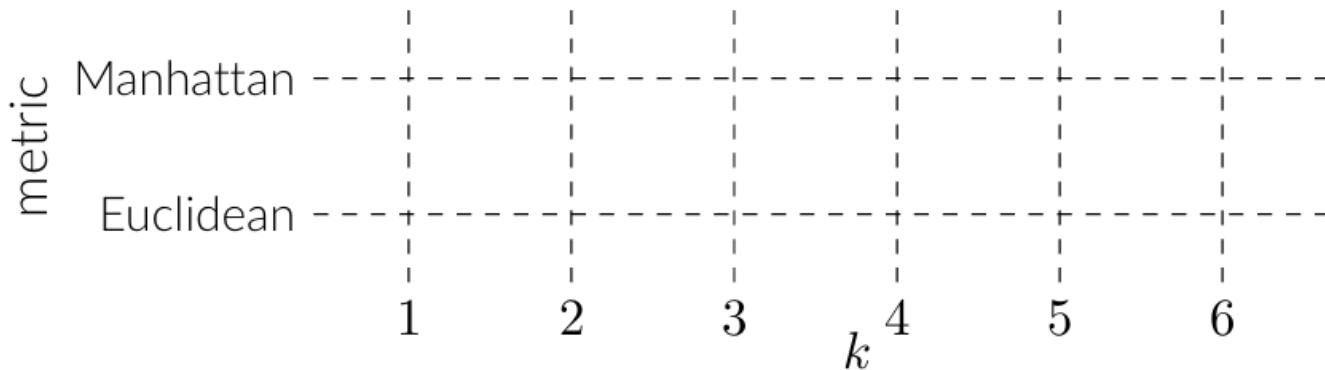
- Model selection
- which input features to include
  - winter rain, summer temp, harvest rain, September. temp, age
- Hyperparameter tuning
- for k-nearest neighbors, hyperparameters include:
  - k
  - metric (e.g. Euclidean distance)



# Hyperparameter tuning

## > Grid search

- We need to try all 12 combinations on the following grid:



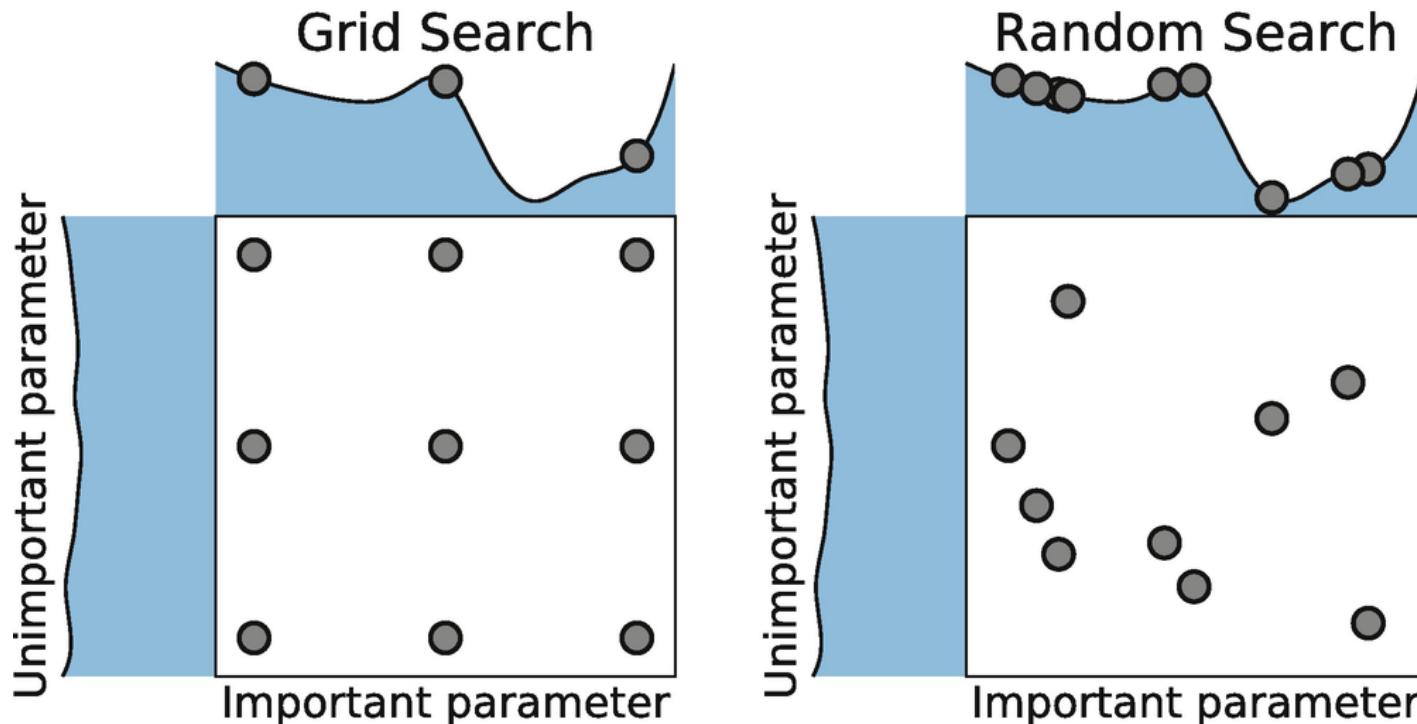
- 2 possible scaler (standard scaler, minmax scaler)
- there were 5 input features in the original data
- $2^5 = 32$  combinations of features we need to try
  
- In total,  $32 * 12 * 2 = 768$  models

# Hyperparameter tuning

- > For large datasets, it is impossible to try every combination of models and parameters
- > So instead we use heuristics, which do not guarantee the best model but tend to work well in practice
  - Randomized search: try random combinations of parameters
  - Coordinate optimization
    - start with guesses for all parameters
    - try all values for one parameter (holding the rest constant) and find the best value of that parameter
    - cycle through the parameters
  - ...

## Random search

- > Explores by choosing a new position at random after each iteration
  - purely random across the search space in each step

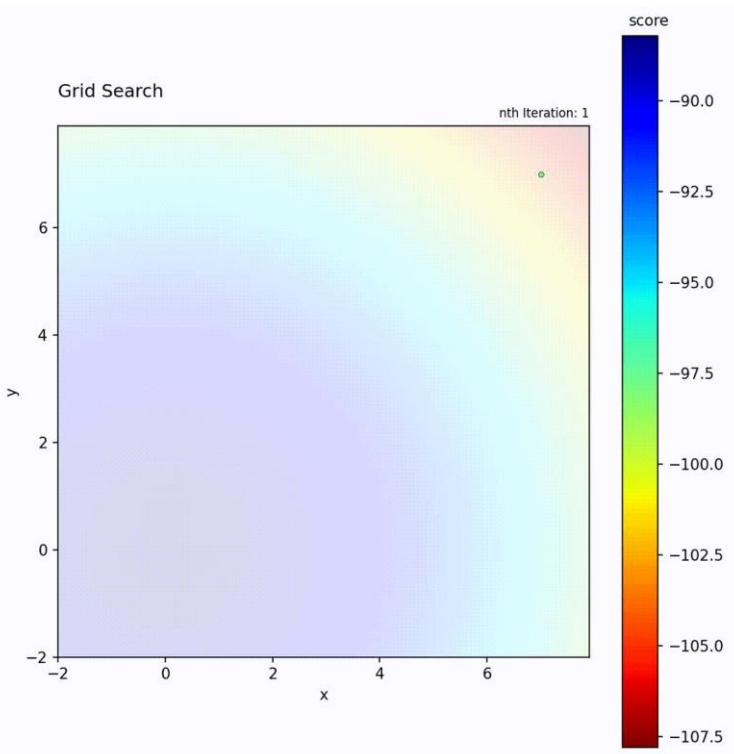
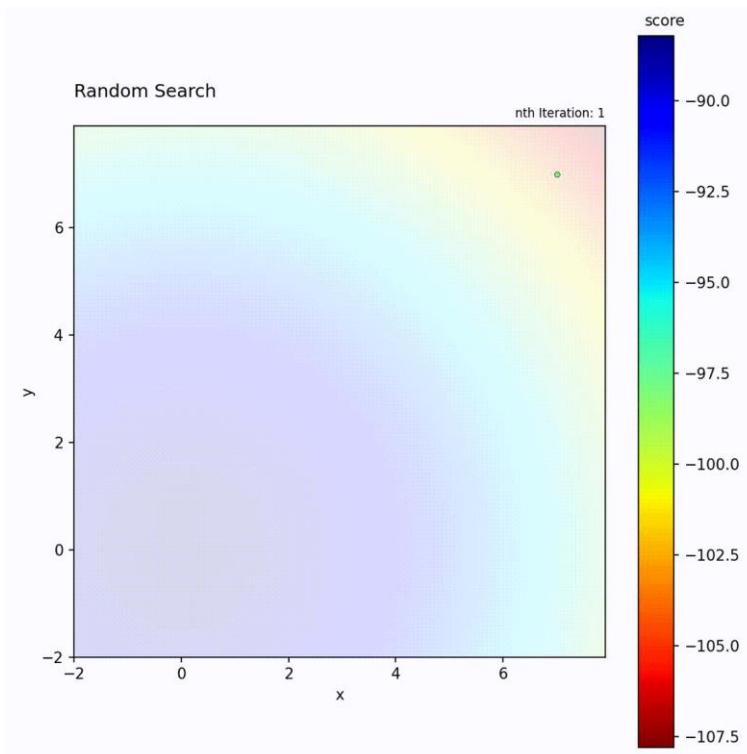


## Random search

- > Appropriate for solving deterministic objectives
  - often have promising theoretical properties
    - convergence in probability
    - hitting probability: to land in an  $\epsilon$ -optimal region with probability at least  $p$
  - Pros
    - extremely easy to implement
    - no smoothness/continuity assumed
    - perfect parallelizable (i.i.d. samples)
    - all points can be generated in advance
  - Cons
    - convergence rate is slow (compared to methods that exploit continuity/gradients)
    - requires large  $N$  to achieve good probability guarantees

## Random search

- > RS samples points independently in each iteration, while grid search systematically explores the search space along predefined directions



## Random search

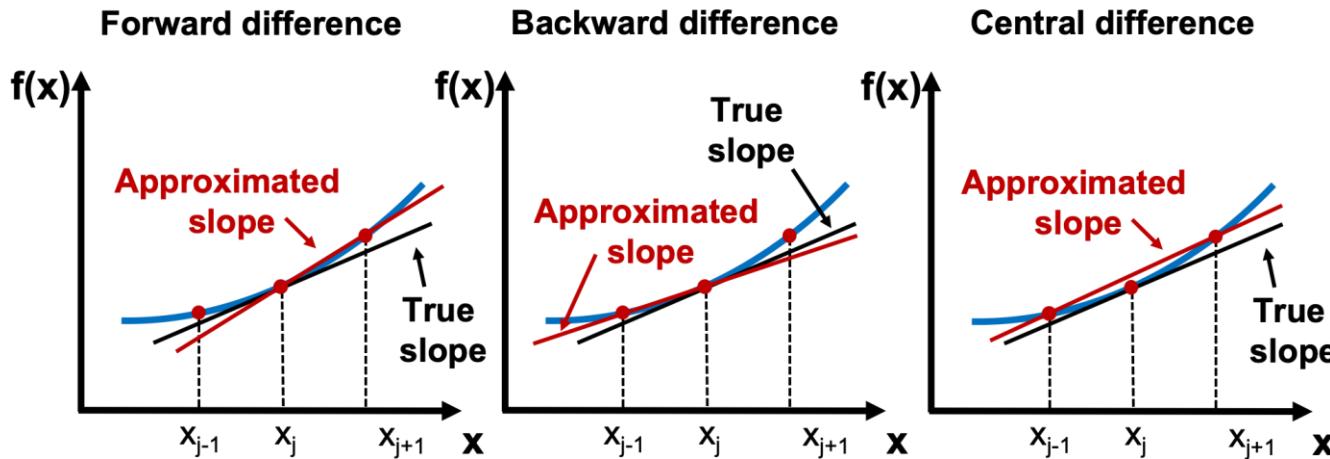
- > How to advance the random search?
  - RS explores blindly
    - ignores state space structure and history
  - Once we collect information, we can guide the search
  - Gradient-free
    - when true gradient is unavailable (or noisy)
    - approximate with sampled perturbations
  - Derivative-free
    - do not compute or approximate gradient at all
    - rely only on function values to guide sampling  
(search more in promising regions)

# Gradient-free optimization

## > Finite difference – coordinate perturbations

- estimate partial derivatives by probing each axis

- gradient estimate:  $f'_j = \frac{f(x+he_j) - f(x)}{h}$  or  $f'_j = \frac{f(x+he_j) - f(x-he_j)}{2h}$



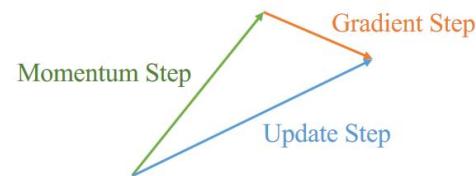
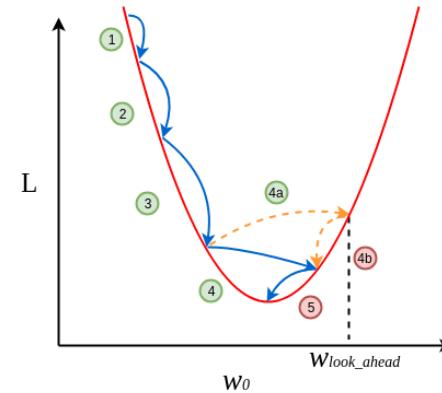
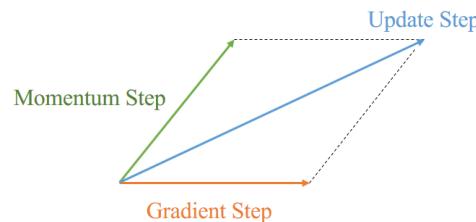
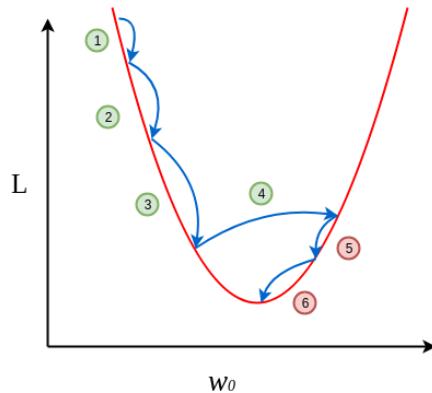
- simple and transparent
- scales poorly with dimension; sensitive to noise/non-differentiability

# Gradient-free optimization

## > Evolution strategy

- sample around the current mean with Gaussian noise
- use reward-weighted averaging to estimate a gradient
- more sample efficient than FD in high dimension
- Nesterov's random search
  - use momentum from past updates to accelerate

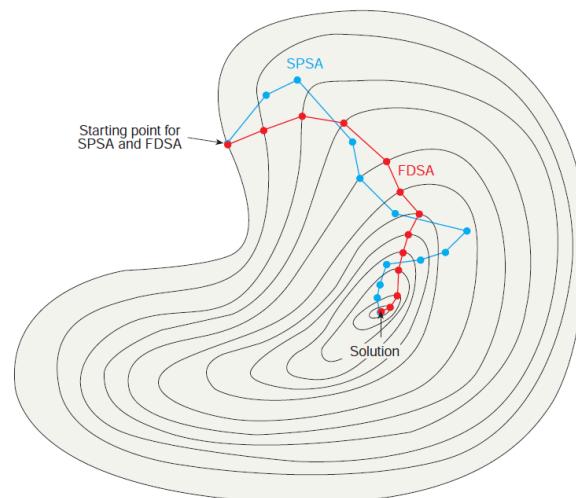
recall  
Nesterov's  
momentum



# Gradient-free optimization

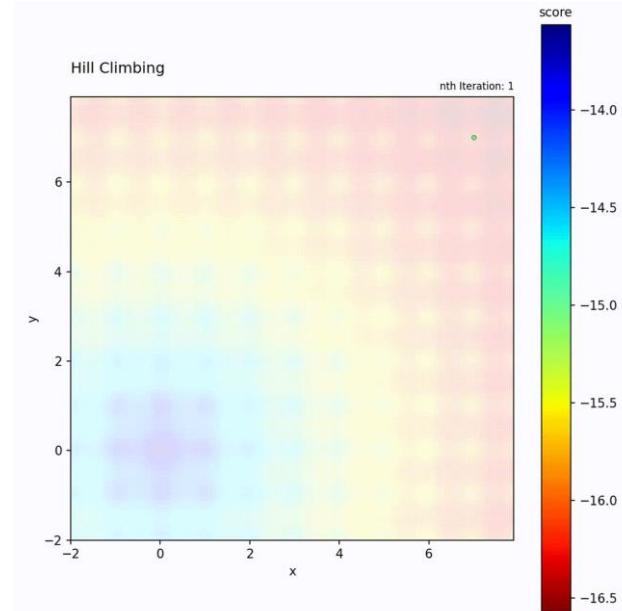
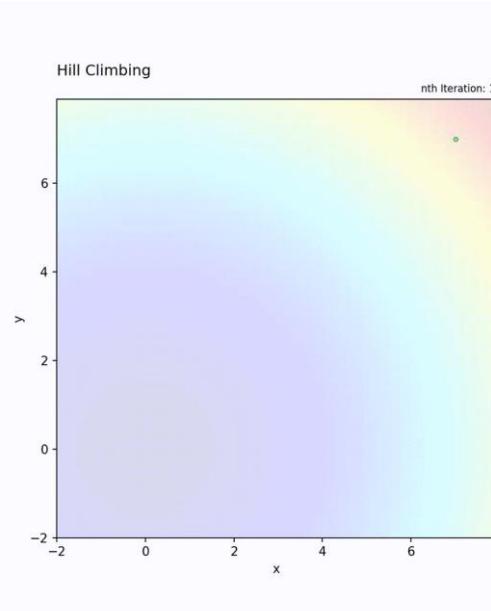
## > Simultaneous Perturbation Stochastic Approximation (SPSA)

- estimate a full gradient with only two function evaluations
- generate a random perturbation vector  $\Delta \in \{\pm 1\}^d$
- Gradient estimate:  $\hat{g}_j = \frac{f(x+c\Delta)-f(x-c\Delta)}{2c\Delta_j}$
- extreme case of gradient approximation
  - ultra sample-efficient per iteration; great for very high-dim., noisy problems
  - higher variance than ES;



# Derivative-free local search

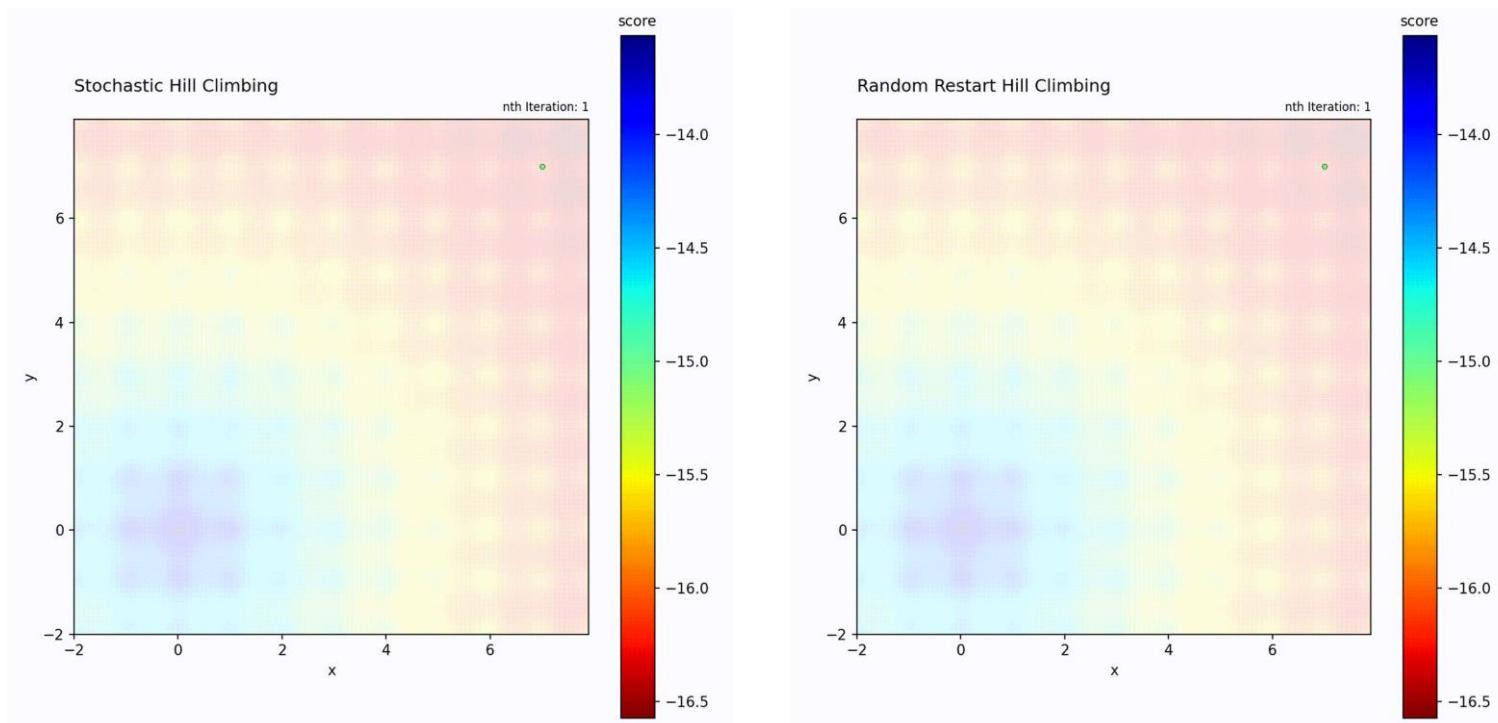
- > Derivative-free
  - methods that do not explicitly estimate a gradient, but still improve solutions via local moves
- > Hill climbing
  - greedy local improvement: probe neighbors, move if better
  - evaluate the score of n neighbors in an epsilon environment and moves to the best one



# Derivative-free local search

## > Hill climbing variations

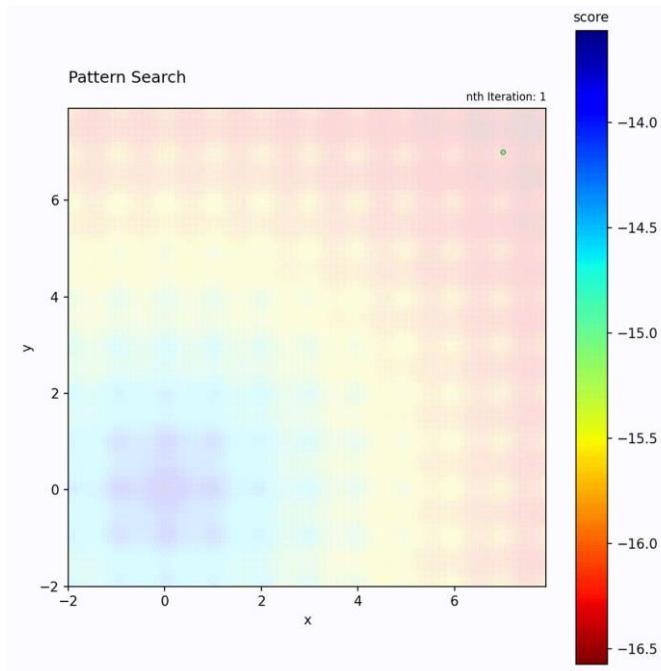
- stochastic: adds a probability to move to a worse position
- random restart: moves to a random position after n iterations.



# Derivative-free local search

## > Pattern search

- probe a set of directions at a given step size
- if any direction yields improvement, move and possibly enlarge steps otherwise shrink step and try again
- robust to noise and non-smoothness
- can be inefficient in ill-conditioned landscapes

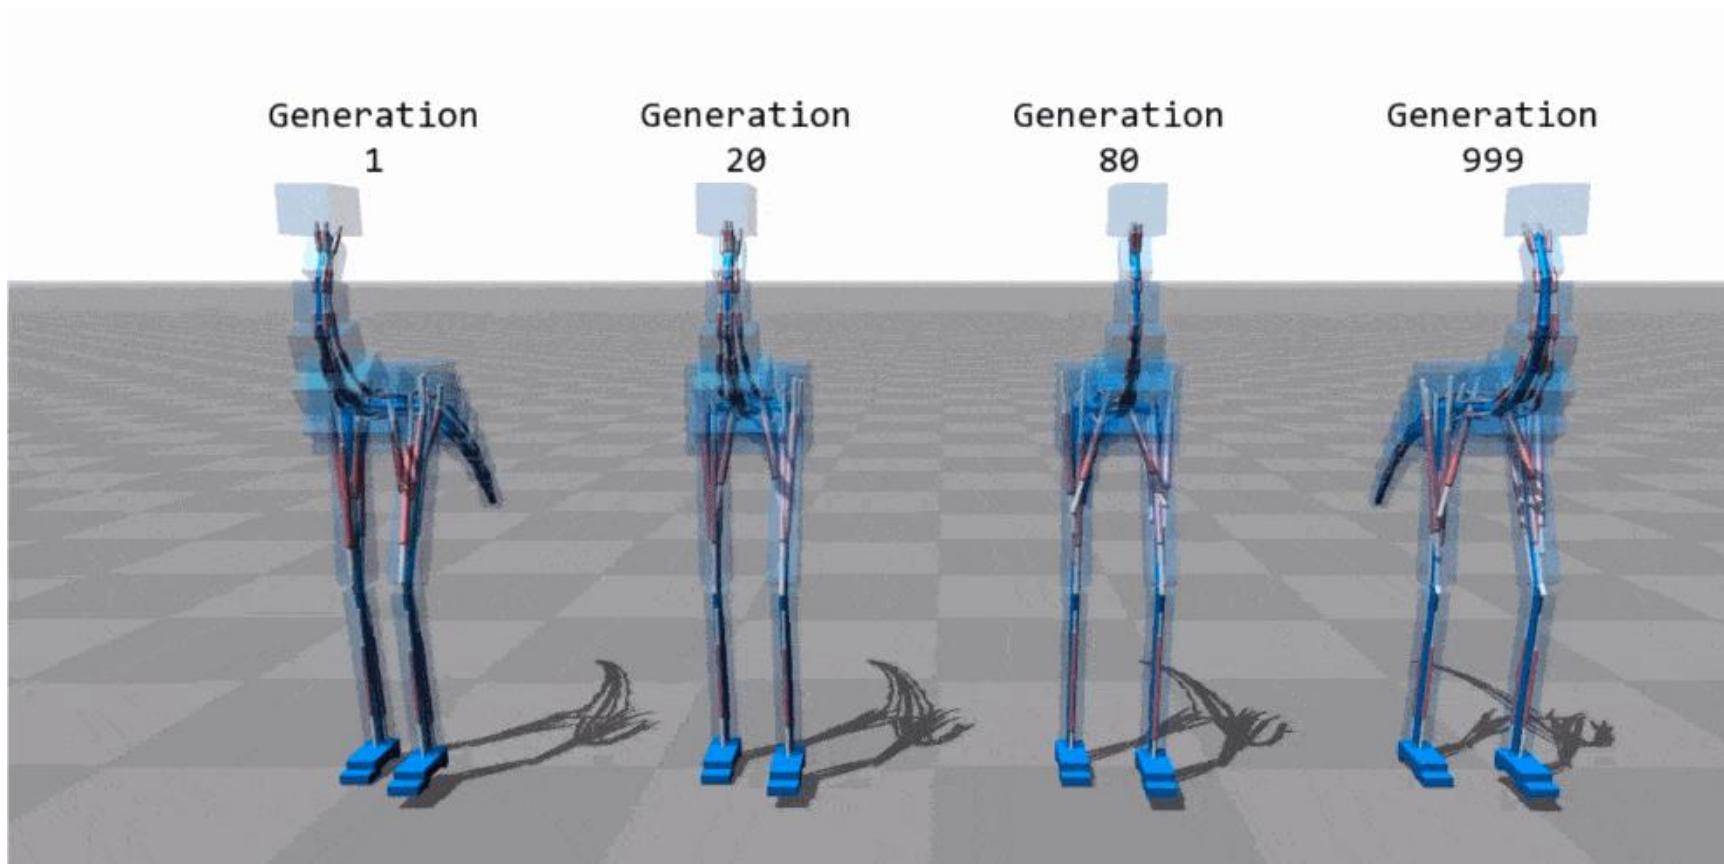


# Population-based metaheuristics

- > Local search
  - explores only the immediate neighborhood of the current solution
  - very efficient for smooth, low-dimensional problems
  - easily trapped in local optima
  - stochastic / random restart methods are partially global but still limited
- > Global search
  - need methods that systematically explore the entire search space
  - population-based and distribution-based approaches

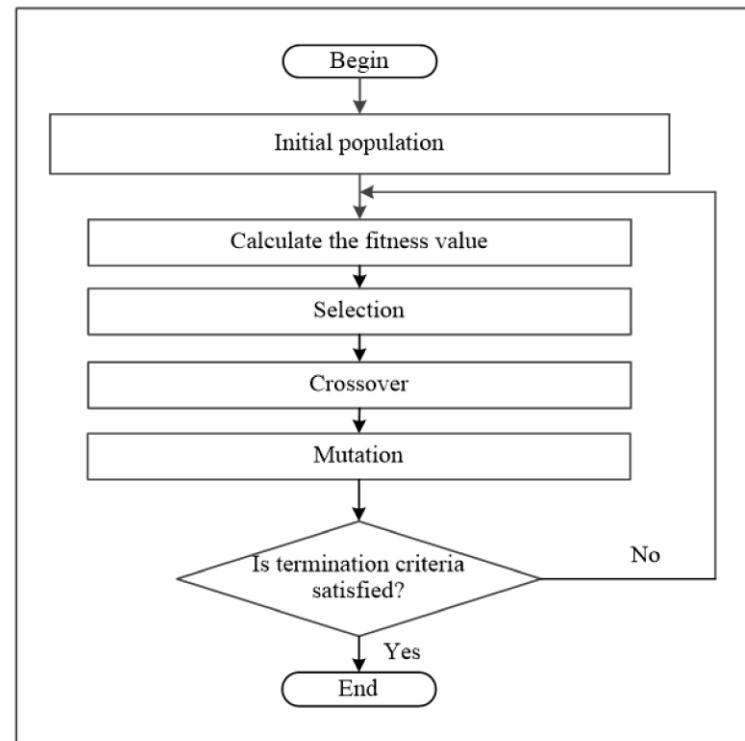
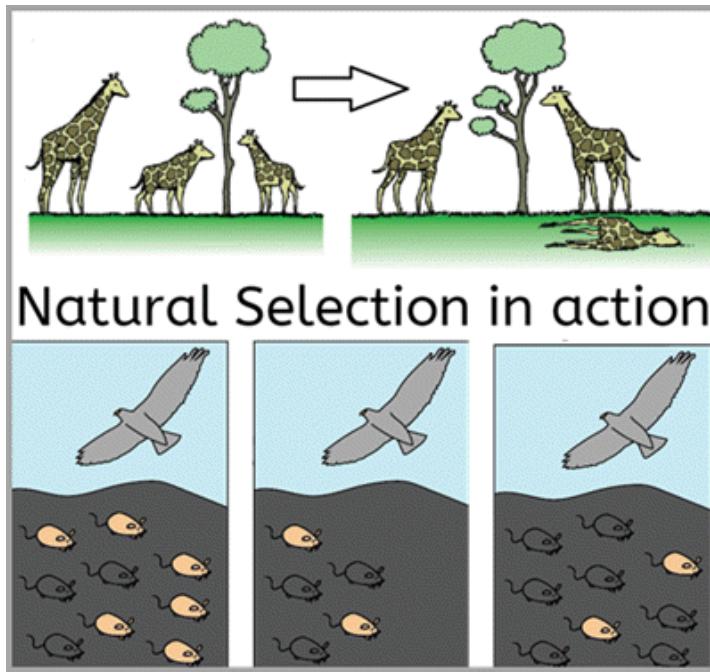
# Population-based metaheuristics

- > Genetic algorithm



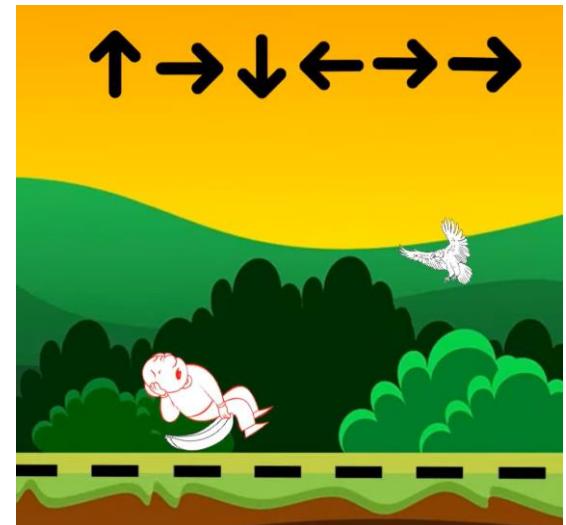
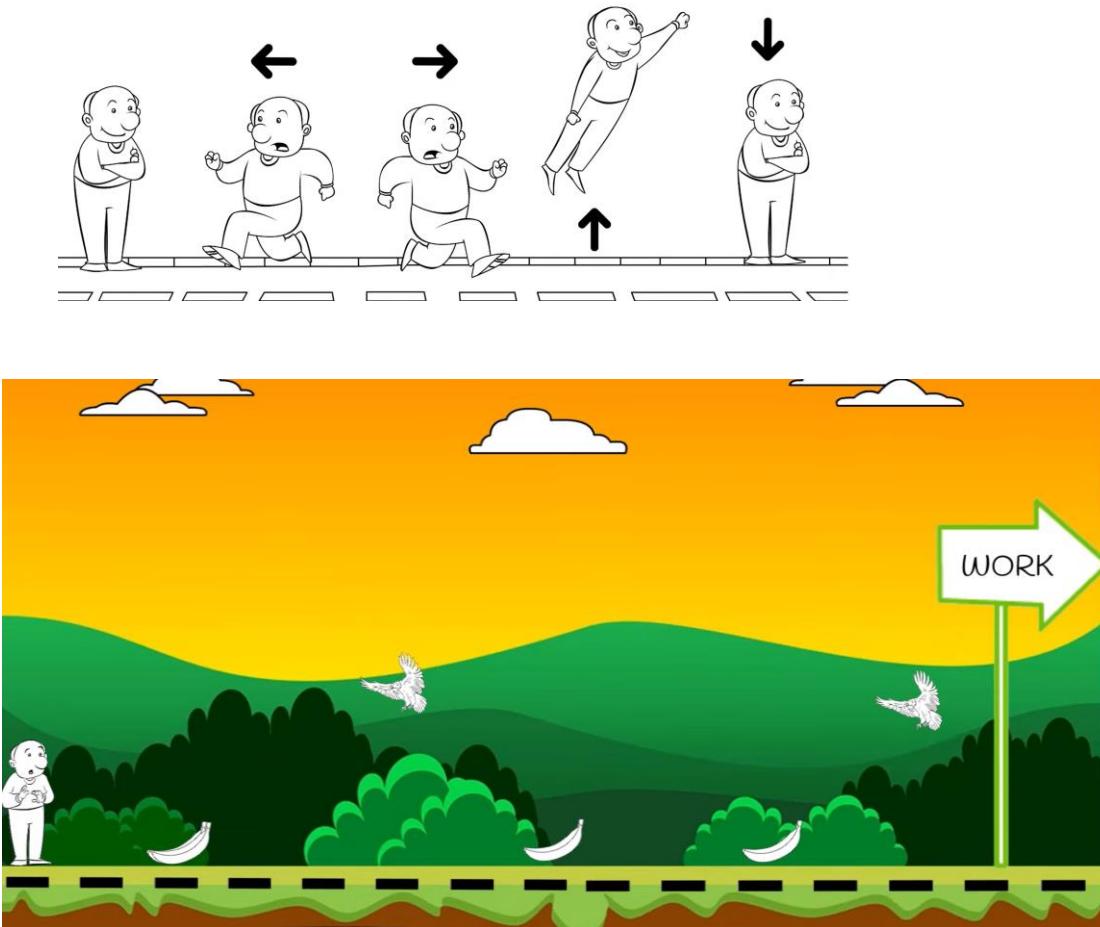
# Population-based metaheuristics

- > GA is a heuristic search that mimics the natural evolution
  - it is a particular class of evolutionary algorithms inspired by evolutionary biology such as mutation, selection, and crossover
  - select the best, discard the rest



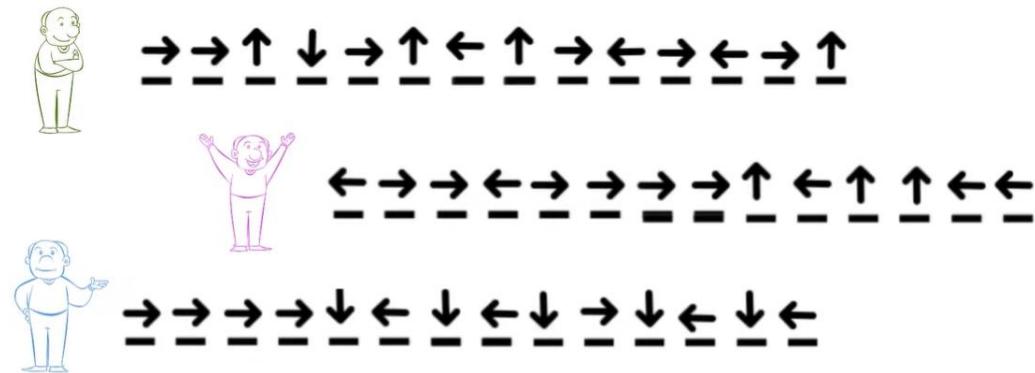
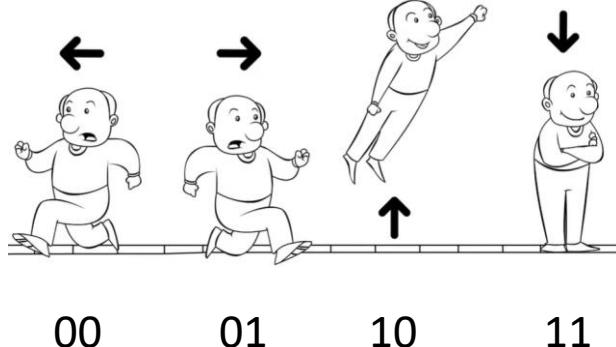
# Population-based metaheuristics

- > Find a path to work safely

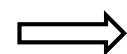


# Population-based metaheuristics

## > GA – binary encoding

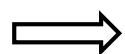


010110110110**00**10...  
000101**00**01010101...  
01**0**1010111001100...  
...



mutation / crossover / ...

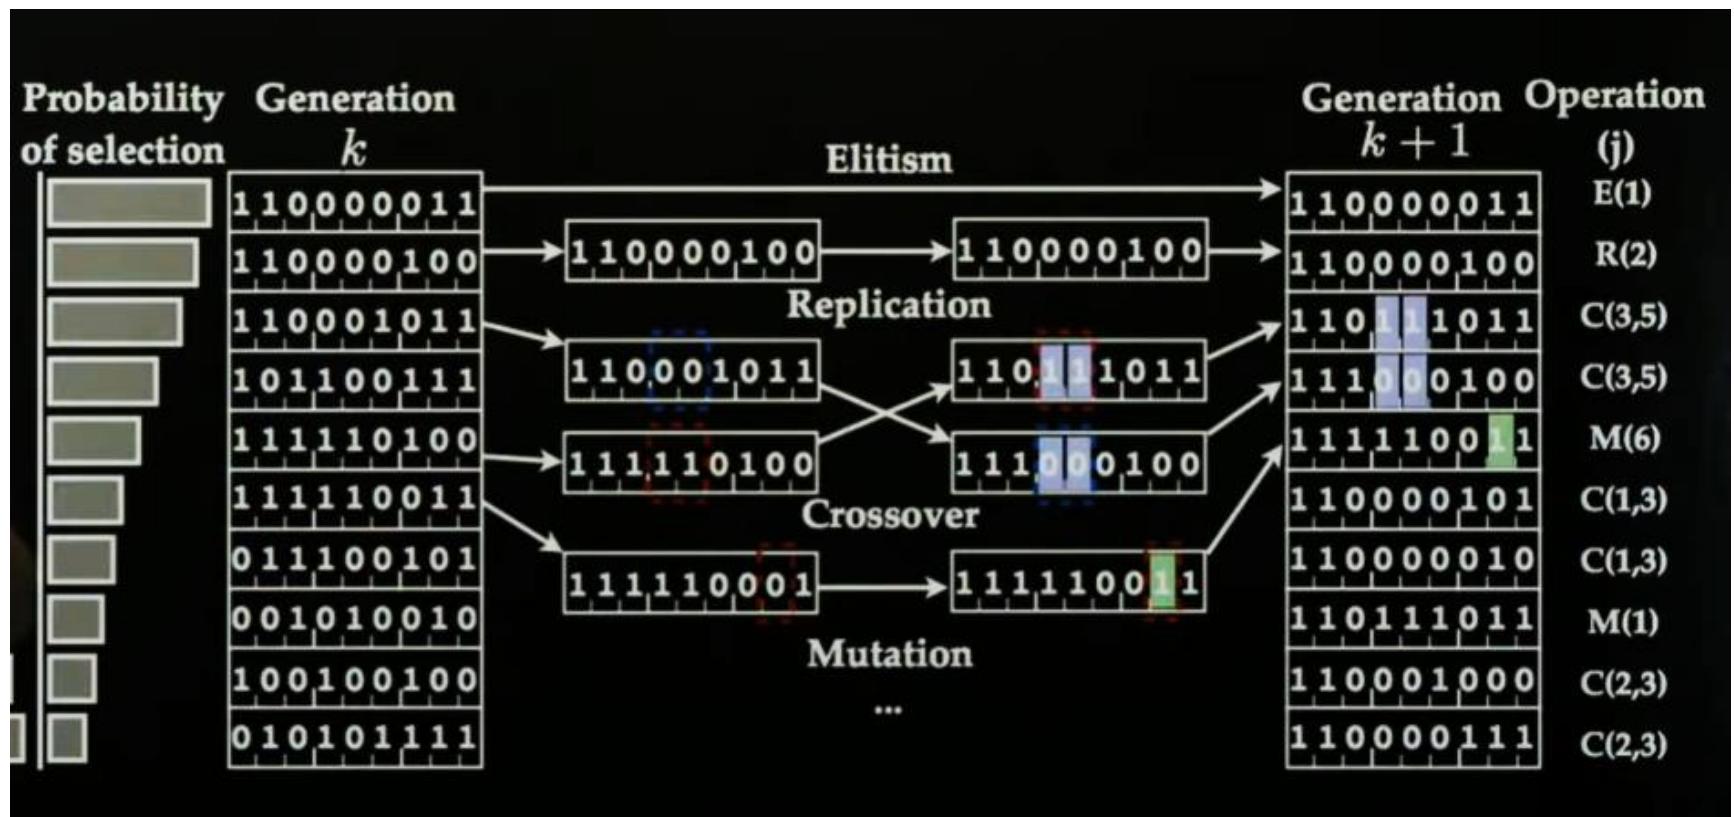
010110110110**10**10...  
000101**11**01010101...  
01**11**01011100**01**00...  
...



calculate the fitness value (evaluate)  
selection

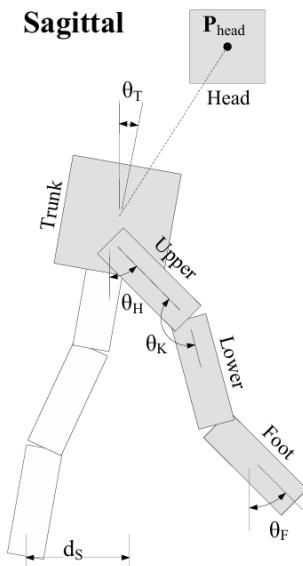
# Population-based metaheuristics

- > GA – binary encoding



# Population-based metaheuristics

- > GA – real-value encoding
  - what if we want to represent the continuous value



1) binning  $\theta$  to 255 levels

01111111 (127/255)  $\rightarrow$  10000000 (128/255)

2) gene encoded as real values = [1.2, 3.5 ...]

crossover: child =  $\alpha x^{(1)} + (1 - \alpha)x^{(2)}$

mutation: child =  $x + \mathcal{N}(0, \sigma^2)$

# Population-based metaheuristics

## > Pros and cons

- GA search a population of points in parallel, not only a single point
  - can be easily parallelized
- It works well on mixed discrete/continuous problems
- Simple to understand and set up
  
- convergence behavior is very dependent on tuning parameters
- cumbersome to take into account constraints
- no clear termination criteria

# Population-based metaheuristics

- > Particle Swarm Optimization (PSO)
  - Inspired by social behavior of bird flocking
  - Suppose a group of birds is searching food in an area
    - birds can remember the best place it has found so far
    - birds can see the best location found by other birds

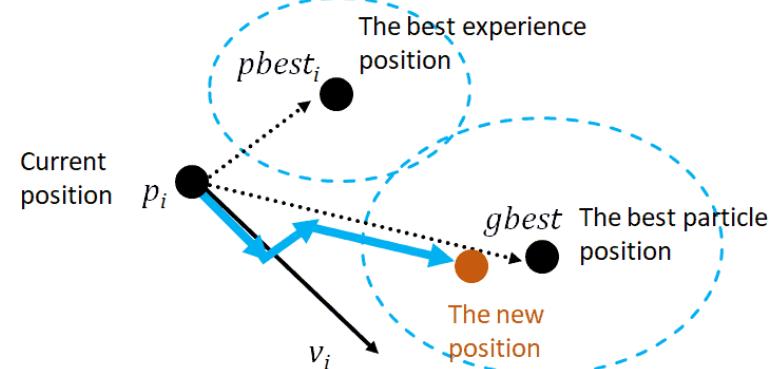


# Population-based metaheuristics

## > Particle Swarm Optimization (PSO)

### > Algorithm

- position  $x_i$
- velocity  $v_i$
- personal best position  $p_i$
- global best position  $g$
- velocity and position update rules
- $x_i(t+1) = x_i(t) + v_i(t+1)$
- $v_i(t+1) = \underline{wv_i(t)} + \underline{c_1r_1(p_i - x_i(t))} + \underline{c_2r_2(g - x_i(t))}$



Inertia – keeps the particle moving in the same direction

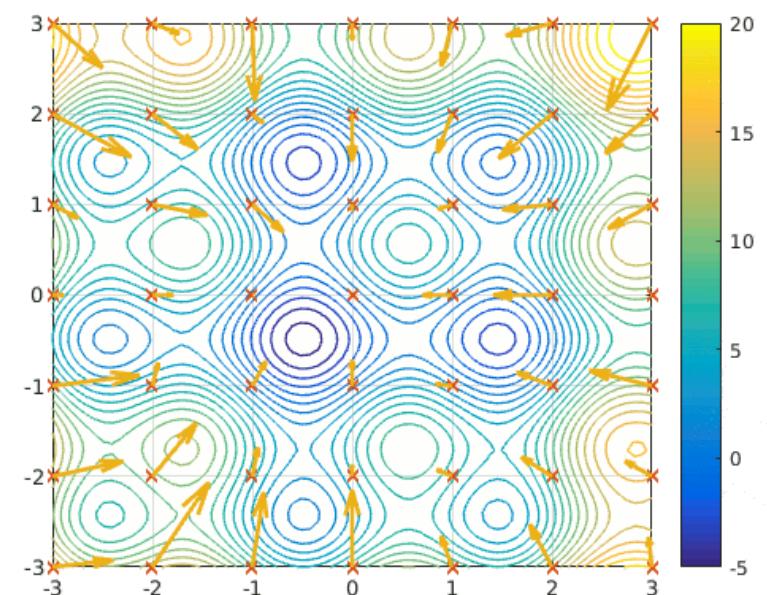
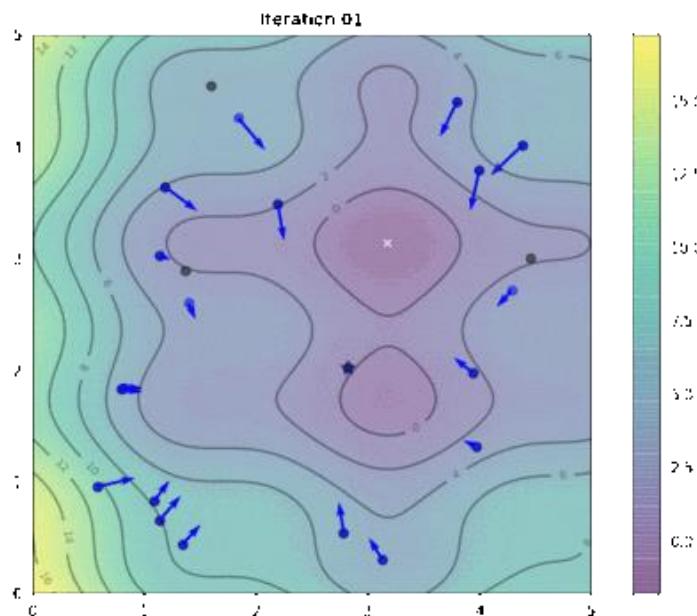
Cognitive – pulls the particle back toward its own best-known position

Social – pulls the particle toward the swarm's best-known position

- $r_1, r_2$  : random numbers introducing stochasticity

# Population-based metaheuristics

## > Particle Swarm Optimization (PSO)



# Population-based metaheuristics

- > Pros and cons
  - simple and easy to implement
  - works well in high-dimensional, nonlinear, or noisy problems
  - may converge prematurely to a local minimum
  - sensitive to parameter tuning
  - no theoretical convergence guarantees
- > Applications
  - control system parameter tuning

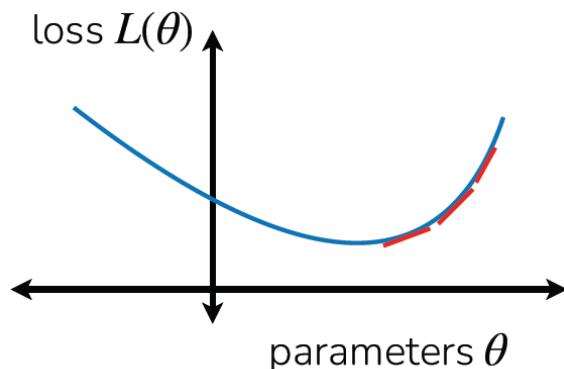
## Population-based metaheuristics

- > Nature-inspired metaheuristic optimization
- > Collective behavior of animals (Swarm intelligence)
  - Ant Colony Optimization (ACO)
    - discrete domain (e.g. route optimization, TSP)
  - Bacterial Foraging Optimization (BFO)
  - Artificial Bee Colony (ABC)
  - Firefly algorithm (FA)
  - ...

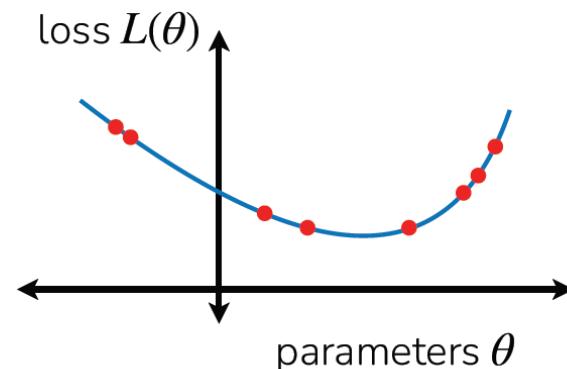
# Stochastic sampling-based optimization

## > Sampling-based optimization

Gradient-based (1st order)



Sampling-based (0th order)



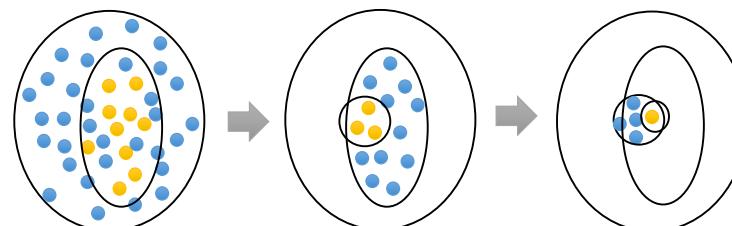
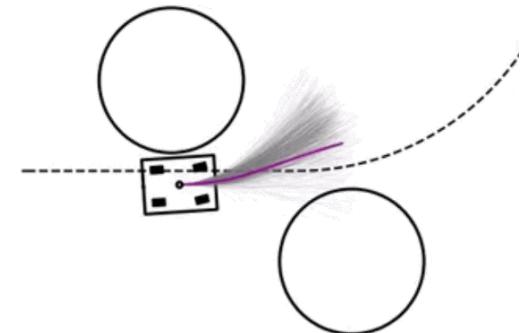
- requires no gradient information (gradient-free optimization)
- use randomness to explore the space
- often better at escaping local minima than gradient-based methods
- parallelizable
- scales poorly to high dimensions

# Stochastic sampling-based optimization

## > Sampling-based optimization

- Random shooting
  - Guess and check
  - sample many action sequences ( $A = \{a_t, \dots, a_{t+H}\}$ )  $A_1, \dots, A_n$  from some distribution (e.g., uniform)
  - choose the best  $A_i$  (max return)
  - can we improve the sampling distribution? (prior knowledge)
- Cross-entropy method (CEM)
  - - sample many action sequences  $A_1, \dots, A_n$  from  $p(A)$
  - evaluate and pick elites  $A_{i_1}, \dots, A_{i_m}$
  - refit  $p(A)$  to the elites  $A_{i_1}, \dots, A_{i_m}$

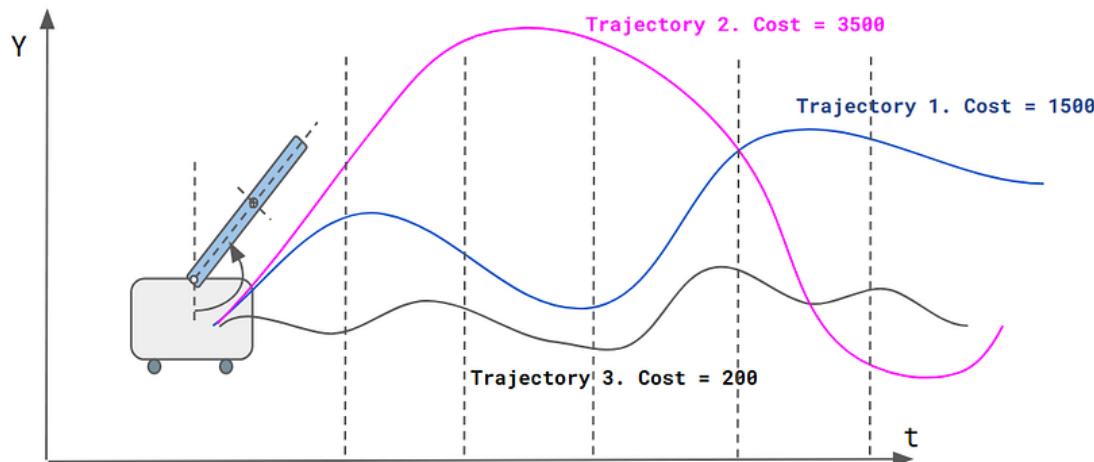
until convergence



# Stochastic sampling-based optimization

## > Sampling-based optimization

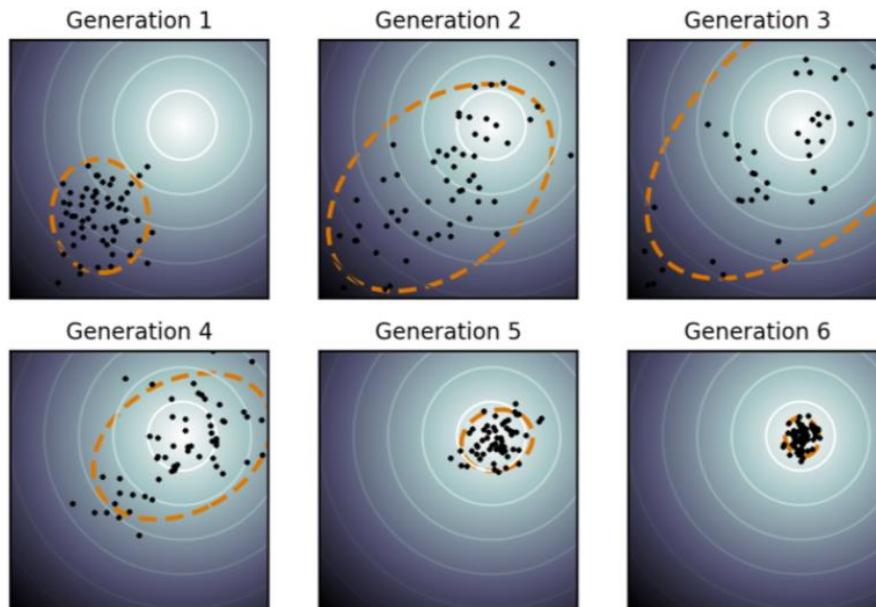
- model predictive path integral control (MPPI)
- - sample many action sequences  $A_1, \dots, A_n$  from  $p(A) \sim \mathcal{N}(\mu, \Sigma)$
- evaluate corresponding costs  $C_1, \dots, C_n$
- compute weights based on costs:  $w_i = \frac{\exp(-\frac{1}{\lambda}C_i)}{\sum_j \exp(-\frac{1}{\lambda}C_j)}$
- update  $\mu \leftarrow \sum_i w_i A_i$



# Stochastic sampling-based optimization

## > Sampling-based optimization

- covariance matrix adaptation evolutionary strategy (CMA-ES)
  - - sample many action sequences  $A_1, \dots, A_n$  from  $p(A) \sim \mathcal{N}(\mu, \Sigma)$
  - evaluate corresponding costs  $C_1, \dots, C_n$
  - compute weights based on costs:  $w_i$  (e.g., exponential, top-K, ...)
  - update  $\mu \leftarrow (1 - \gamma_\mu)\mu + \gamma_\mu \sum_i w_i A_i$ ,  $\Sigma \leftarrow (1 - \gamma_\Sigma)\Sigma + \gamma_\Sigma \sum_i (A_i - \mu)(A_i - \mu)^\top$



## Reference

- > <https://web.stanford.edu/class/datasci112/lectures/lecture13.pdf>
- > <https://simonblanke.github.io/gradient-free-optimizers-documentation/1.5/>
- > [https://www.deisenroth.cc/teaching/2020-21/ml-seminar/lecture\\_bayesian\\_optimization.pdf](https://www.deisenroth.cc/teaching/2020-21/ml-seminar/lecture_bayesian_optimization.pdf)