

ECE7121 Learning-based control – 2025 Fall

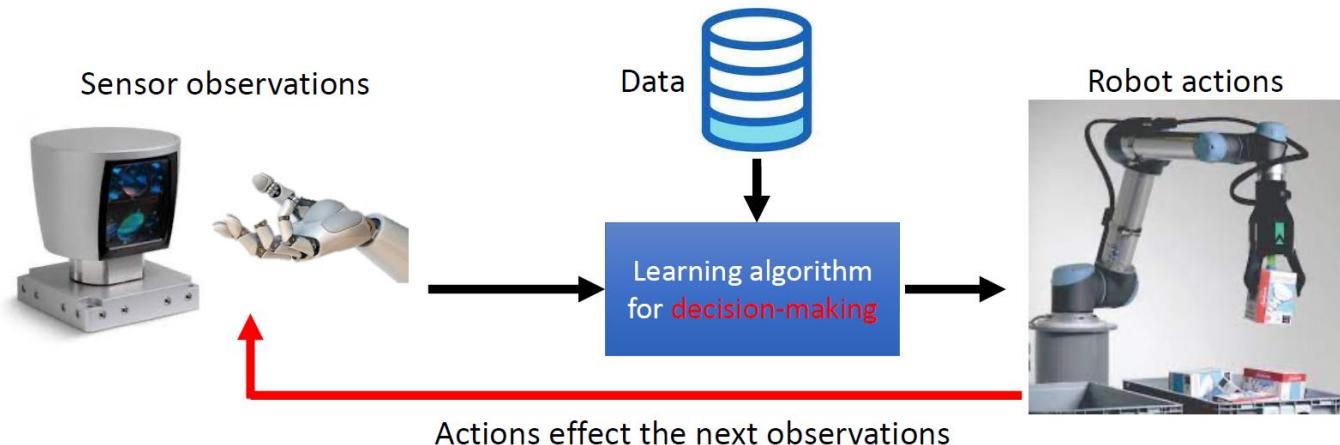
Reinforcement learning basics



INHA UNIVERSITY

Robot learning

- > Learning to make sequential decisions in the physical world
 - A system need to make multiple decisions based on stream of information
- > The solutions to such problems
 - imitation learning - offline & online RL
 - model-free & model-based RL - multi-task & meta RL



RL Success

> Game



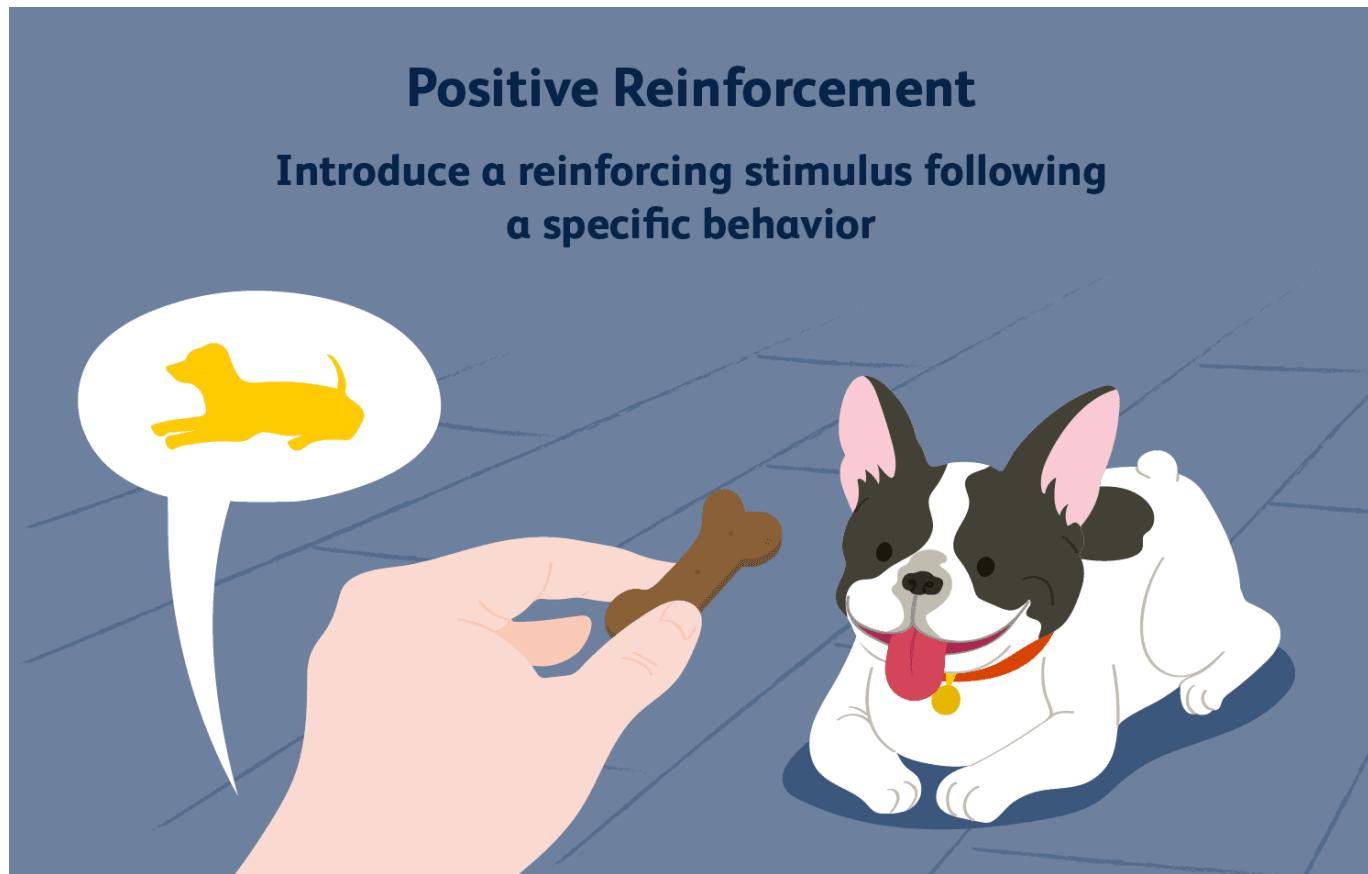
AlphaGo: Go World champion



AlphaStar: Grandmaster (99.8%)

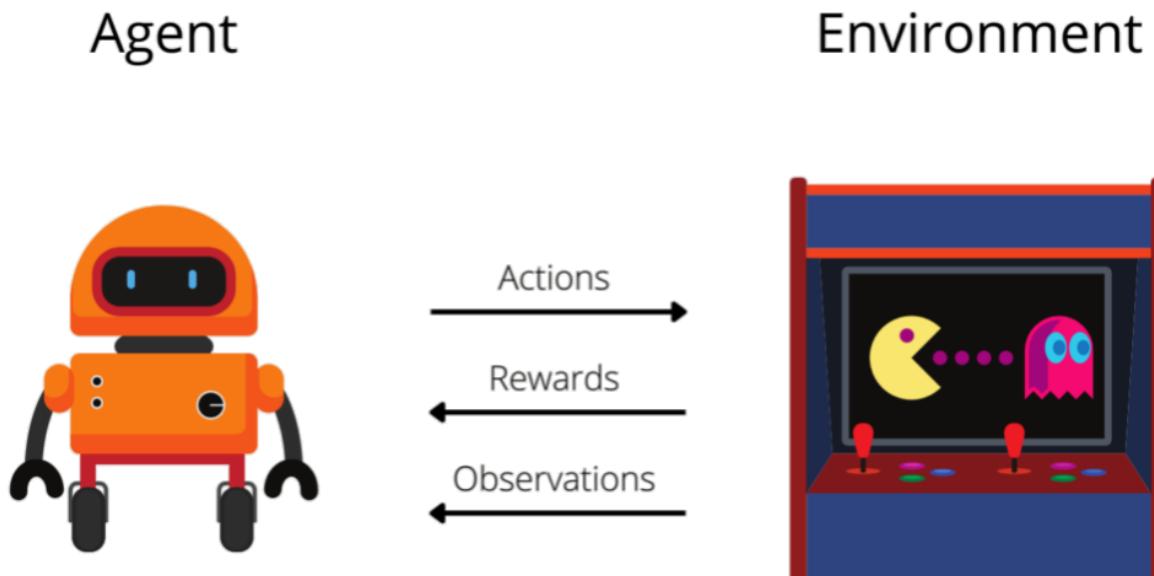
RL = trial-and-error learning

- > Reinforcement in educational psychology
 - Big advantage: AI can learn autonomously



RL = trial-and-error learning

- > Learning a policy that maximizes rewards by interacting with the environment



- Each action results in an immediate reward.
- We want to choose actions that maximize our immediate reward in expectation.

Multi-armed bandits

> bandit =



One-armed bandit: slot machine



Multi-armed bandit: multiple slot machines

Multi-armed bandits

- > The state does not change!
 - We don't move. We use the same slot machines.
- > We have N slot machines and select one to pull.
 - We have N possible actions.
- > We will get an immediate reward from the pulled slot machine.
 - Each slot machine gives a random reward.
 - The reward probability of each machine is fixed but unknown.
- > Objective: maximize cumulative rewards

How to earn money?

- > Strategy 1: pull each once, exploit the best



Reward: 100₩ 0₩ 200₩ 30₩

- > Next, we only pull the third machine

How to earn money?

- > Strategy 2: pull each 4 times, exploit the best



Reward:

100₩	0₩	200₩	30₩
100₩	300₩	0₩	500₩
100₩	0₩	0₩	20₩
100₩	400₩	0₩	40₩
400₩	700₩	200₩	590₩

- > Next, we only pull the second machine

Achieving a balance

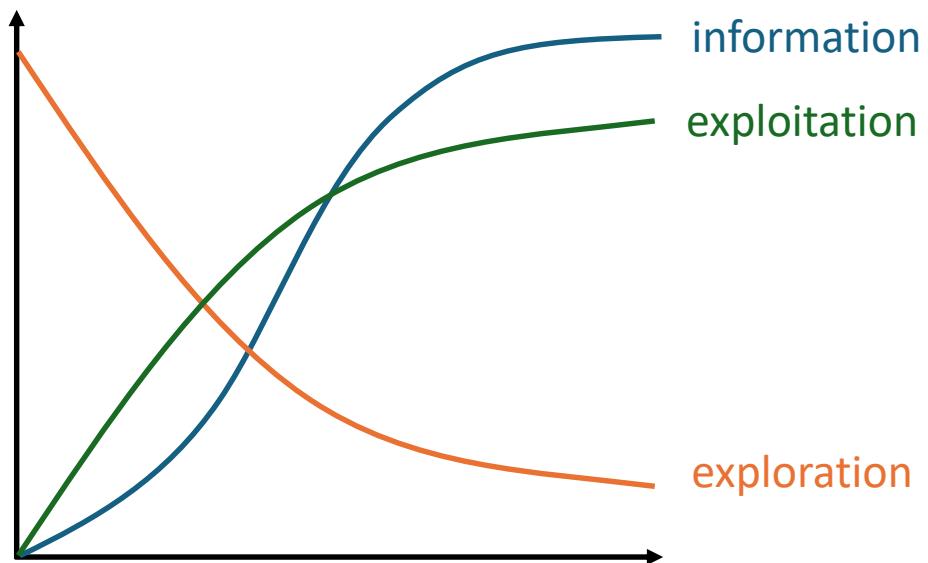
- > Pulling the same machine several times
 - = learning reward probability distribution and its mean
 - = **exploration (collecting information)**
- > Pulling the best machine
 - = exploiting the best to earn money (given current information)
 - = **exploitation (collecting reward)**
- > Action-value for action a is its mean reward
 - $Q_t^*(a) = \mathbb{E}[R_{t+1} | A_t = a]$, action-value estimate: $Q_t(a) \approx Q_t^*(a)$
 - $A_t^* = \arg \max Q_t(a)$
 - If $A_t = A_t^*$: exploiting
 - If $A_t \neq A_t^*$: exploring
- > We need to do both

Exploration dilemma

- > Exploration vs Exploitation dilemma
 - The best long-term strategy may involve short-term sacrifices
 - This is not a problem unique to RL; it is a fundamental issue in the decision making of any intelligent agent.
- > Restaurant selection
 - exploitation: go to your favorite restaurant
 - exploration: try a new restaurant
- > Studying
 - exploitation: solve example problems
 - exploration: read additional materials

Exploration dilemma

- > $\epsilon - \text{greedy}$ algorithm

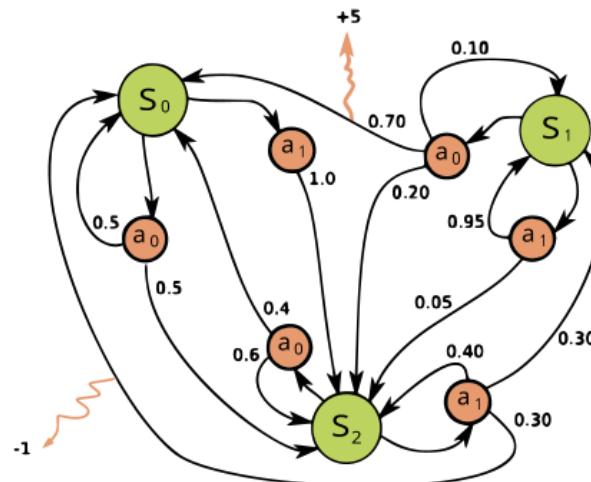
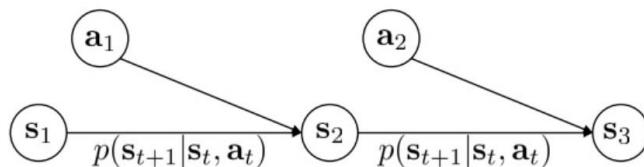


Markov decision process

- > Multi-armed bandit
 - $\tau: (A_t, R_{t+1}, A_{t+1}, R_{t+2}, A_{t+2}, R_{t+3}, \dots)$
- > Markov decision process
 - $\tau: (S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}, R_{t+2}, S_{t+2}, A_{t+2}, R_{t+3}, \dots)$
- > Markov property
 - $p[R_{t+1} = r, S_{t+1} = s' | S_0, A_0, R_0, S_1, A_1, R_1, \dots, S_t, A_t]$
 $= p[R_{t+1} = r, S_{t+1} = s' | S_t, A_t]$
 - Only the present determines the future; we can ignore the history.

Markov decision process

- > Finite Markov decision process is a tuple (S, A, T, r, γ)
 - S is a finite set of states $s \in S$
 - A is a finite set of actions $a \in A$
 - T is one step transition/dynamics function $p(s'|s, a)$
 - $r(s, a, s')$ is a reward function
 - γ is a discount factor ($0 \leq \gamma \leq 1$)
- > $\pi(a|s)$: a policy is a distribution over actions given states

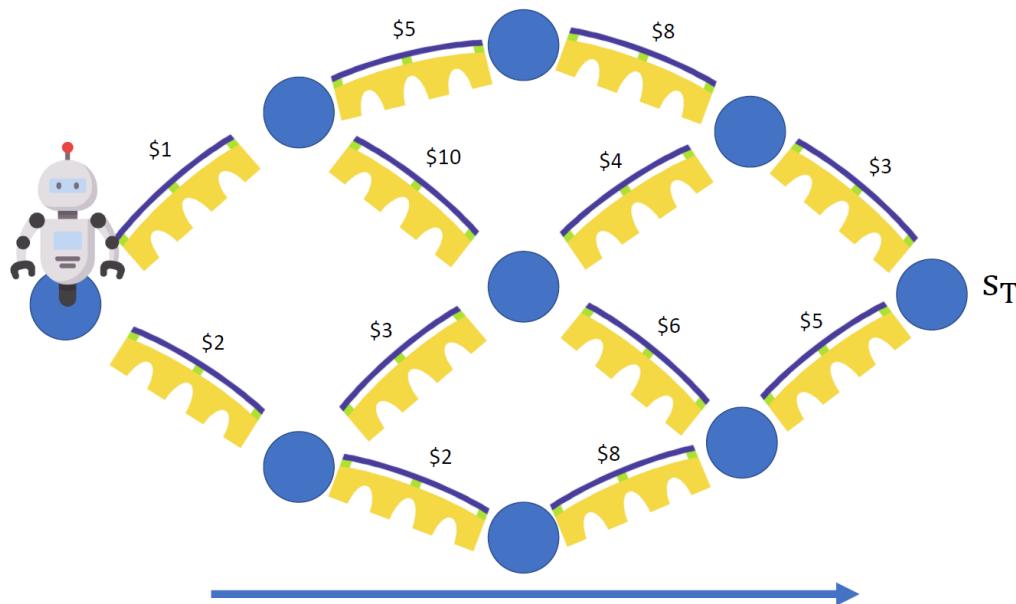


Markov decision process

> Maximize your sum of future rewards (cumulative reward)

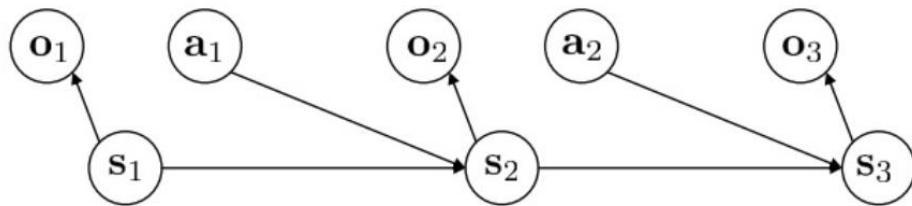
- $G = R(\tau) = R_1 + R_2 + R_3 + R_4 \dots$
- Future rewards may less important
- $G = R(\tau) = R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots$

> The robot collects toll on every bridge



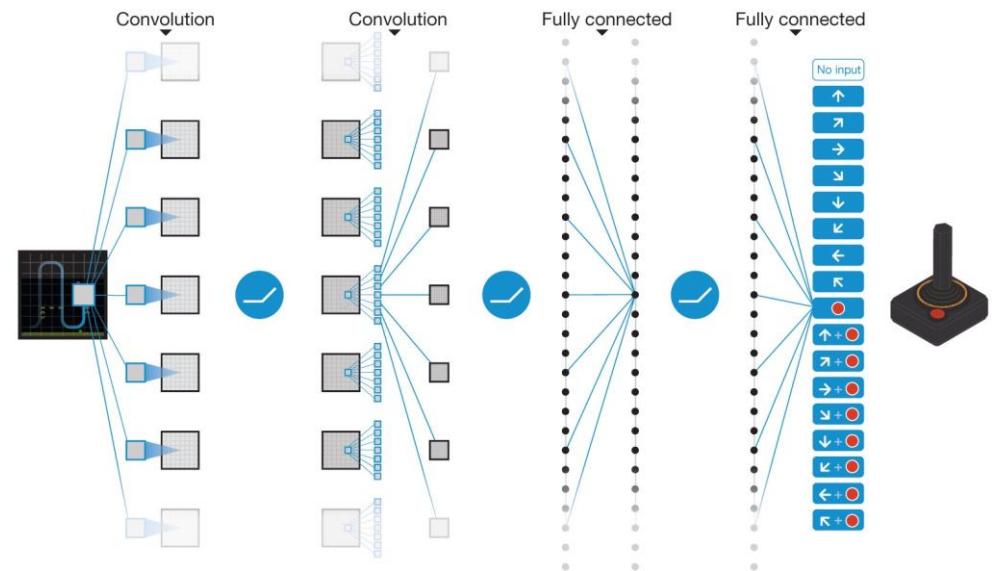
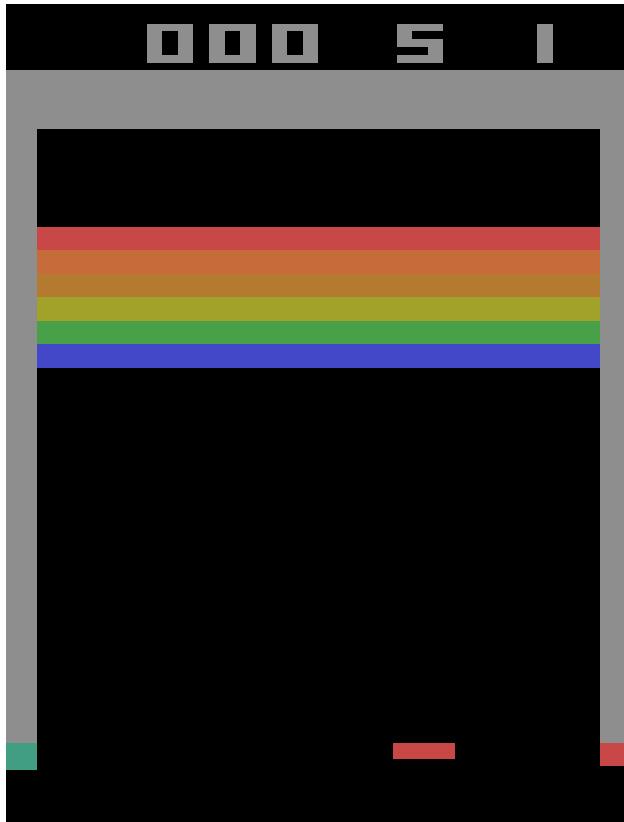
POMDP

- > Partially Observable MDP
 - Finite POMDP is a tuple $(S, A, O, T, h, r, \gamma)$
 - h is the observation model $p(o|s)$
- > Learn a policy $\pi(a|o)$



POMDP

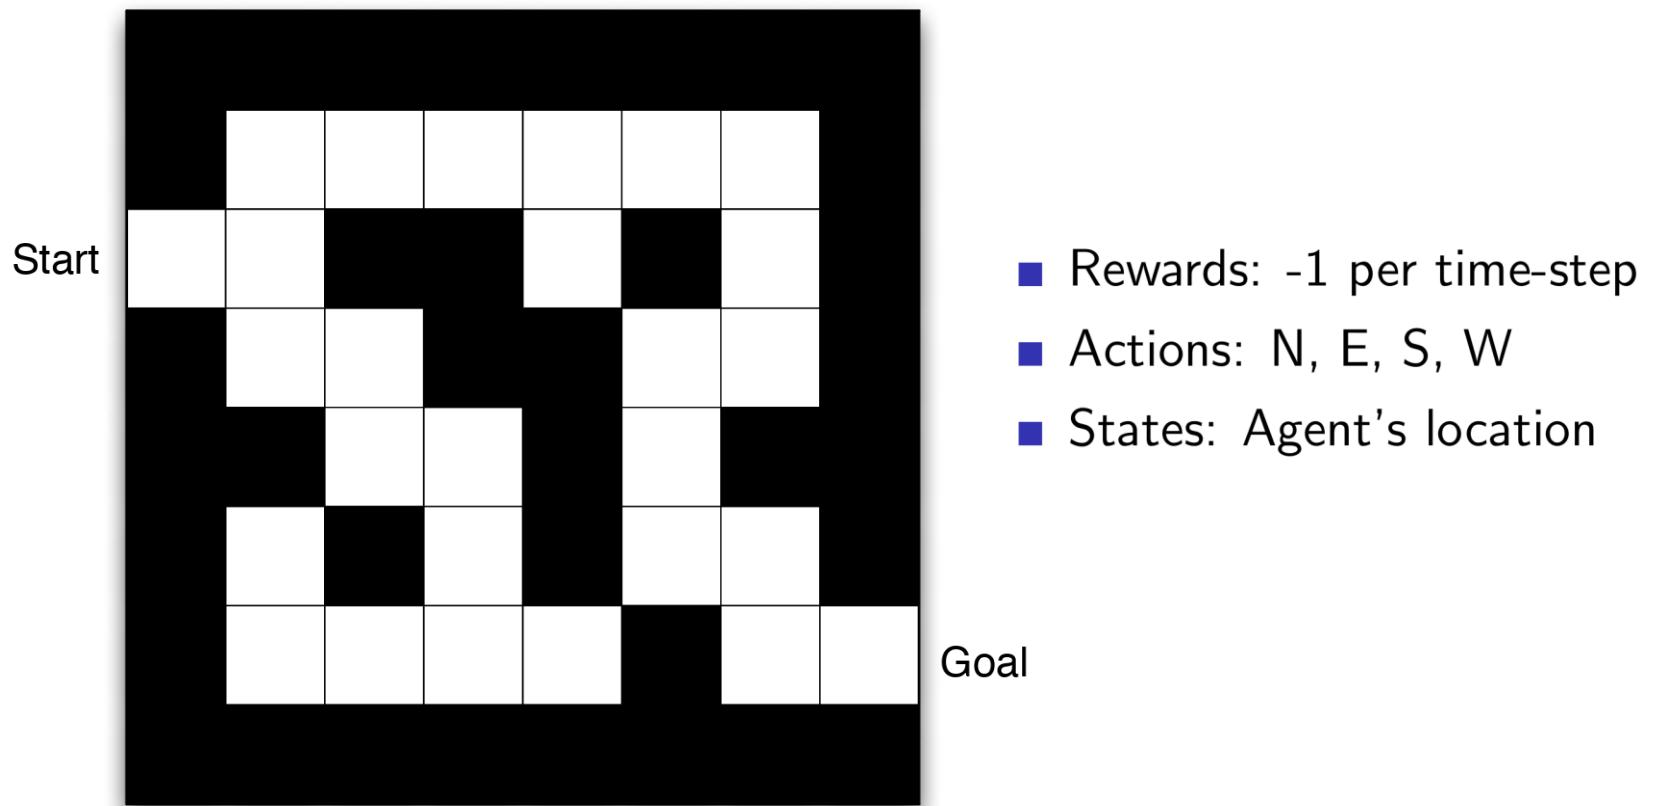
> Playing Atari with DQN



Major components of an RL agent

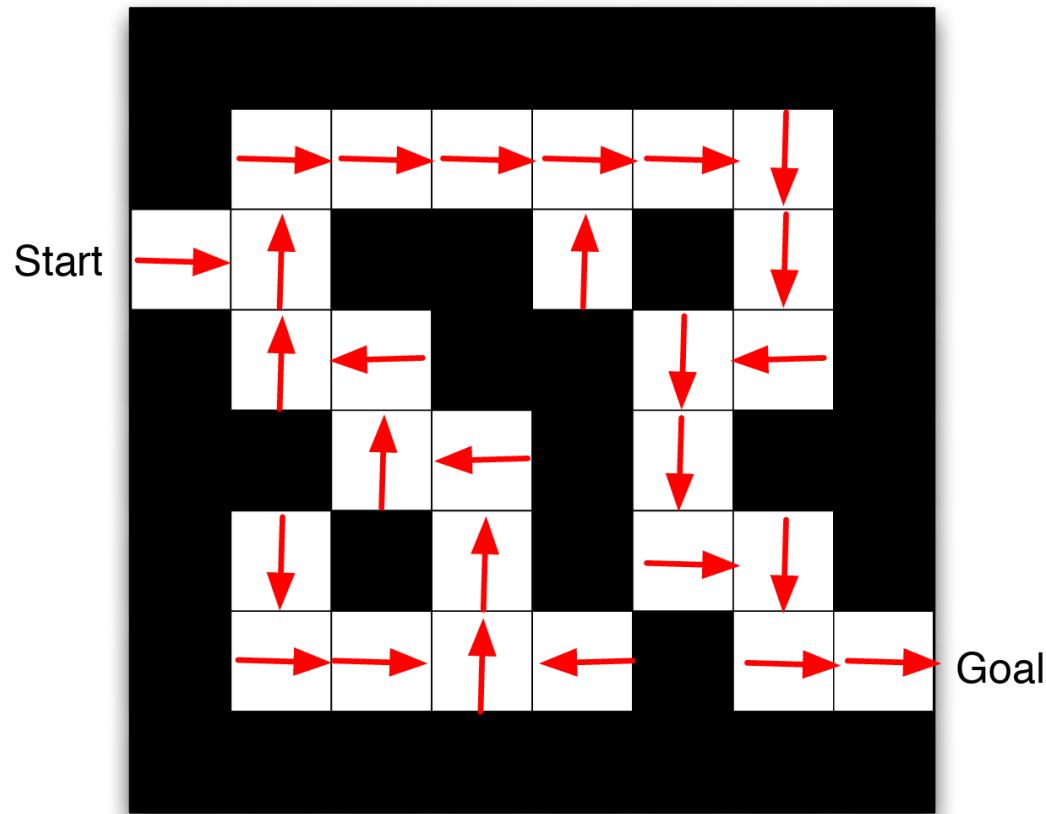
- > An RL agent may include those components:
 - policy: agent's behavior function
 - deterministic policy $A_t = \pi(s)$
 - stochastic policy $A_t = \pi(a|s) = p(A_t = a|S_t = s)$
 - value function: how good is each state and/or action
 - model: agent's representation of the environment
 - predict what the environments will do next
 - next state: $p(S_{t+1} = s' | S_t = s, A_t = a)$
 - next reward: $p(R_{t+1} | S_t = s, A_t = a)$

Maze example



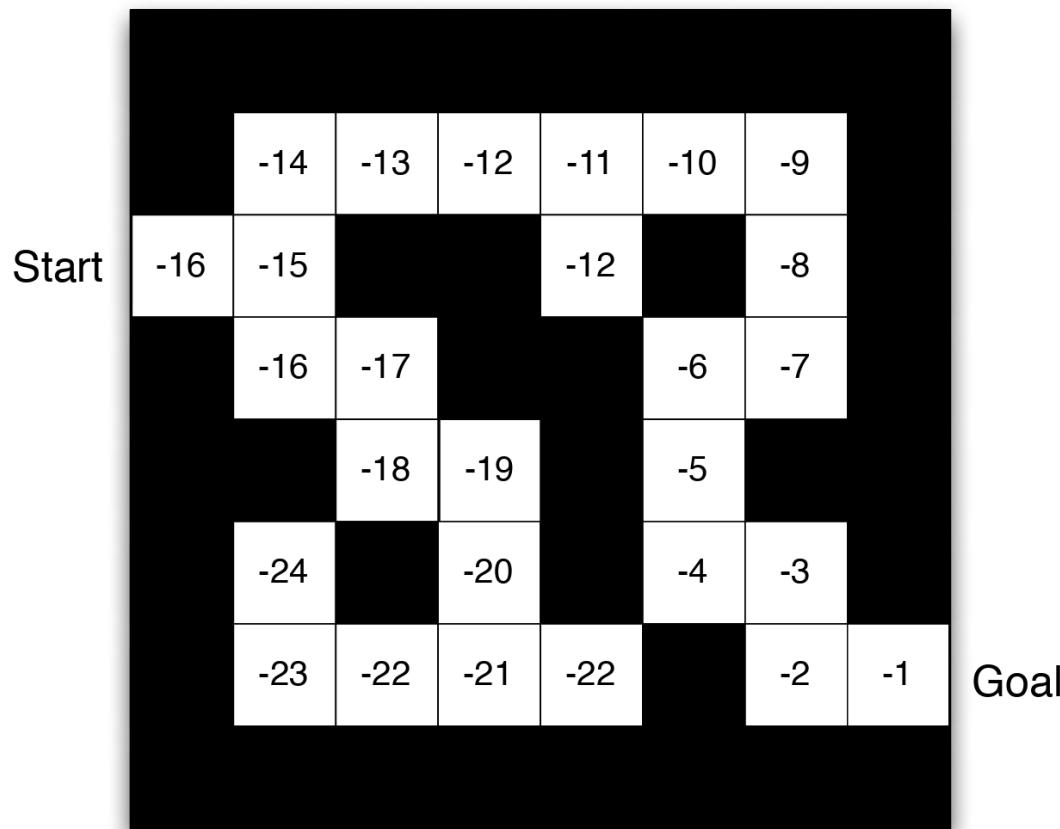
Maze example

> Policy



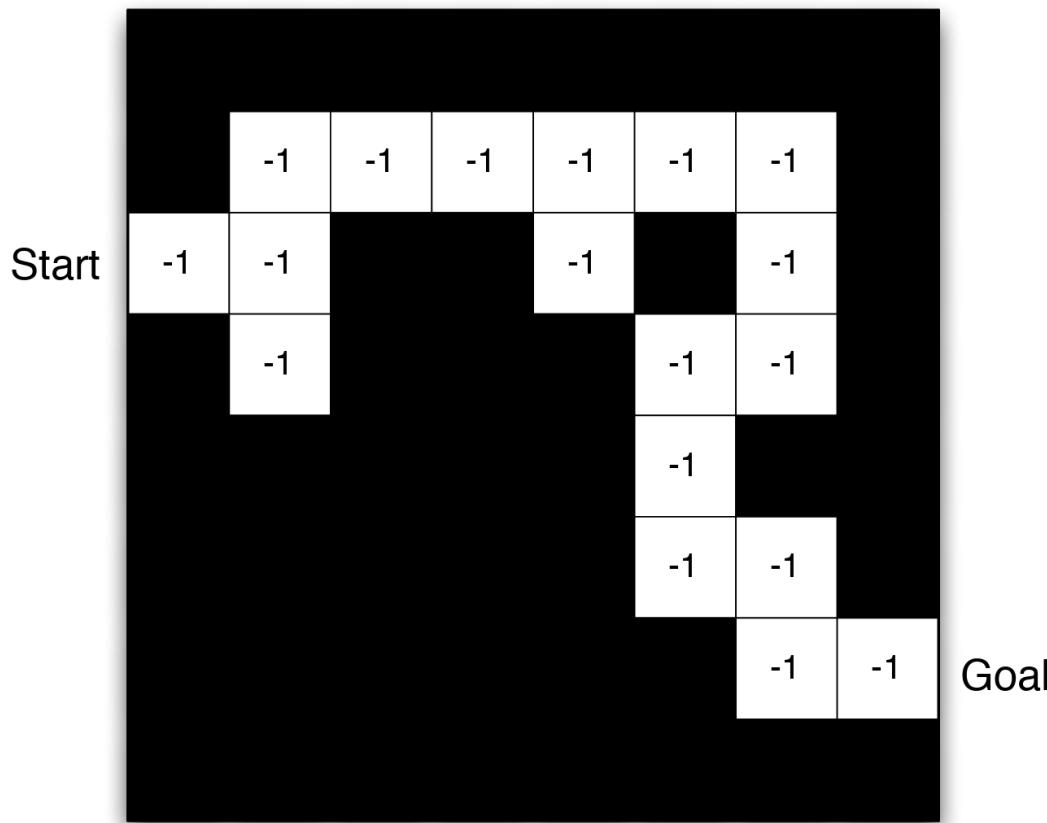
Maze example

> Value function



Maze example

> Model

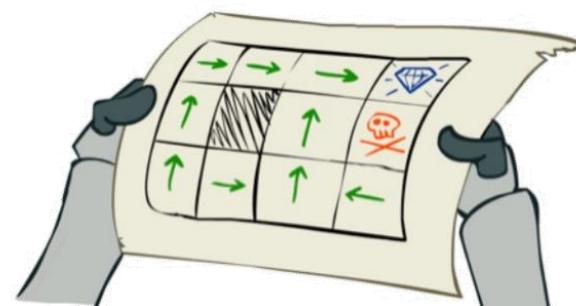
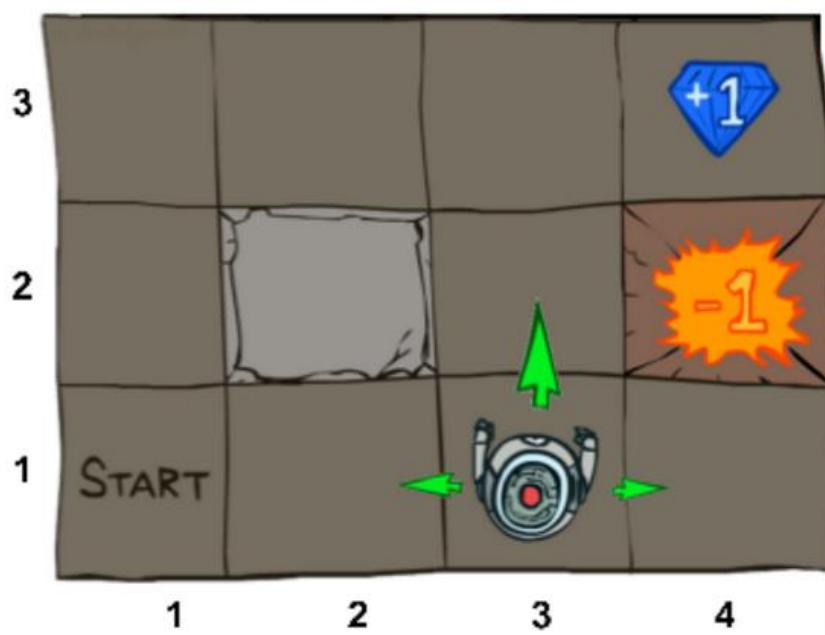


- Dynamics: how actions change the state
- Rewards: how much reward from each state
- Grid layout represents transition model
- Numbers represent immediate reward

Goal

The goal of RL

- > Finite horizon case: T is finite
- > Infinite horizon case: $T = \infty$
- > The cumulative reward is often discounted: $G_t = \sum_t \gamma^t r(s_t, a_t)$
- > Goal: find a policy to maximize the cumulative reward



Value functions

- > State-value function (V): the expected return starting from state s , and then following policy π
 - $V_\pi = \mathbb{E}_\pi[G_t | S_t = s]$
 - optimal: $V^*(s) = \max_\pi V_\pi(s)$
(maximum value function over all policies)
(the expected return when acting optimally)
- > Action-value function (Q): the expected return starting from state s , taking action a , and then following policy π
 - $Q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$
 - optimal: $Q^*(s, a) = \max_\pi Q_\pi(s, a)$
- > Value functions capture the knowledge of the agent regarding how good is each state for the goal the agent is trying to achieve.

Value functions

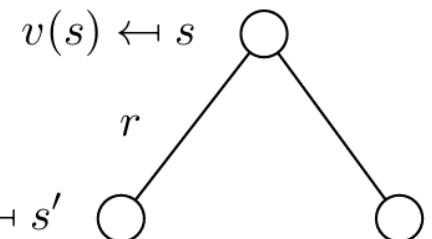
- > Once we obtain the optimal value function, we can find an optimal policy
- > Maximizing over $Q^*(s, a)$
 - $\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_a Q^*(s, a) \\ 0 & \text{otherwise} \end{cases}$
- > Maximizing over $V^*(s)$ with the model dynamics
 - $\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_a \sum p(s', r|s, a)(r + \gamma V^*(s')) \\ 0 & \text{otherwise} \end{cases}$
 - We need the dynamics to do one step lookahead to choose the optimal a

Roadmap

- > Computing state and state-action value functions by solving linear systems of equations
- > The matrix inversion is too costly
 - > iterative estimation is required (Bellman backup operation)
- > We cannot visit every state
 - > selective backups on state-actions that the agent visits
- > We may not know dynamics
 - > Monte Carlo learning or TD learning

Recursive relationships for returns

$$\begin{aligned}> G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots \\&= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \dots) \\&= R_{t+1} + \gamma G_{t+1}\end{aligned}$$



> By taking expectations

- $\mathbb{E}[G_t | S_t = s] = \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s]$
- $V_\pi(s) = \mathbb{E}[R_{t+1} + \gamma V_\pi(s') | S_t = s]$
 $= \sum_{s'} p(s', r | s)[r + \gamma V_\pi(s')]$ (Bellman expectation equation)

- For all states, $V_\pi = R_\pi + \gamma P_\pi V_\pi$
where v is a column vector with one entry per state

$$\begin{bmatrix} V(1) \\ \vdots \\ V(n) \end{bmatrix} = \begin{bmatrix} R_1 \\ \vdots \\ R_n \end{bmatrix} + \gamma \begin{bmatrix} P_{11} & \cdots & P_{1n} \\ \vdots & \ddots & \vdots \\ P_{n1} & \cdots & P_{nn} \end{bmatrix} \begin{bmatrix} V(1) \\ \vdots \\ V(n) \end{bmatrix}$$

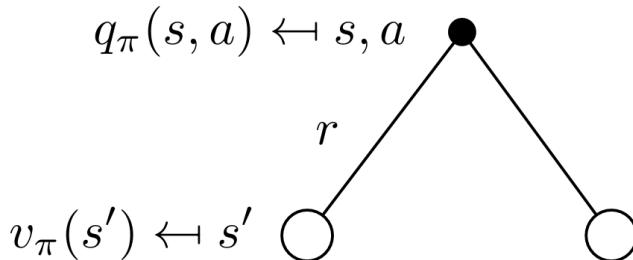
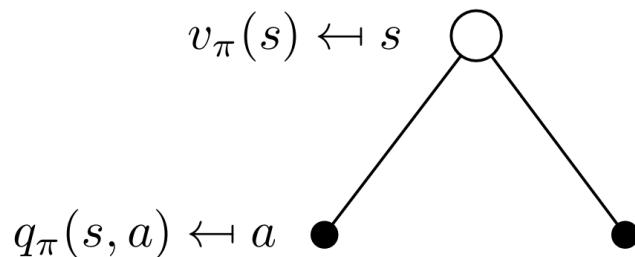
Solving the Bellman expectation equation

- > The Bellman expectation equation is a linear equation
- > It can be solved directly:
 - $$V_\pi = R_\pi + \gamma P_\pi V_\pi$$
$$(I - \gamma P_\pi)V_\pi = R_\pi$$
$$V_\pi = (I - \gamma P_\pi)^{-1}R_\pi$$
- > Computational complexity is $O(n^3)$ for n states
- > There are many iterative methods for large state system
 - Dynamic programming
 - Monte-Carlo evaluation
 - Temporal-Difference learning

Relating state and state-action value functions

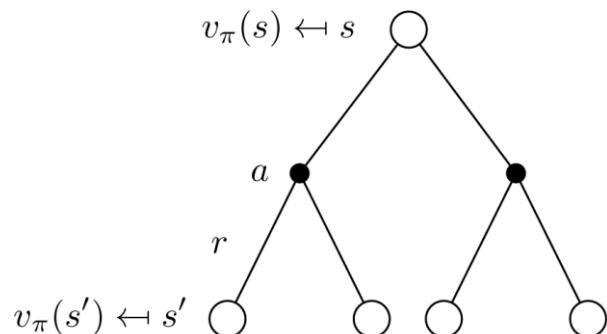
- > The action-value function can similarly be decomposed

$$Q_\pi(s, a) = \mathbb{E}[R_{t+1} + \gamma Q_\pi(s', a') | S_t = s, A_t = a]$$



$$V_\pi(s) = \sum_a \pi(a|s) Q_\pi(s, a)$$

$$Q_\pi(s, a) = \sum_{s'} p(s', r|s, a) [r + \gamma V_\pi(s')]$$



$$V_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_\pi(s')]$$

$$Q_\pi(s, a) = \sum_{s'} p(s', r|s, a) [r + \gamma \sum_a \pi(a'|s') Q_\pi(s', a')]$$

Bellman optimality equations

> For the Bellman expectation equations, we sum over all the possibilities.

> Now, we choose only the best action.

- $V_\pi(s) = \sum_a \pi(a|s) Q_\pi(s, a) \rightarrow V^* = \max_a Q^*(s, a)$

- $Q_\pi(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma V_\pi(s')]$

$$\rightarrow Q^*(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma V^*(s')]$$

- $V_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a)[r + \gamma V_\pi(s')]$

$$\rightarrow V^*(s) = \max_a \sum_{s'} p(s', r|s, a)[r + \gamma V^*(s')]$$

- $Q_\pi(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma \sum_a \pi(a'|s') Q_\pi(s', a')]$

$$\rightarrow Q^*(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma \max_a Q^*(s', a')]$$

Solving the Bellman optimality equation

- > Bellman optimality equation is nonlinear
- > No closed form solution (in general)
- > Many iterative solution methods
 - Using models / dynamic programming
 - Value iteration
 - Policy iteration
 - Using samples
 - Monte Carlo
 - Q-learning
 - SARSA

Extensions to MDPs

- > Infinite and continuous MDP
- > Continuous state and/or action spaces
 - closed form for linear quadratic model (LQR)
- > Continuous time
 - requires partial differential equations
 - Hamilton-Jacobi-Bellman (HJB) equation

Planning by dynamic programming

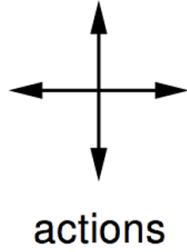
- > Dynamic programming assumes full knowledge of the MDP
- > It is used for planning in an MDP
- > For prediction:
 - input: MDP and policy
 - output: value function V_π
- > For control:
 - input: MDP
 - output: optimal value function V^* and optimal policy π^*

Policy Evaluation

- > Problem: evaluate a given policy π (prediction)
- > Solution: iterative application of Bellman expectation backup
 - $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$
 - Synchronous backups – for all states
- > Iterative policy evaluation
- >
$$V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_k(s')]$$

Policy Evaluation

- > Evaluating a random policy in the small GridWorld



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$
on all transitions

- undiscounted episodic MDP ($\gamma = 1$)
- one terminal state (shown twice as shaded squares)
- actions leading out of the grid leave state unchanged
- reward is -1 until the terminal state is reached
- agent follows uniform random policy
(each action has a probability of 0.25)

Policy Evaluation

$$> V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_k(s')]$$

v_k for the
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

Policy Evaluation

> Overall algorithm

Input π , the policy to be evaluated

Initialize an array $V(s) = 0$, for all $s \in \mathcal{S}^+$

Repeat

$\Delta \leftarrow 0$

 For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

 until $\Delta < \theta$ (a small positive number)

Output $V \approx v_\pi$

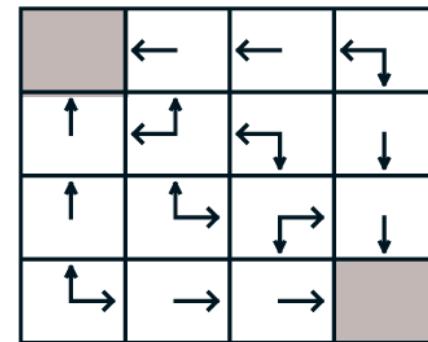
> It will converge to the fixed value function

Policy Evaluation

- Once we found the converged (optimal) value function, we can get a better (optimal) policy

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



Policy Evaluation

- > Convergence proof
 - Define the Bellman expectation backup operator
 - $T_\pi(V) = R_\pi + \gamma P_\pi V$
 - This operator is a γ -contraction; it makes value functions closer by at least γ
 - $$\begin{aligned} \|T_\pi(U) - T_\pi(V)\|_\infty &= \|R_\pi + \gamma P_\pi U - (R_\pi + \gamma P_\pi V)\|_\infty \\ &= \|\gamma P_\pi(U - V)\|_\infty \\ &\leq \gamma \|P_\pi\|_\infty \|U - V\|_\infty \\ &= \gamma \|U - V\|_\infty \end{aligned}$$
 - T_π converges to a unique fixed point at a linear convergence rate γ

Policy iteration

- > Now, we want to move to the control problem
- > Policy iteration
 - Evaluate the policy
 - Improve the policy by acting greedily with respect to V_π

$$\pi' = \text{greedy}(V_\pi)$$

- It always converges to optimal policy π^*

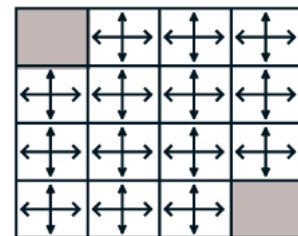
Policy iteration

v_k for the
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

Greedy Policy
w.r.t. v_k

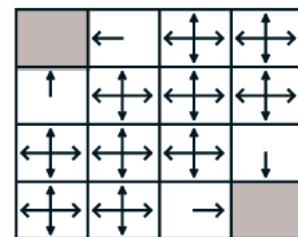


$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

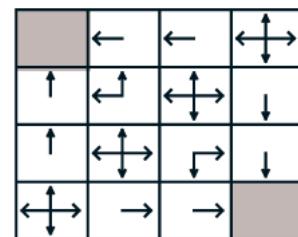
$k = 2$

0.0	-1.0	-2.0	-2.0
-1.0	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.0
-2.0	-2.0	-1.0	0.0

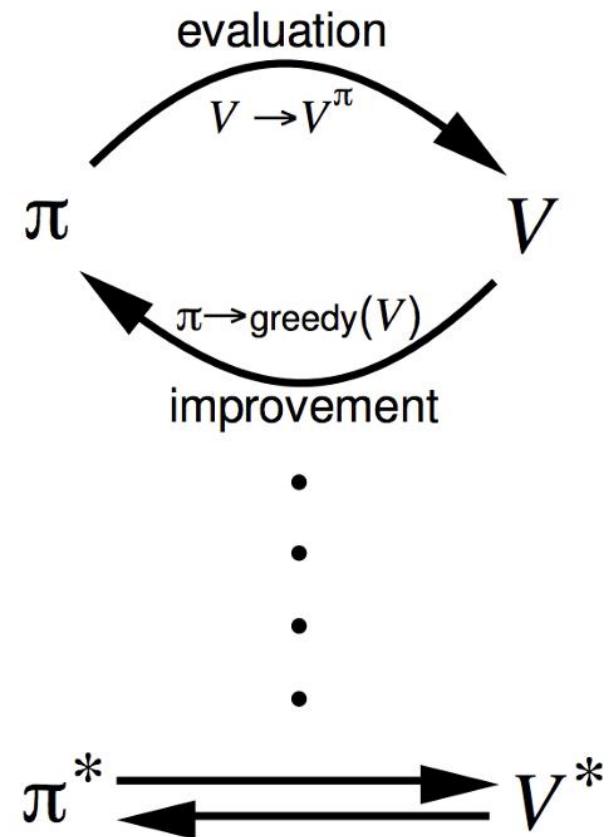
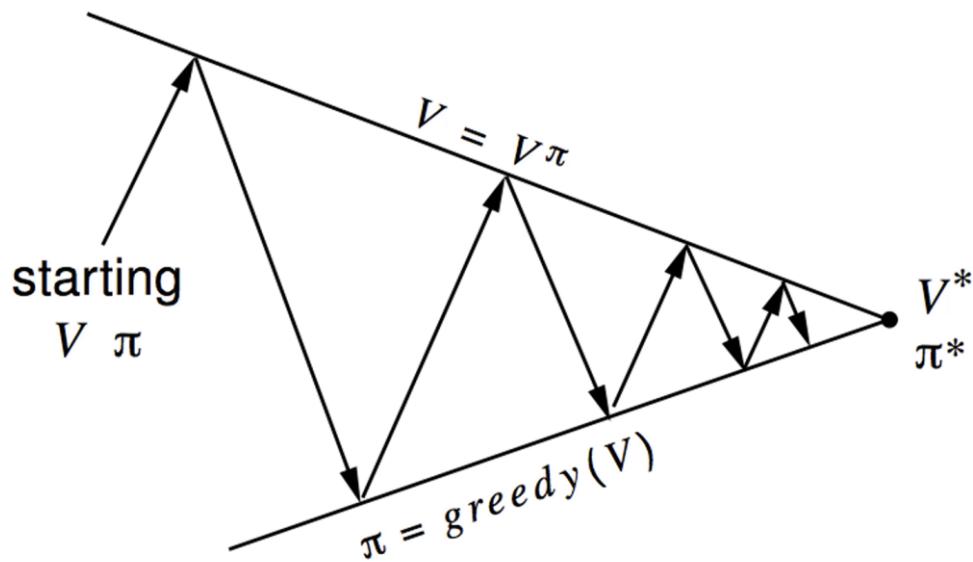


$k = 3$

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0



Policy iteration



Policy iteration

- > Consider a deterministic policy, $a = \pi(s)$
- > Improve the policy by acting greedily, $\pi'(s) = \operatorname{argmax}_a Q_\pi(s, a)$
- > It improves the value function

- $$Q_\pi(s, \pi'(s)) = \max_a Q_\pi(s, a) \geq Q_\pi(s, \pi(s)) = V_\pi(s)$$
- $$\begin{aligned} V_\pi(s) &\leq Q_\pi(s, \pi'(s)) = \mathbb{E}_{\pi'}[R_{t+1} + \gamma V_\pi(s') | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma Q_\pi(s', \pi'(s')) | S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma \mathbb{E}_{\pi'}[R_{t+2} + \gamma V_\pi(s'')] | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 Q_\pi(s', \pi'(s'')) | S_t = s] \\ &\quad \cdots \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma R_{t+3} + \cdots | S_t = s] = \\ &V_{\pi'}(s) \end{aligned}$$

Policy iteration

- > If improvements stop, $Q_\pi(s, \pi'(s)) = V_\pi(s)$
- > Then, the Bellman optimality equation has been satisfied

$$V^* = \max_a Q^*(s, a)$$

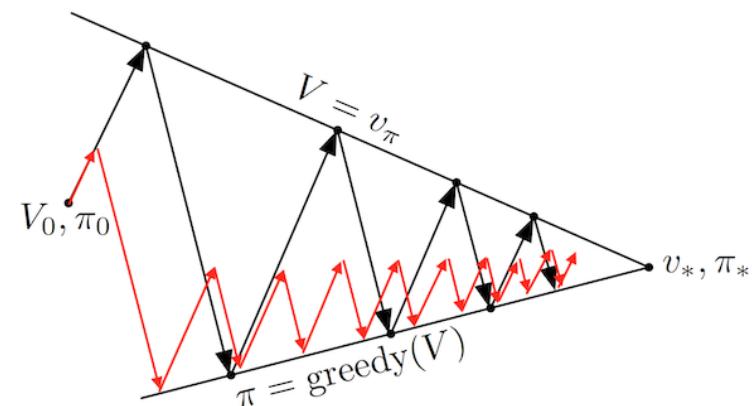
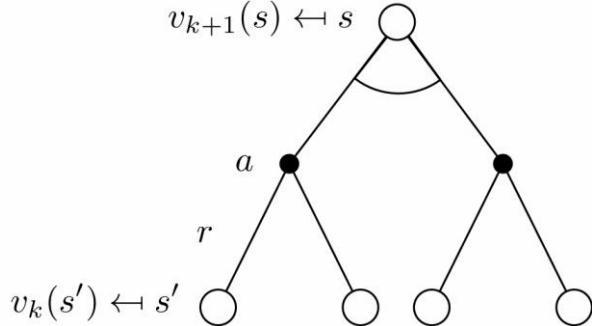
- > Therefore, V and π are optimal

Generalized policy iteration

- > Any interleaving of policy evaluation and policy improvement
 - > Evaluation and improvement do not need to be exact or complete at each step
 - partial updates, approximations, or asynchronous updates are allowed
 - > All RL methods are a form of GPI
-
- > Policy iteration: full evaluation + full improvement
 - using Bellman expectation eqn.
 - > Value iteration: one-step evaluation
 - using Bellman optimality eqn.
-
- > How GPI leads optimality?
 - Over time, even with approximate steps the policy becomes progressively better

Value iteration

- > Problem: find optimal policy π
- > Solution: iterative application of Bellman optimality backup
- > Using synchronous backups
- > Unlike policy iteration, there is no explicit policy
- >
$$V_{k+1}(s) = \max_a [R(s, a) + \gamma \sum_{s'} p(s', r|s, a)V_k(s')]$$



Value iteration

$$> V_{k+1}(s) = \max_a [R(s, a) + \gamma \sum_{s'} p(s', r|s, a)V_k(s')]$$

g				

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

V_1

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

V_2

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

V_3

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

V_4

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

V_5

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

V_6

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

V_7

Synchronous dynamic programming algorithms

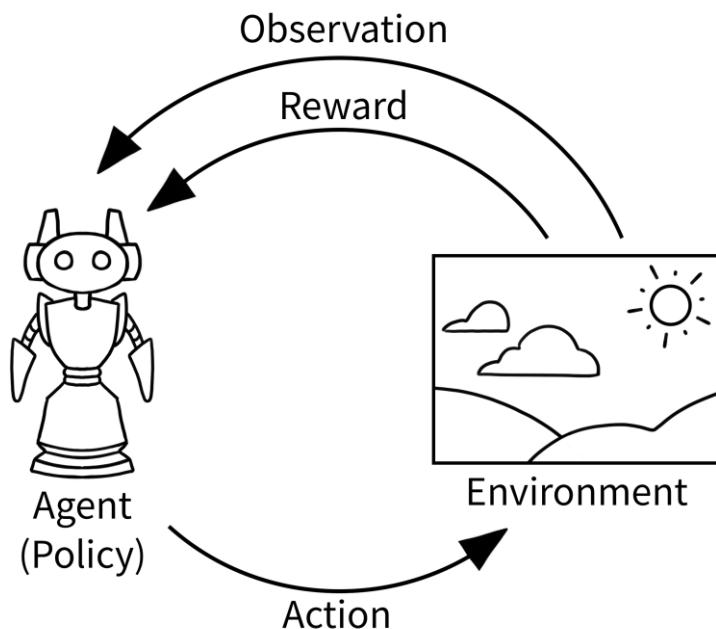
Problem	Bellman Equation	Algorithm
Prediction	Bellman expectation equation	Iterative policy evaluation
Control	Bellman expectation equation +Greedy policy improvement	Policy iteration
Control	Bellman optimality equation	Value iteration

- > Algorithms are based on state-value function $V_\pi(s)$ or $V^*(s)$
- > Complexity $O(mn^2)$ per iteration, for m actions and n states
- > Could also apply to action-value function $Q_\pi(s, a)$ or $Q^*(s, a)$
- > Complexity $O(m^2n^2)$ per iteration

Environment for RL

> Gymnasium

- Provide an API for all single agent RL envs.
- Maintained by OpenAI



```
env = gym.make  
observation, info = env.reset  
action = env.action_space.sample()  
observation, reward, terminated, truncated,  
info = env.step(action)
```

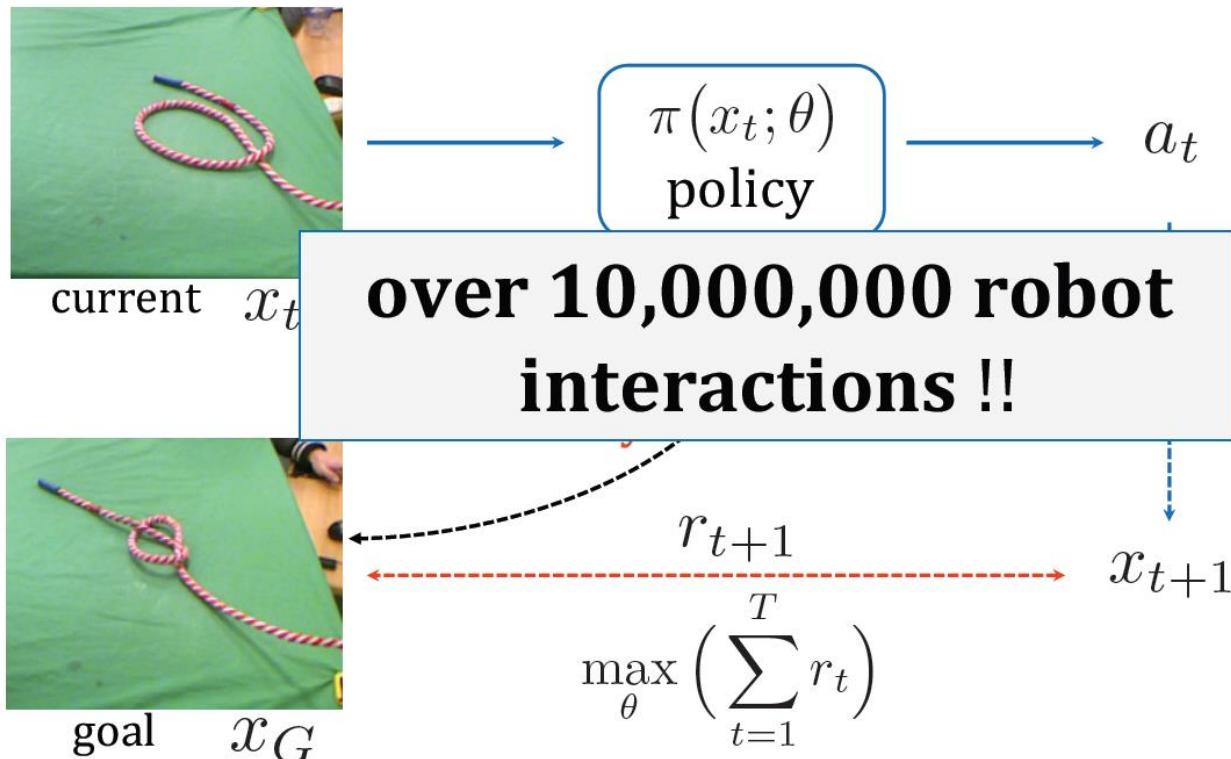
IssacSim

- > Developed by NVIDIA
 - End-to-end GPU accelerated
 - Massive parallelization of thousands of environments
 - PhysX 5 physics engine, including fluid dynamics



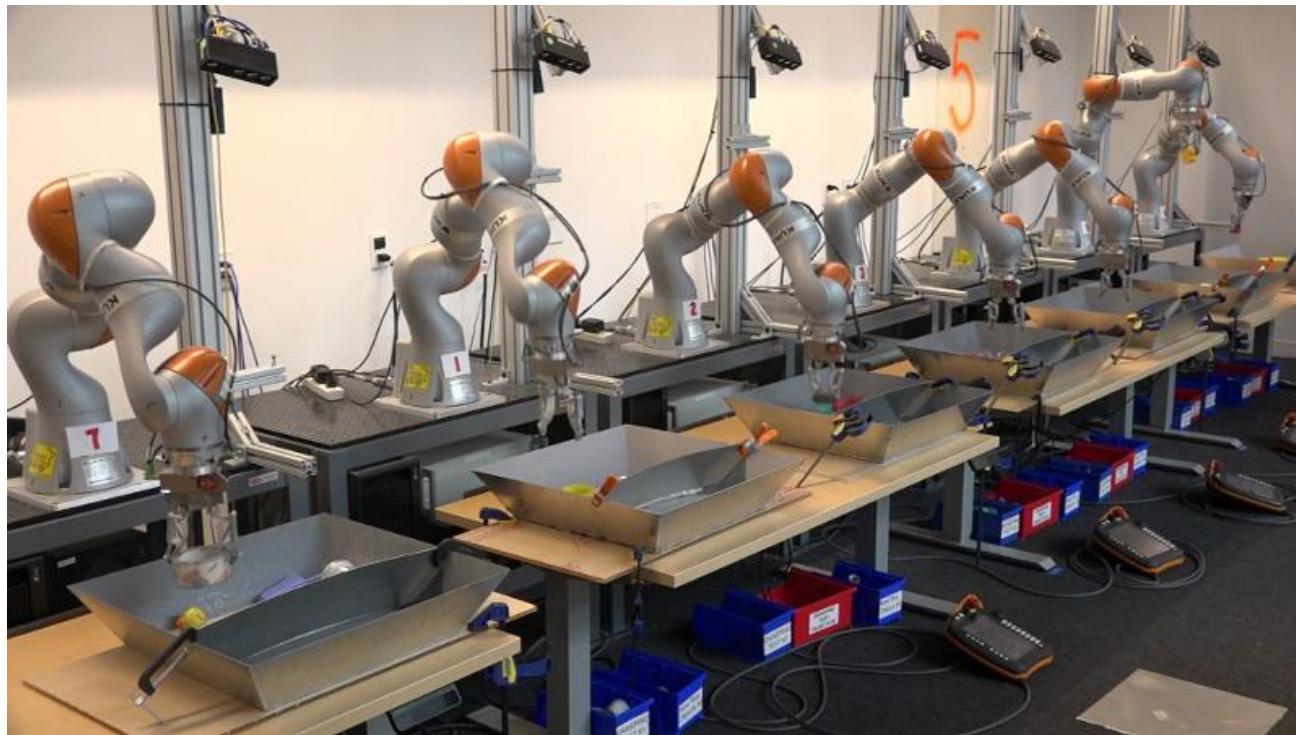
Why simulators for robot learning

- > RL is very sample inefficient



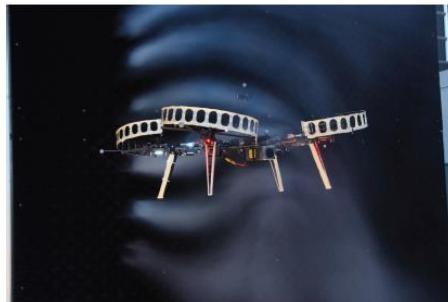
Why simulators for robot learning

- > Google QT-Opt
 - 4 months, 800 robot hours, 7 robots, 580,000 attempts

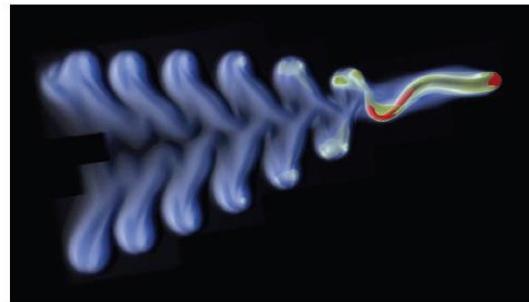


Why simulators for robot learning

- > Advantages of using simulated data
 - Cheap, fast, and scalable
 - Safe
 - Labeled (we have access to ground truth)
 - No physical harm or deformation (wear and tear)
- > Disadvantages
 - Not exactly same with real env. (sim2real)



Aerodynamics in wind, *Neural-Fly*



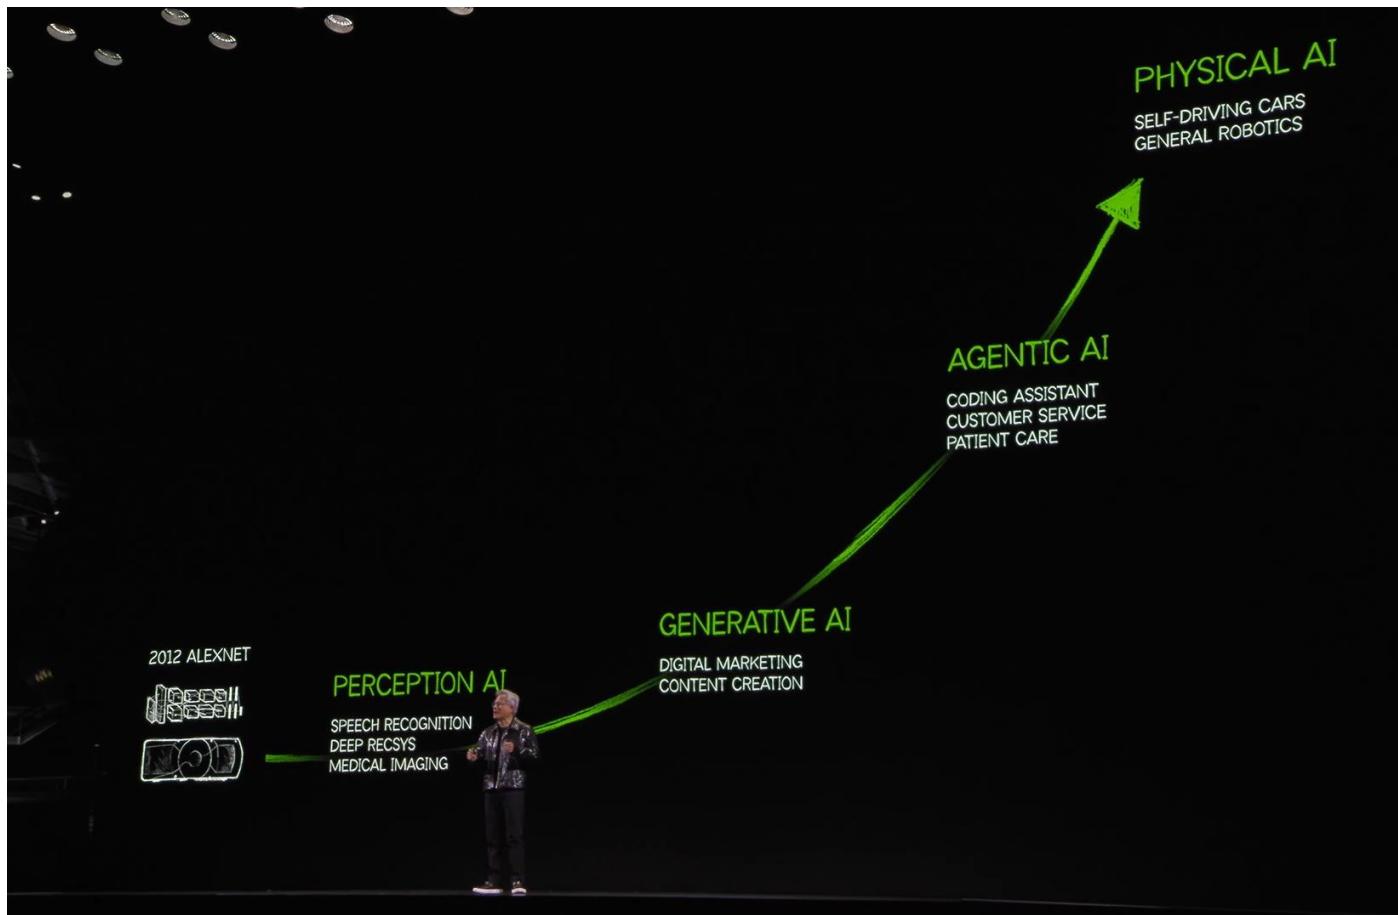
Fluid dynamics, MIT van Rees Lab



Offroad vehicle dynamics, UW Racer team

Ultimate goal

- > Build general-purpose embodied intelligence by learning to make sequential decisions in the physical world.



Where are we today: non-learning method

- > trajectory optimization and control: optimal control + robust control



Where are we today: non-learning method

- > trajectory optimization + MPC



Where are we today: learning method

- > Sim2Real - NVIDIA



Where are we today: learning method

- > Collect real-world data efficiently – Mobile ALOHA



Where are we today: learning method

- > Control foundation model: A general navigation model (GNM)

